

GStreamer-integrated HLS-based JPEG Encoder for Edge FPGA SoCs

Yuri Guimaraes Pereira Primo da Silva¹, Shinya Honda¹,
Sugako Otani¹², Masato Edahiro¹, Abraham Monrroy Cano³

¹Nagoya University,
²Renesas Electronics, ³Map IV Inc.

Motivation

1. Lack of freely available image encoder designed for edge use case.
- 2.

Approach

Architecture

Results & Comparisons

Future Work

GStreamer-integrated HLS-based JPEG Encoder for Edge FPGA SoCs

Yuri Guimaraes Pereira Primo da Silva¹, Shinya Honda¹, Sugako Otani^{1,2}, Masato Eda¹, Abraham Monrroy Cano³

¹Nagoya University, ²Renesas Electronics, ³Map IV Inc.

E-mail: yuri_gpps1@ert1.jp

Motivation

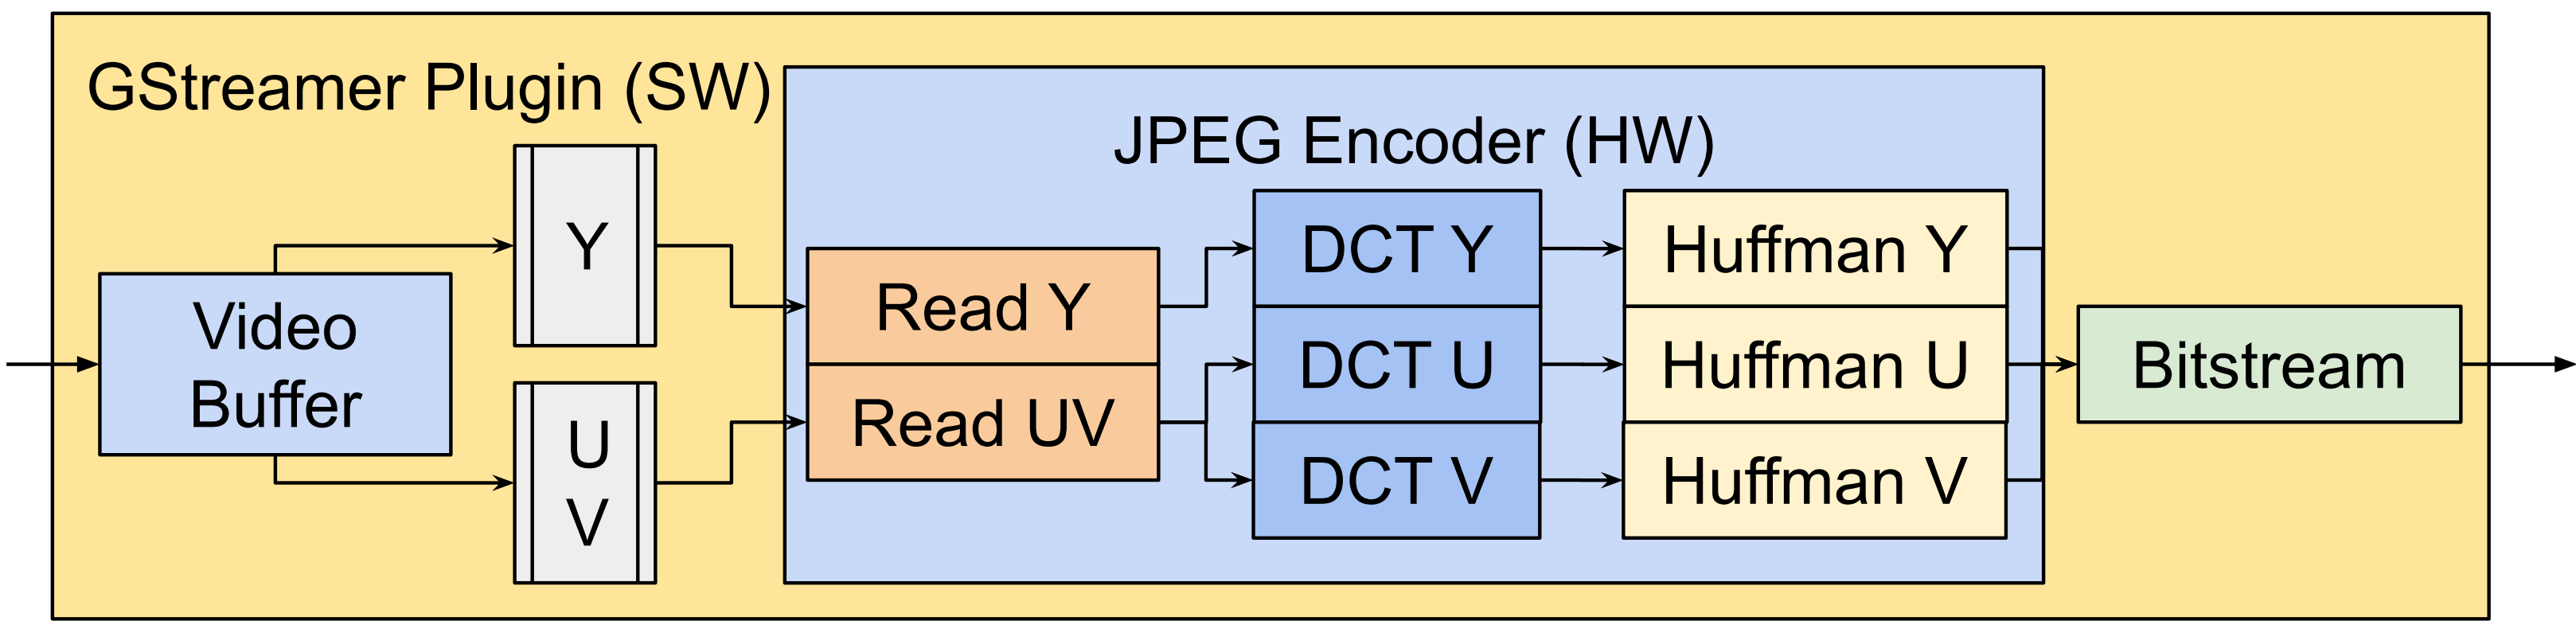
- 1. Lack of freely available image encoder designed for edge use case.
- 2.

Approach

JPEG HLS Encoder

GStreamer Plugin Integration

Architecture



Experiments

Future Work

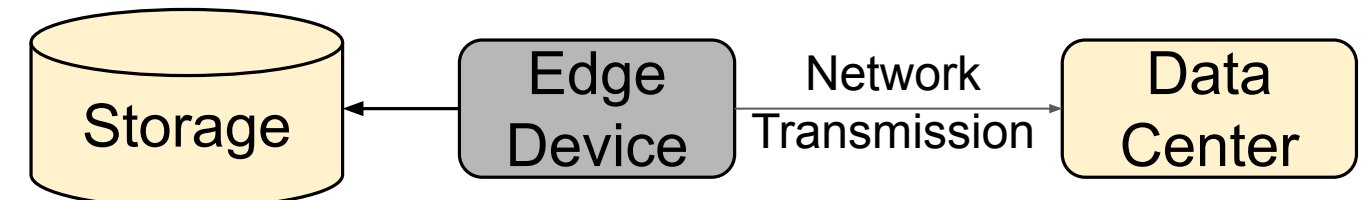
GStreamer-integrated HLS-based JPEG Encoder for Edge FPGA SoCs

Yuri Guimaraes Pereira Primo da Silva¹, Shinya Honda¹,
Sugako Otani¹², Masato Eda¹, Abraham Monrroy Cano³

¹Nagoya University, ²Renesas Electronics,
³Map IV Inc.

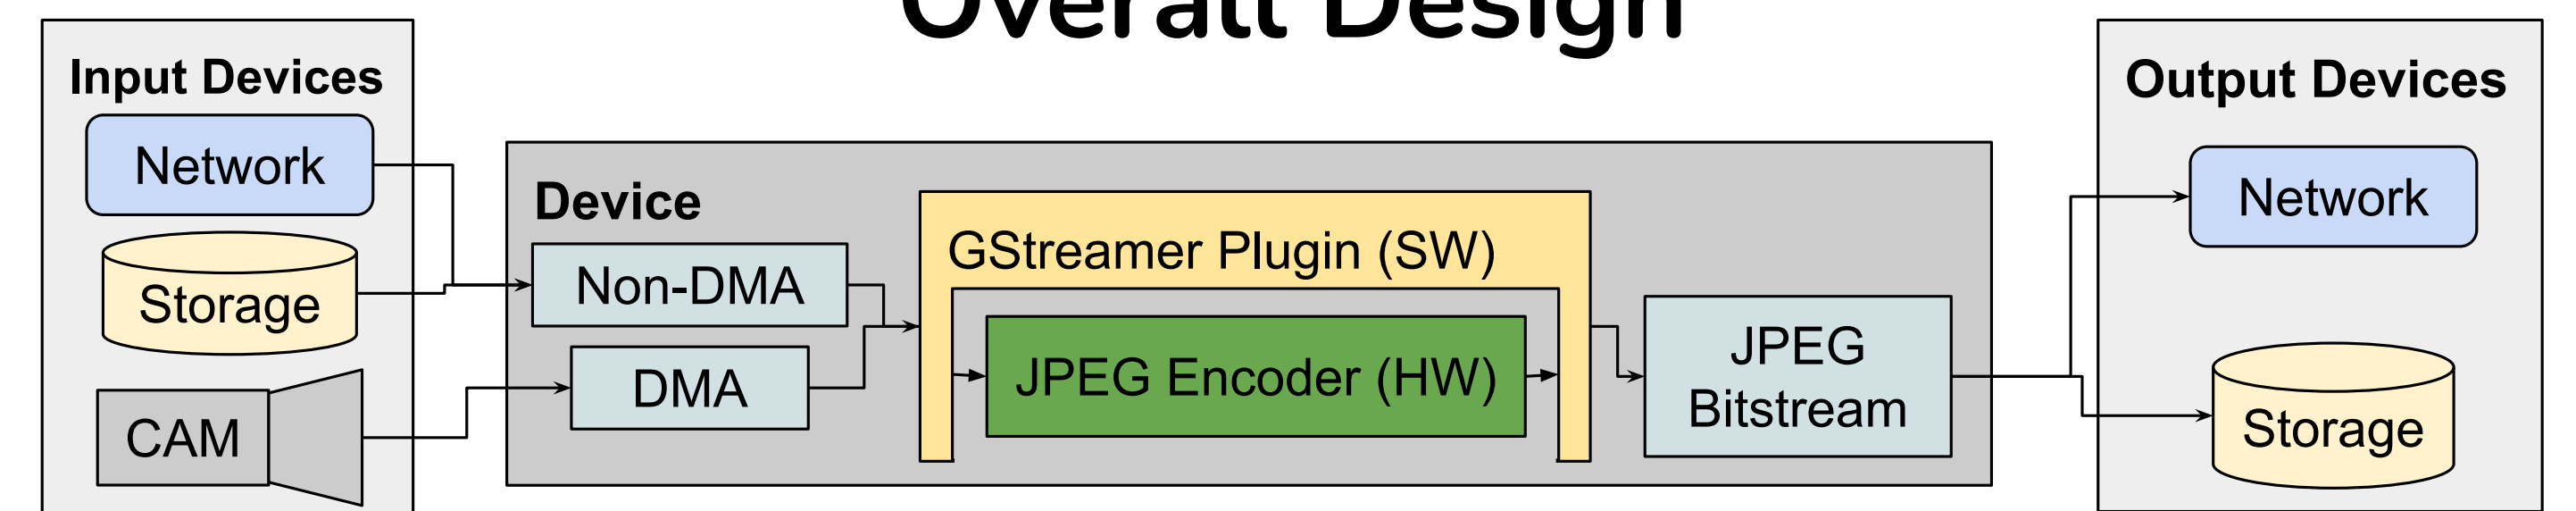
E-mail: yuri_gpps1@ert1.jp

Motivation

- 
- Increasing importance of image encoding at the edge.
 - Image encoding** efforts have been mostly concentrated on accelerator cards.
 - This research focuses on the implementation and integration aspect of an **HLS-based JPEG encoder** on Kria KV260 SoC [1] integrated with **GStreamer**.

Approach

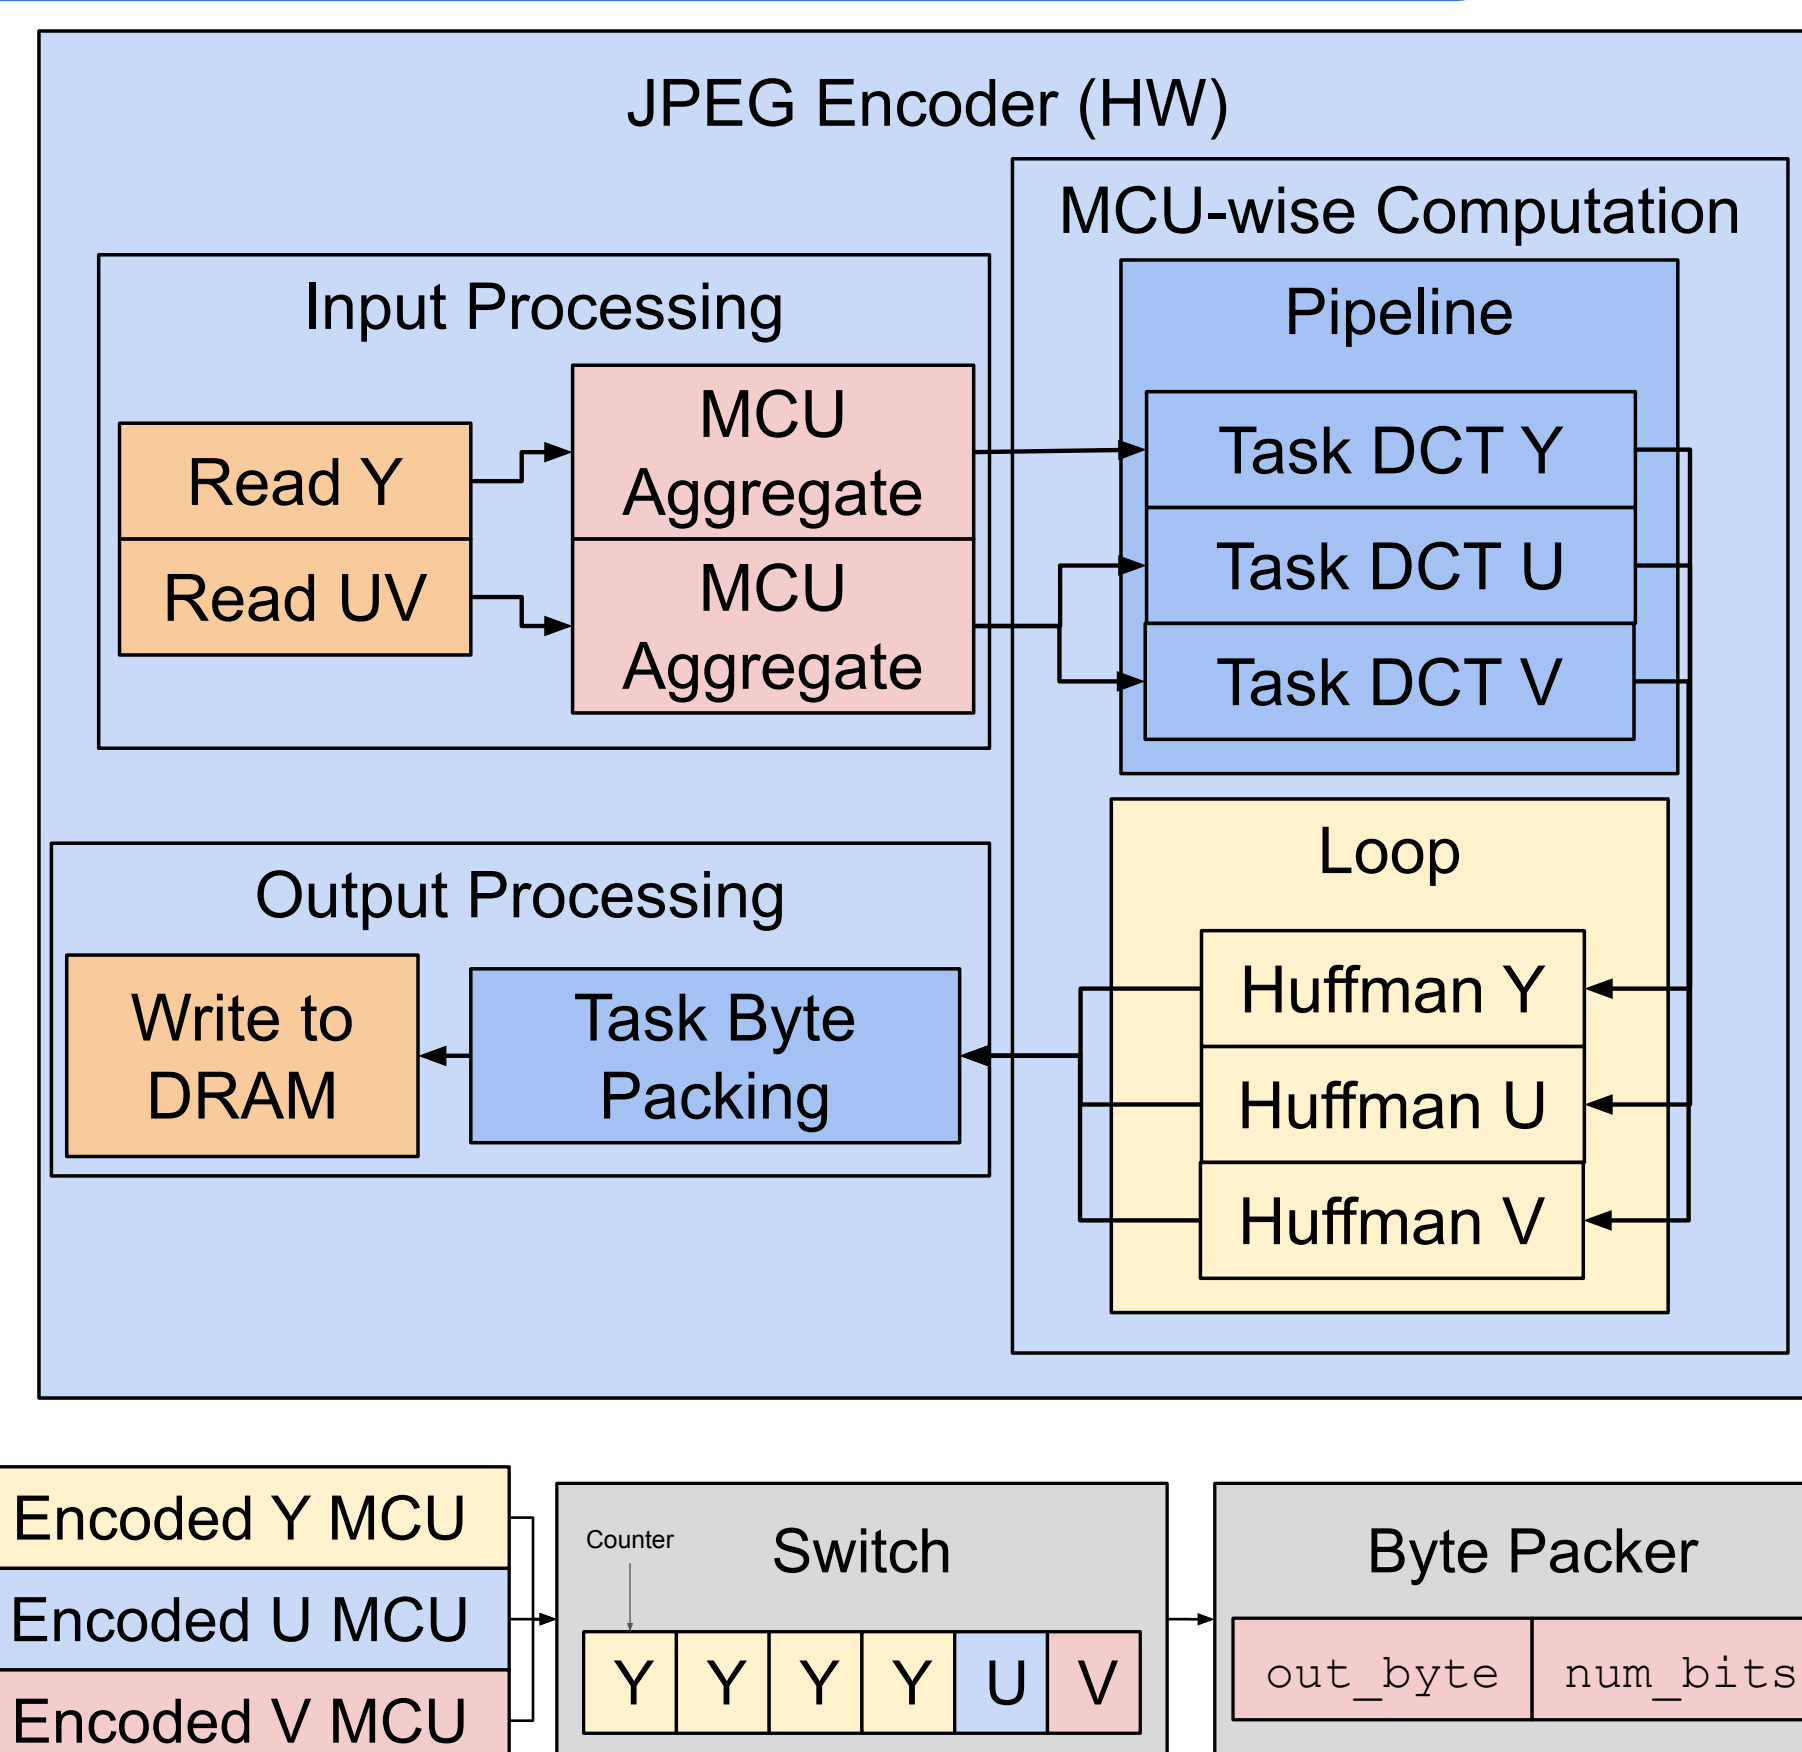
Overall Design



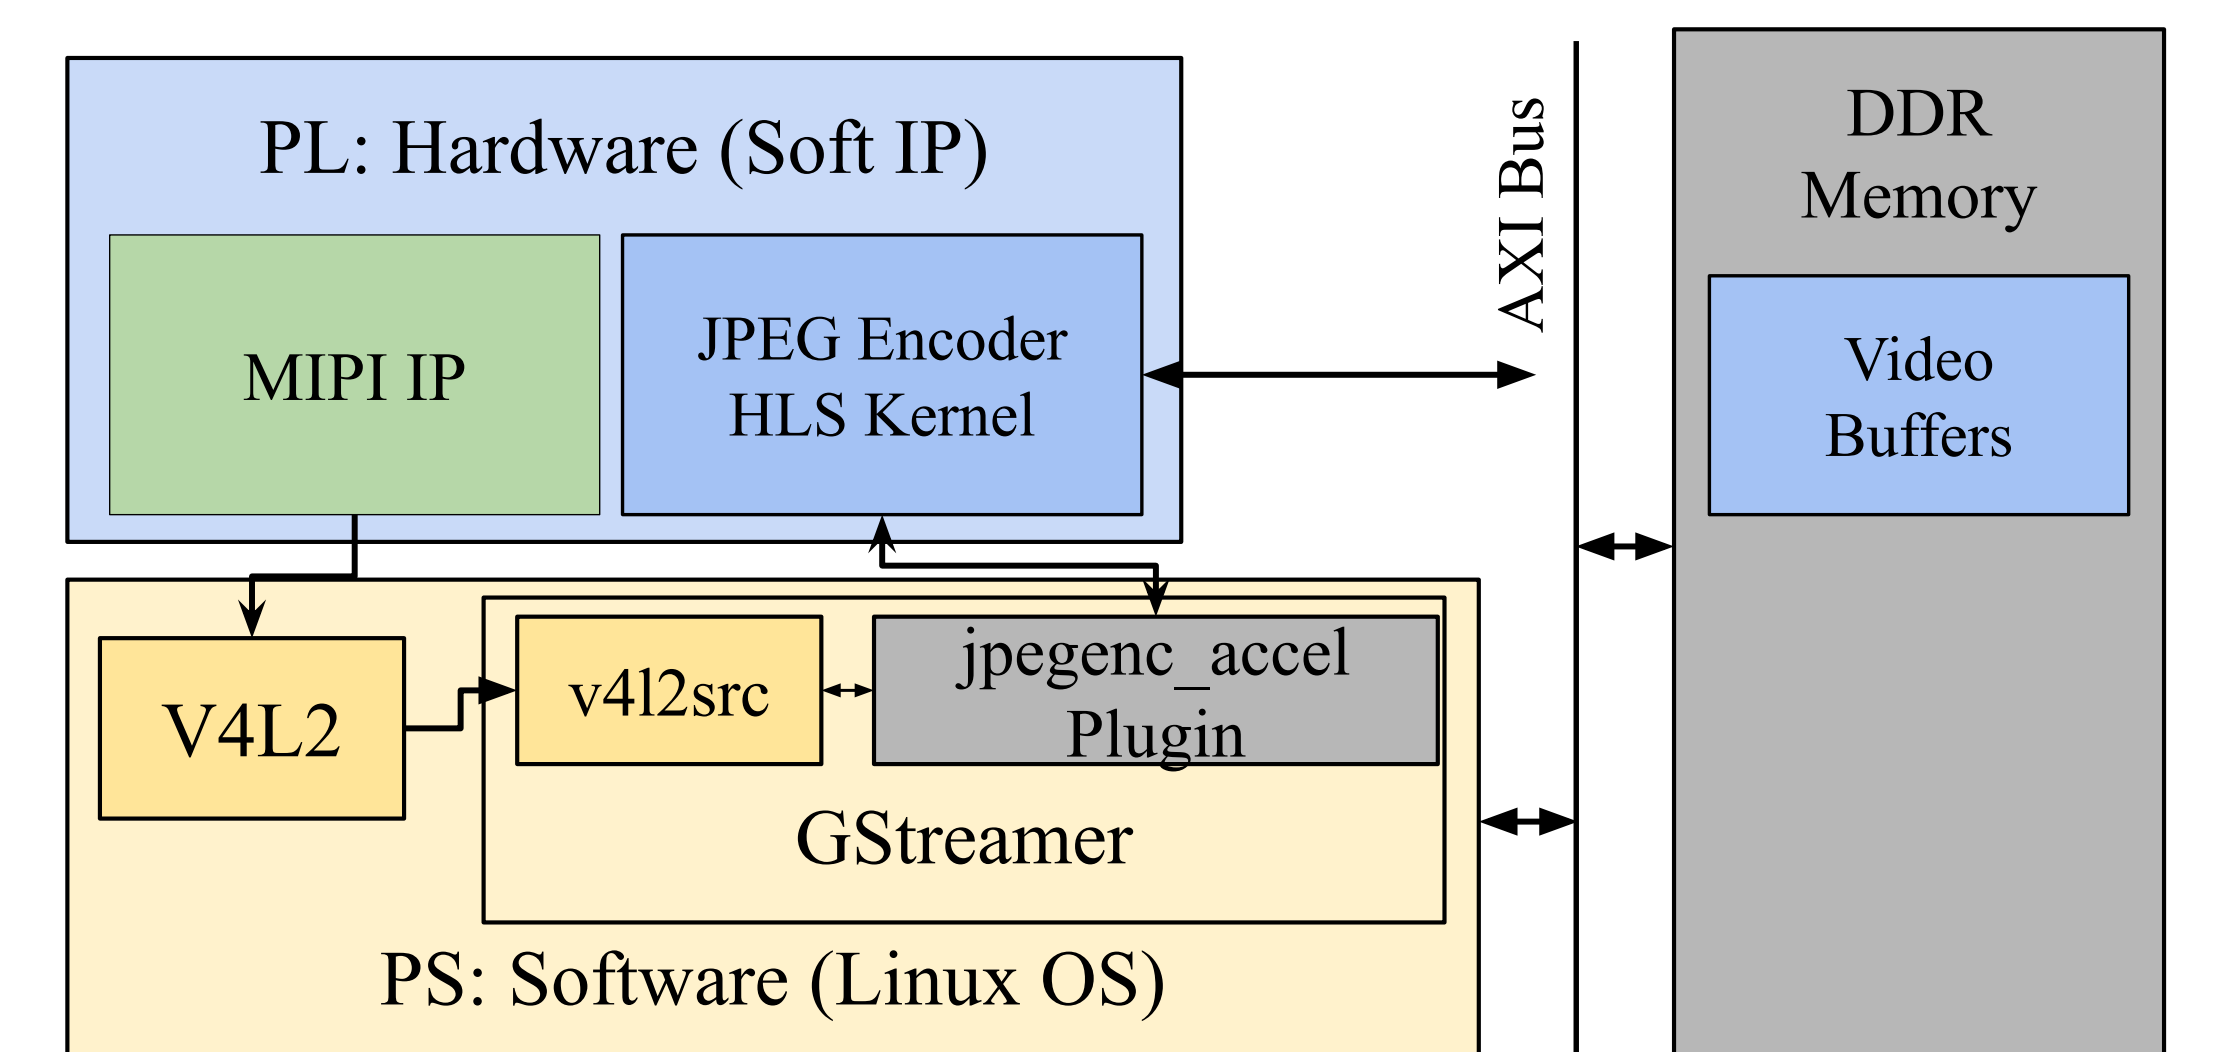
- HW JPEG Encoder implemented in *High-level Synthesis* [2].
- GStreamer plugin interfaces with encoder.
 - Allows for mix and matching of pipeline elements for arbitrary media processing.
 - Easy usage of DMA capabilities.

Hardware Architecture

- Conventional JPEG algorithm design implemented C++ HLS.
- Streaming design:
 - Data is packed into MCU streams at input processing.
 - Each MCU in a stream goes through pipelined DCT and RLE pipelined, and sequential Huffman.
 - MCUs are packed into the expected JPEG bitstream format at the end.
- Encoded bits are packed via a stateful switch that shifts and ORs according to the encoding order (e.g. YYYYUV for YUV420).



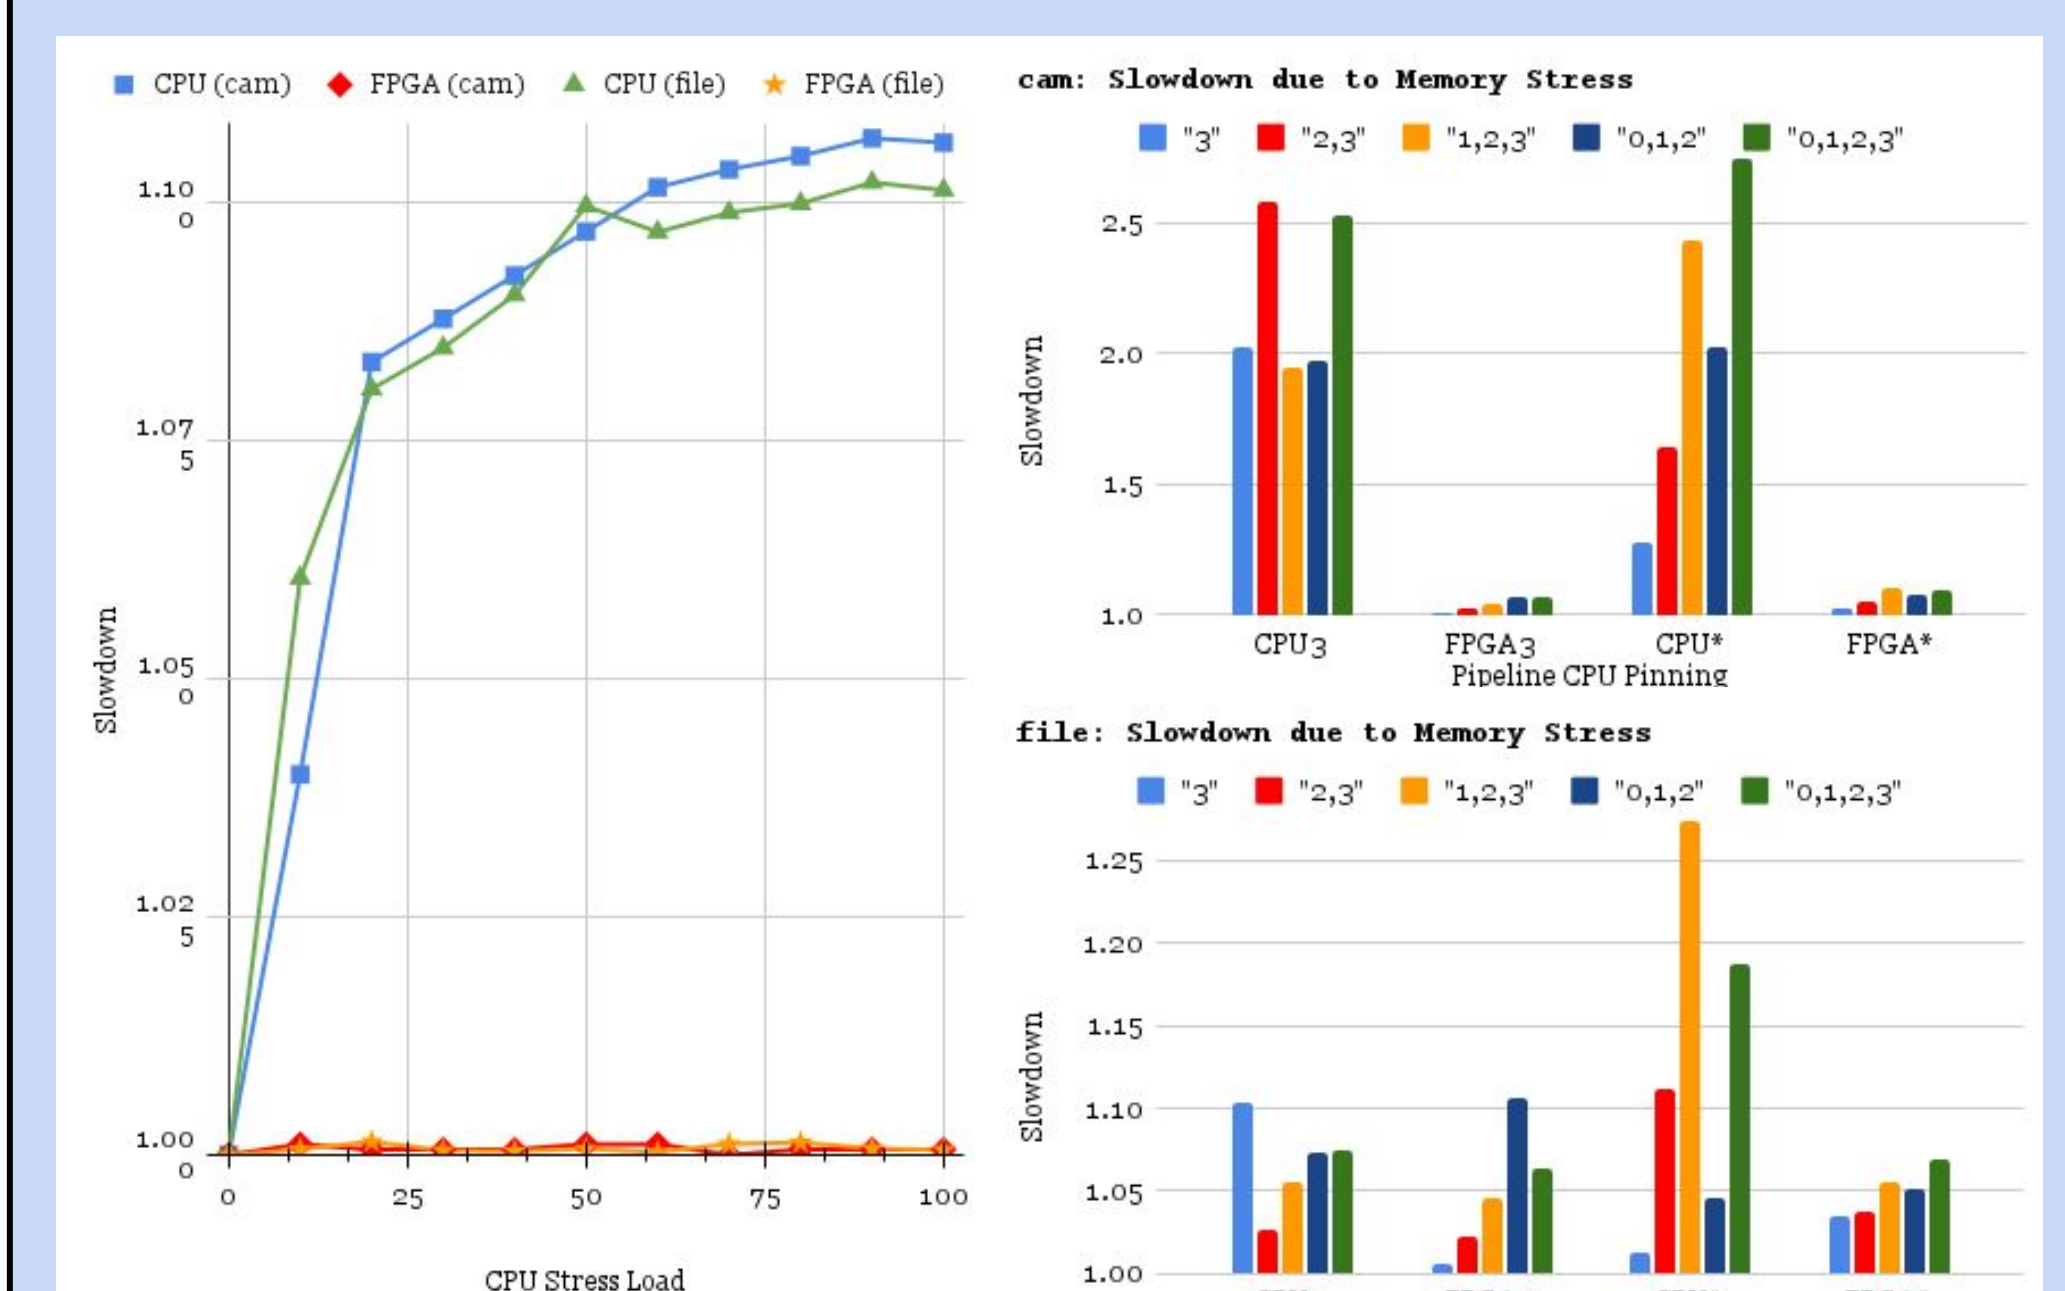
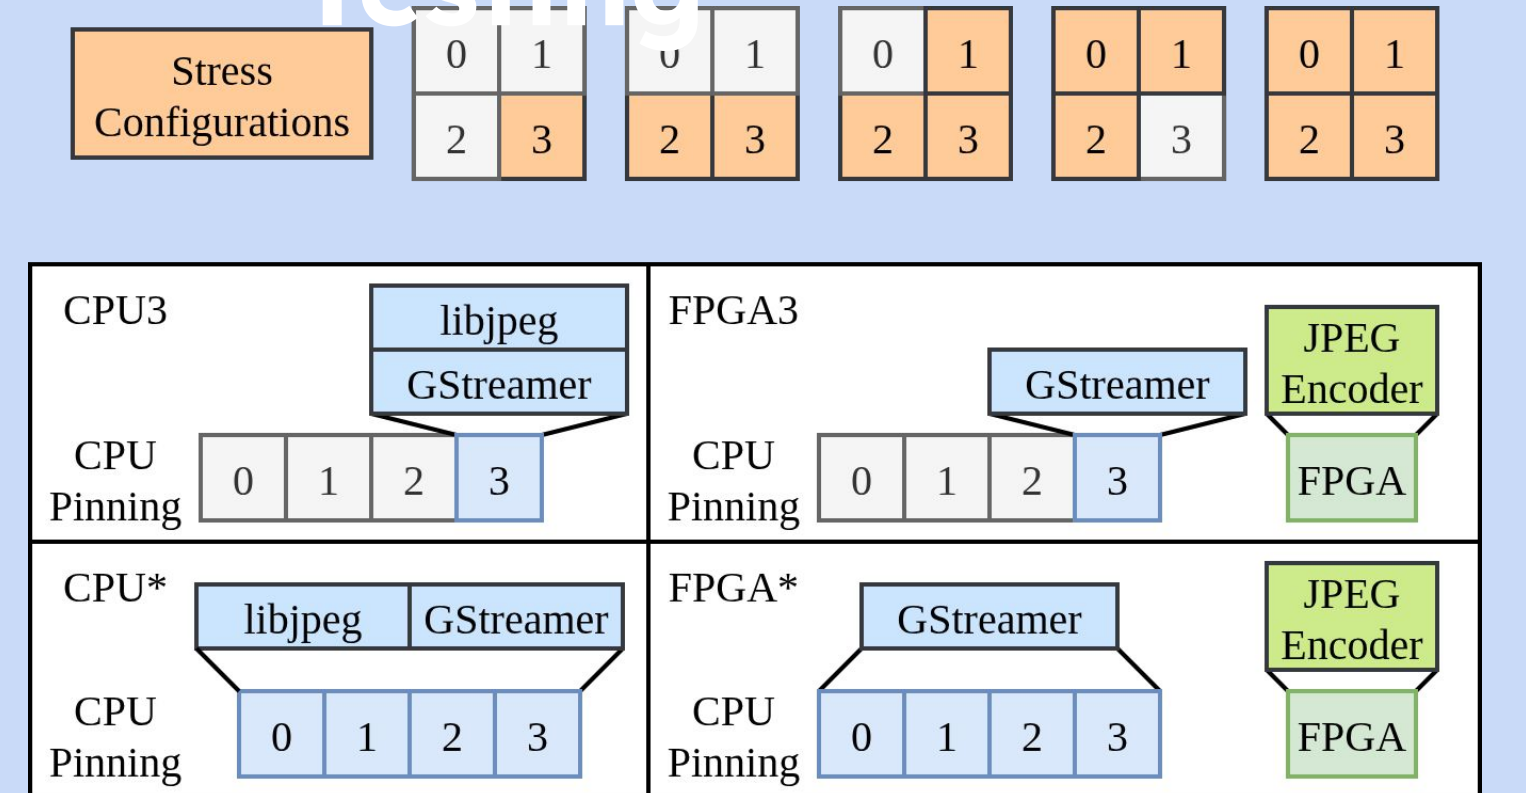
Top-level Block Design



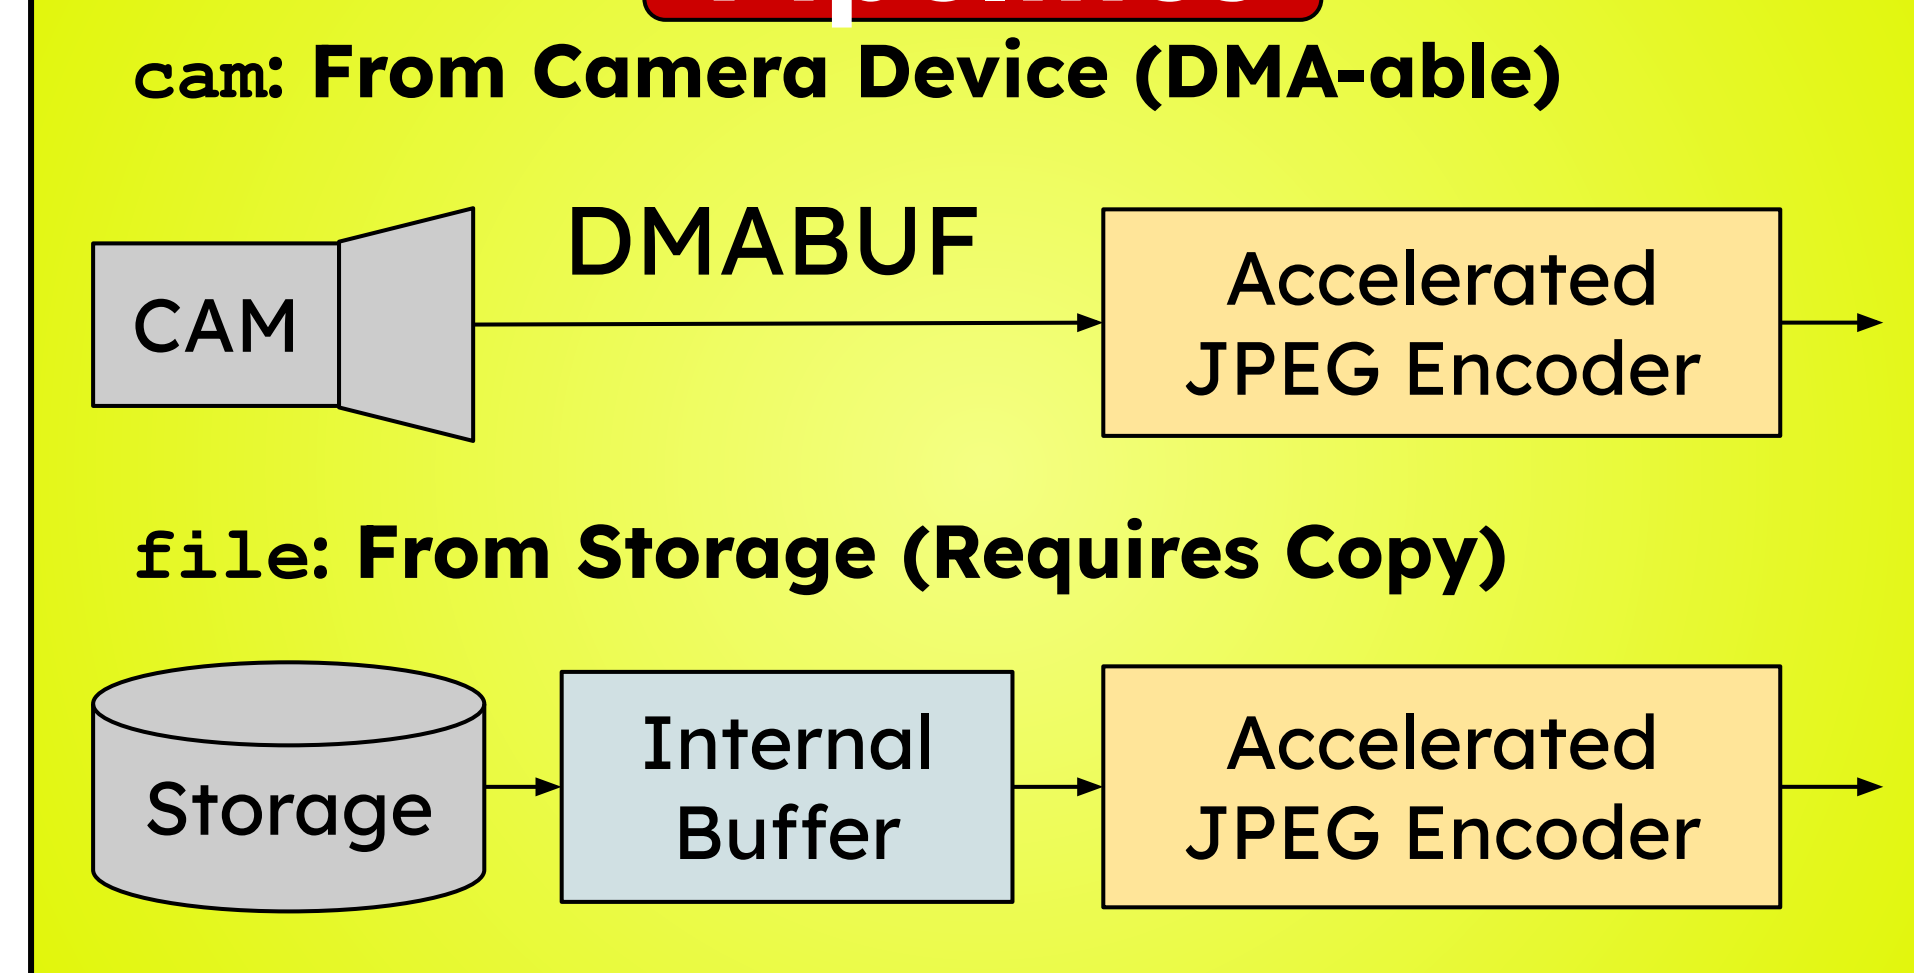
Experiments

CPU/Memory Stress Testing

Objective: Test the slowdown behavior of both the **jpegenc-based CPU encoder** and our **FPGA encoder** when the system is under load.



Pipelines



Performance Evaluation FPGA vs. CPU (jpegenc)

		CPU	FPGA
cam	FPS	2.6	19.5
	CPU %	93	5
file	FPS	22.6	21.7
	CPU %	107	14

Resource Utilization

Generally, *Vitis Libraries* [4] codec kernels **exceed resource limitations** for edge boards, such as the KV260. Thus, employing these kernels on such boards often requires **lowering the target frequency**. Our encoder is able to achieve **7~25%** of the resource usage of even the smallest Vitis encoding kernel (Webp at 100MHz), as shown below. This is done possible by reducing the **HW functionality** as much as possible, and thus lowering complexity.

	Freq.	BRAM	URAM	DSP	FF	LUT
JPEG ¹	250MHz	6	12	85	18150	13868
Webp ²	100MHz	81	46	834	71227	68906
Webp ²	250MHz	229	10	414	92030	68755
JXL ²	260MHz	584	122	644	270565	199537
PIK ²	200MHz	614	473	2398	549987	449995

¹Our work

²Vitis Codec Libraries [4]

Discussion and Conclusions

- Suitability for Edge
-
- Increased Predictability

References

- [1] AMD. "Kria KV260 Vision AI Starter Kit". <https://www.amd.com/en/products/system-on-modules/kria/k260/kv260-vision-starter-kit.html>. Accessed 2025-08-20.
- [2] AMD. "Vitis HLS". <https://www.amd.com/en/products/software/adaptive-socs-and-fpgas/vitis/vitis-hls.html>. Accessed 2025-05-30.
- [4] AMD. "Vitis Libraries". https://github.com/Xilinx/Vitis_Libraries. Accessed 2025-05-30.

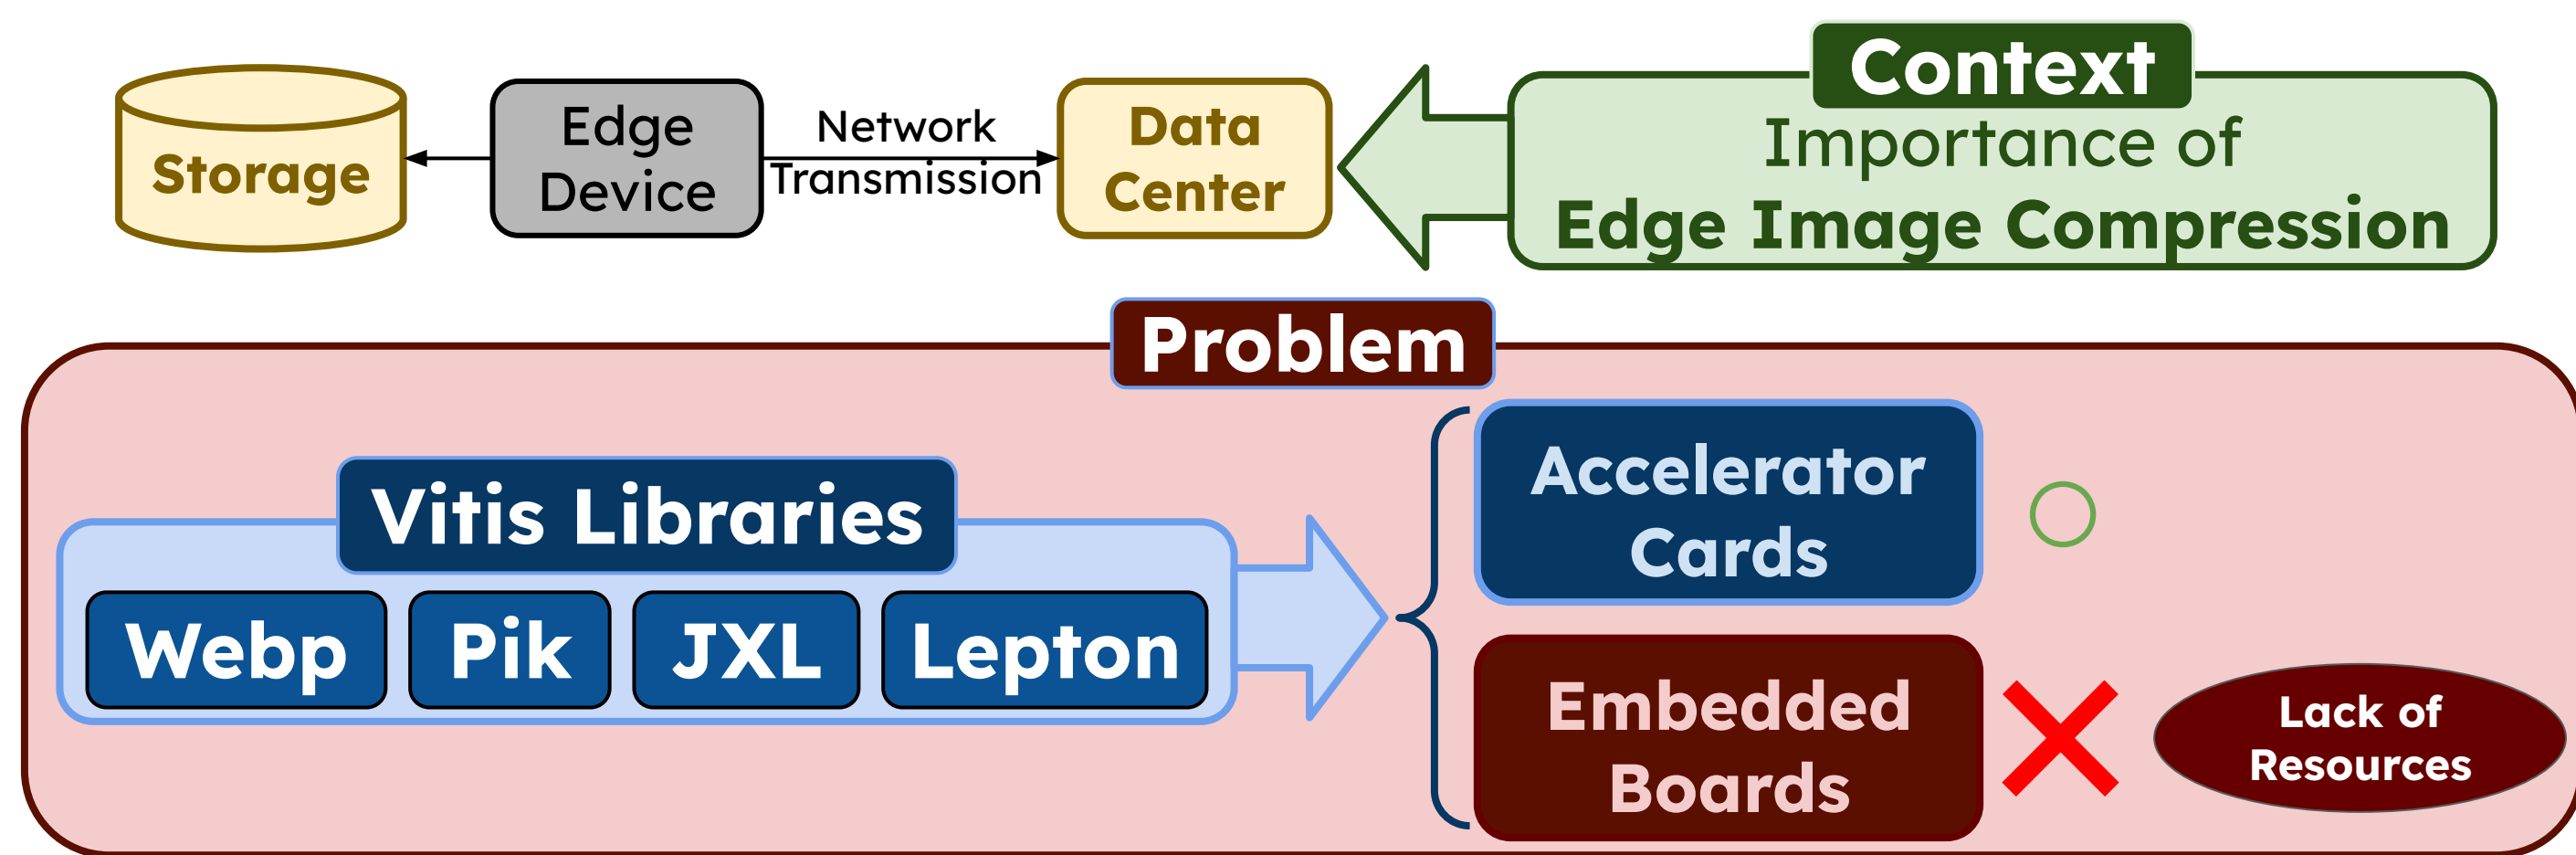
GStreamer-integrated HLS-based JPEG Encoder for Edge FPGA SoCs

Yuri Guimaraes Pereira Primo da Silva¹, Shinya Honda¹,
Sugako Otani¹², Masato Eda¹, Abraham Monrroy Cano³

¹Nagoya University, ²Renesas Electronics,
³Map IV Inc.

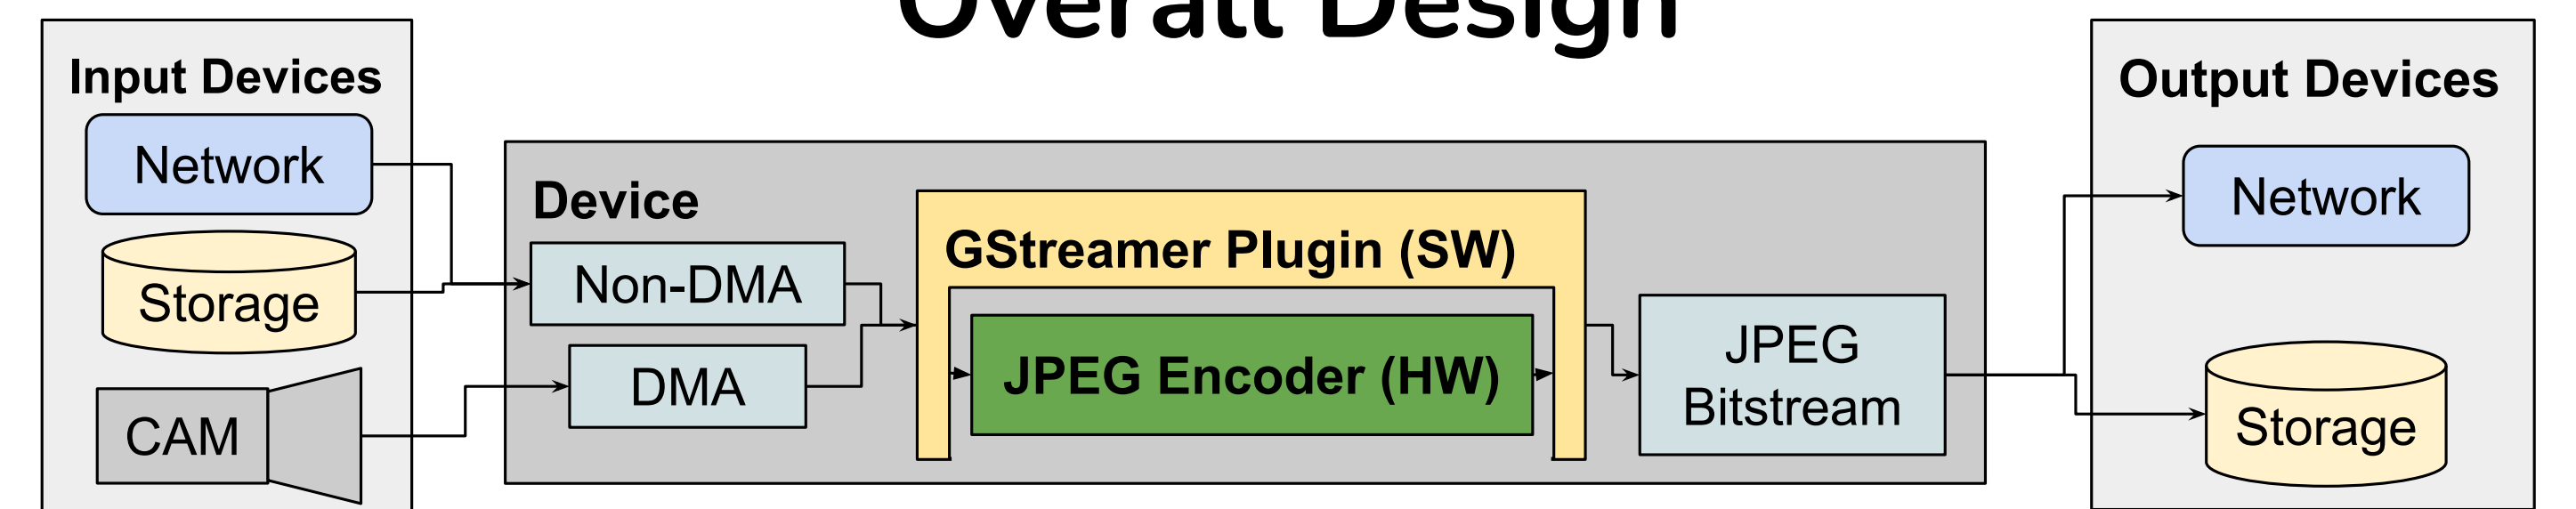
E-mail: yuri_gpps1@ert1.jp

Motivation



Approach

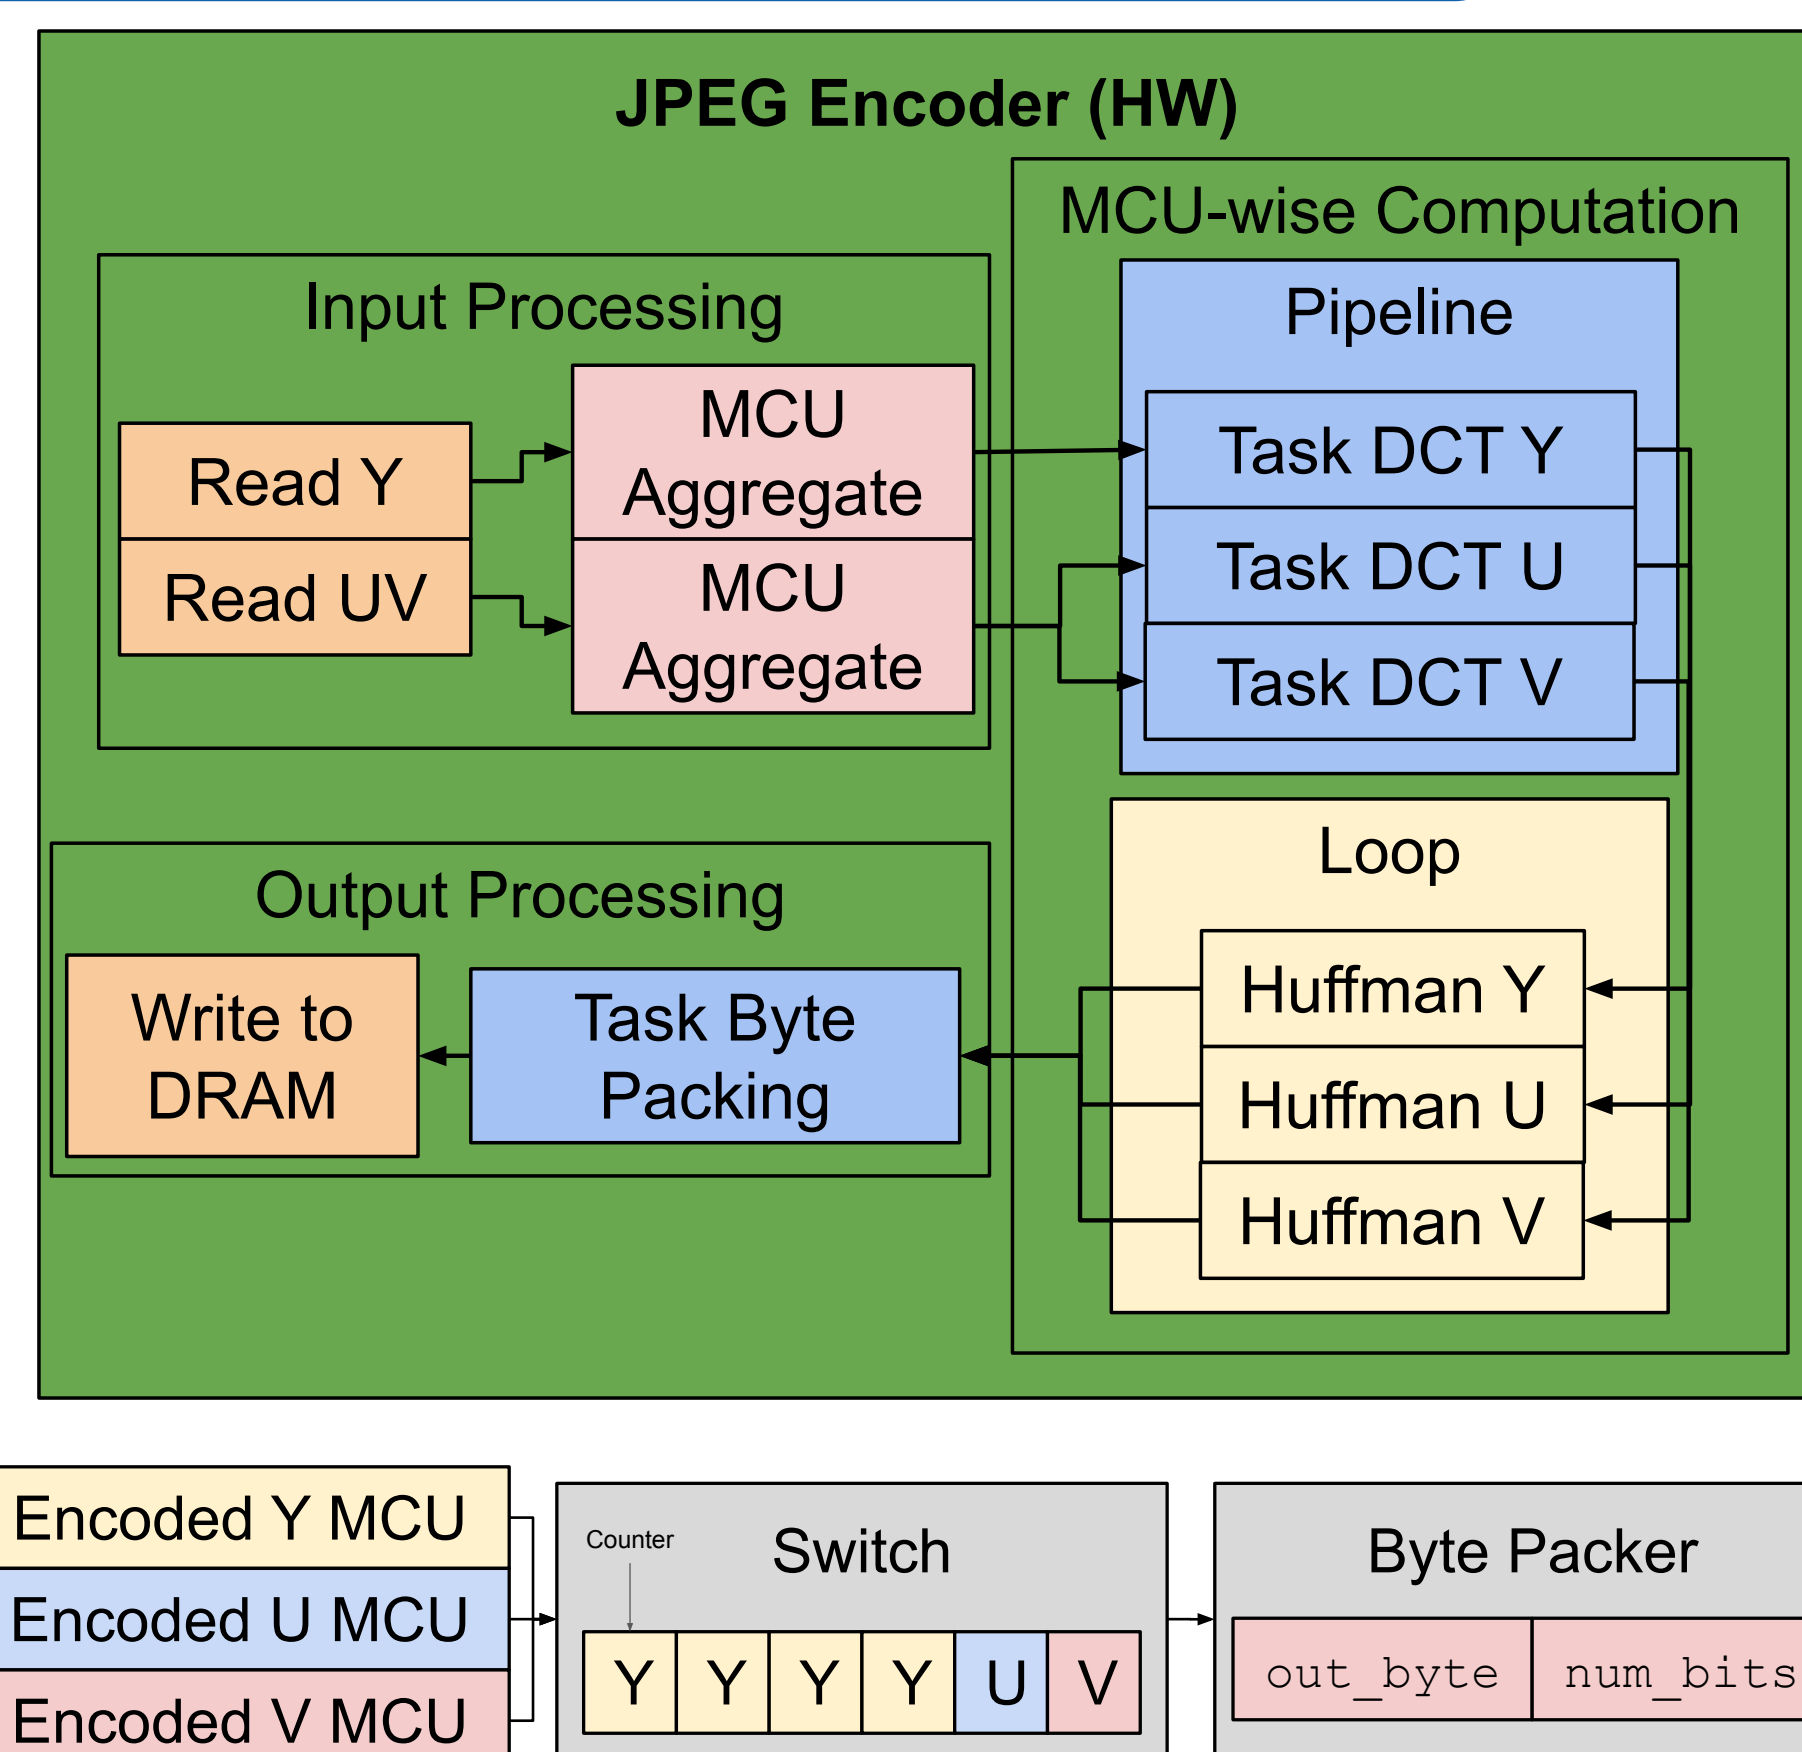
Overall Design



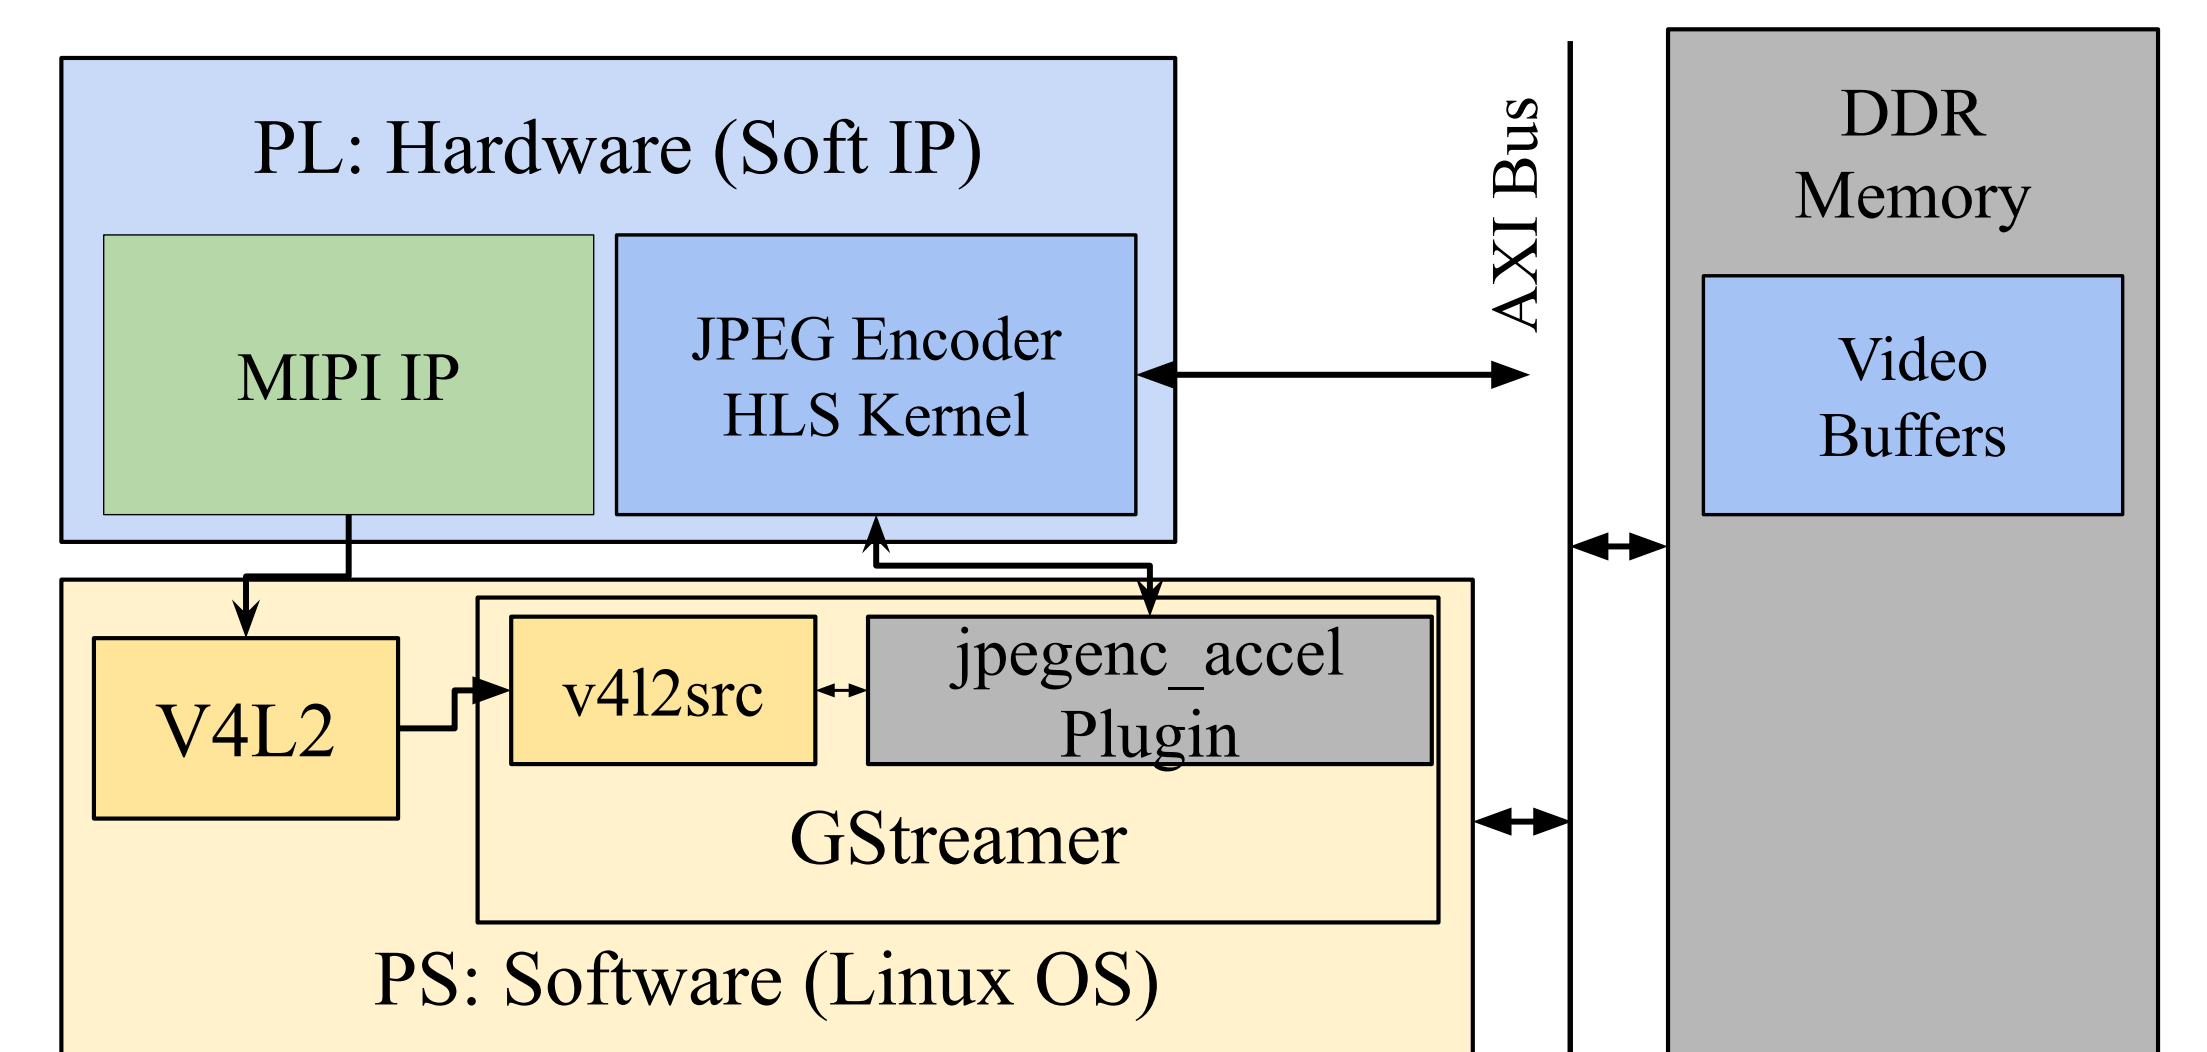
- HW JPEG Encoder implemented in *High-level Synthesis* [2].
- GStreamer plugin interfaces with encoder.
 - Allows for mix and matching of pipeline elements for arbitrary media processing.
 - Easy usage of DMA capabilities.

Hardware Architecture

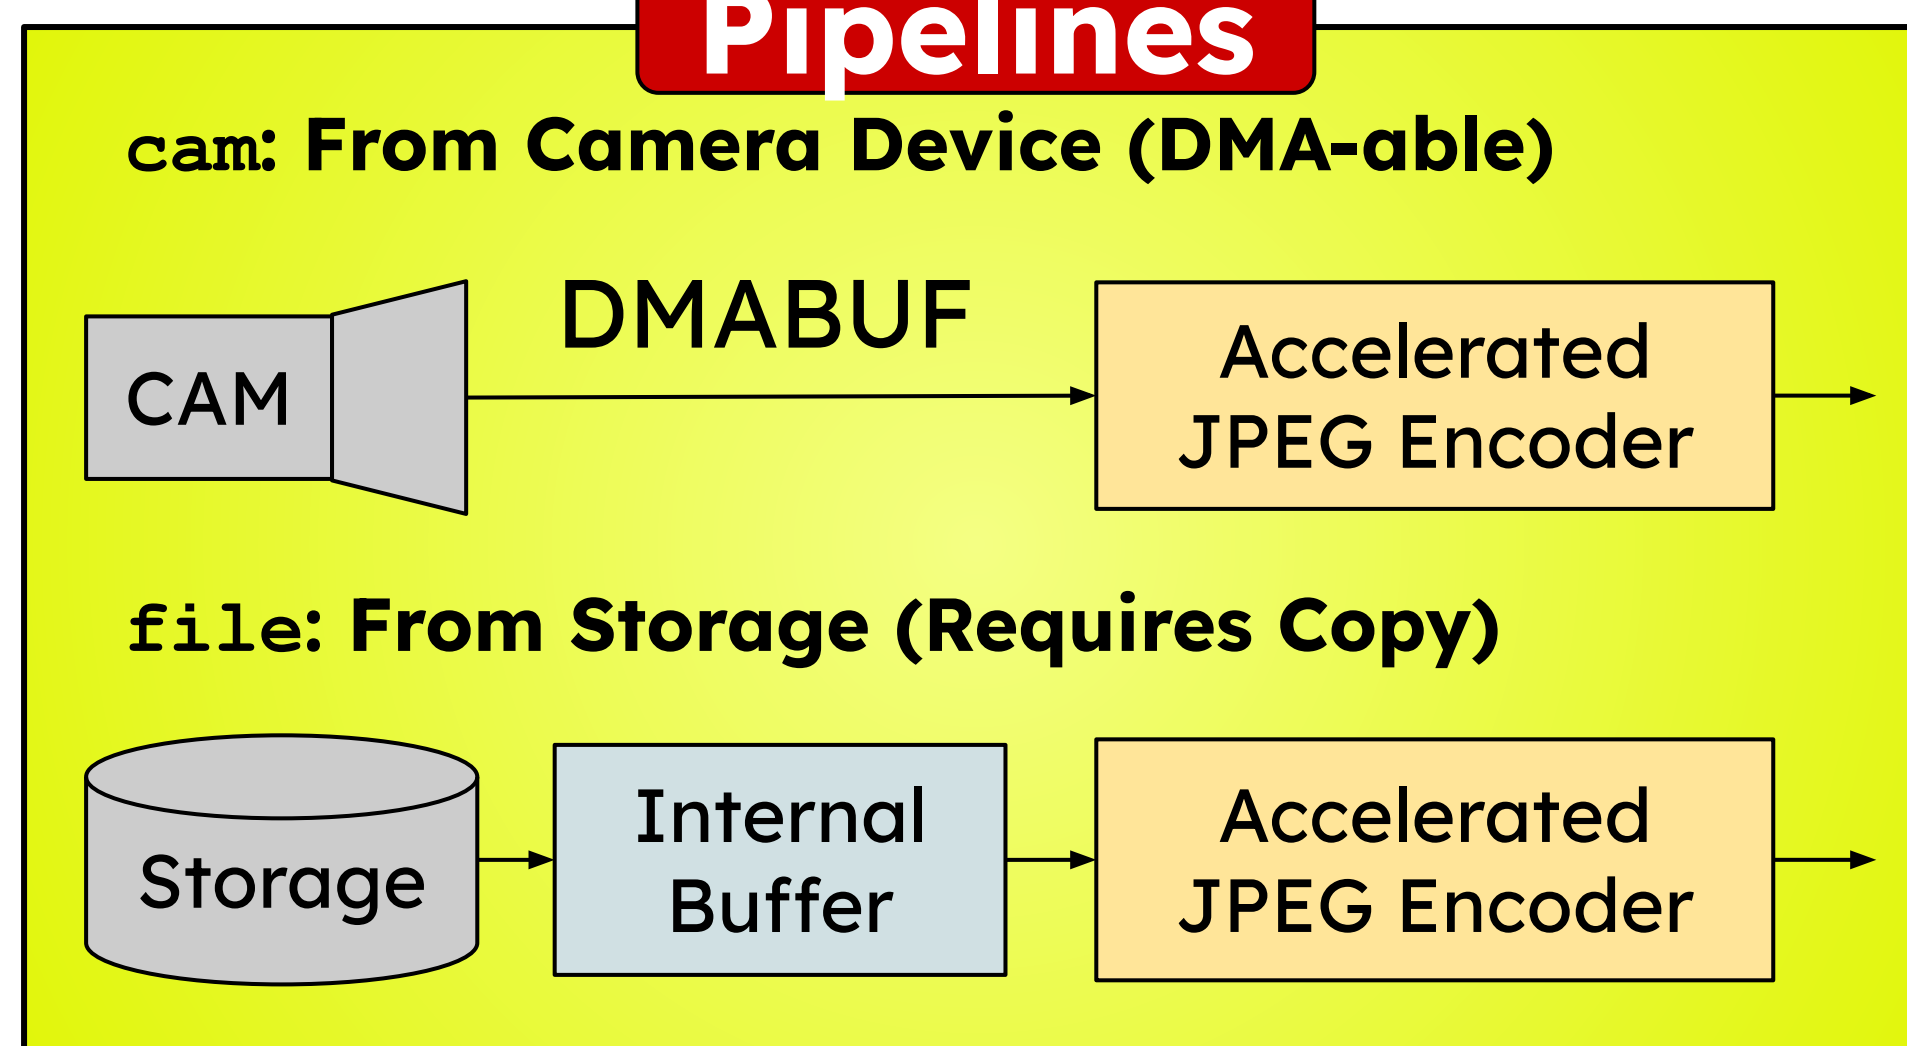
- Conventional JPEG algorithm design implemented C++ HLS.
- Streaming design:
 - Data is packed into MCU streams at input processing.
 - Each MCU in a stream goes through pipelined DCT and RLE pipelined, and sequential Huffman.
 - MCUs are packed into the expected JPEG bitstream format at the end.
- Encoded bits are packed via a stateful switch that shifts and ORs according to the encoding order (e.g. YYYYUV for YUV420).



Top-level Block Design



Pipelines



Performance Evaluation FPGA vs. CPU (jpegenc)

		CPU	FPGA
cam	FPS	2.6	19.5
	CPU %	93	5
file	FPS	22.6	21.7
	CPU %	107	14

Resource Utilization

Generally, *Vitis Libraries* [4] codec kernels **exceed resource limitations** for edge boards, such as the **KV260**. Thus, employing these kernels on such boards often requires **lowering the target frequency**. Our encoder is able to achieve **7~25%** of the resource usage of even the smallest Vitis encoding kernel (Webp at 100MHz), as shown below. This is done possible by reducing the **HW functionality** as much as possible, and thus lowering complexity.

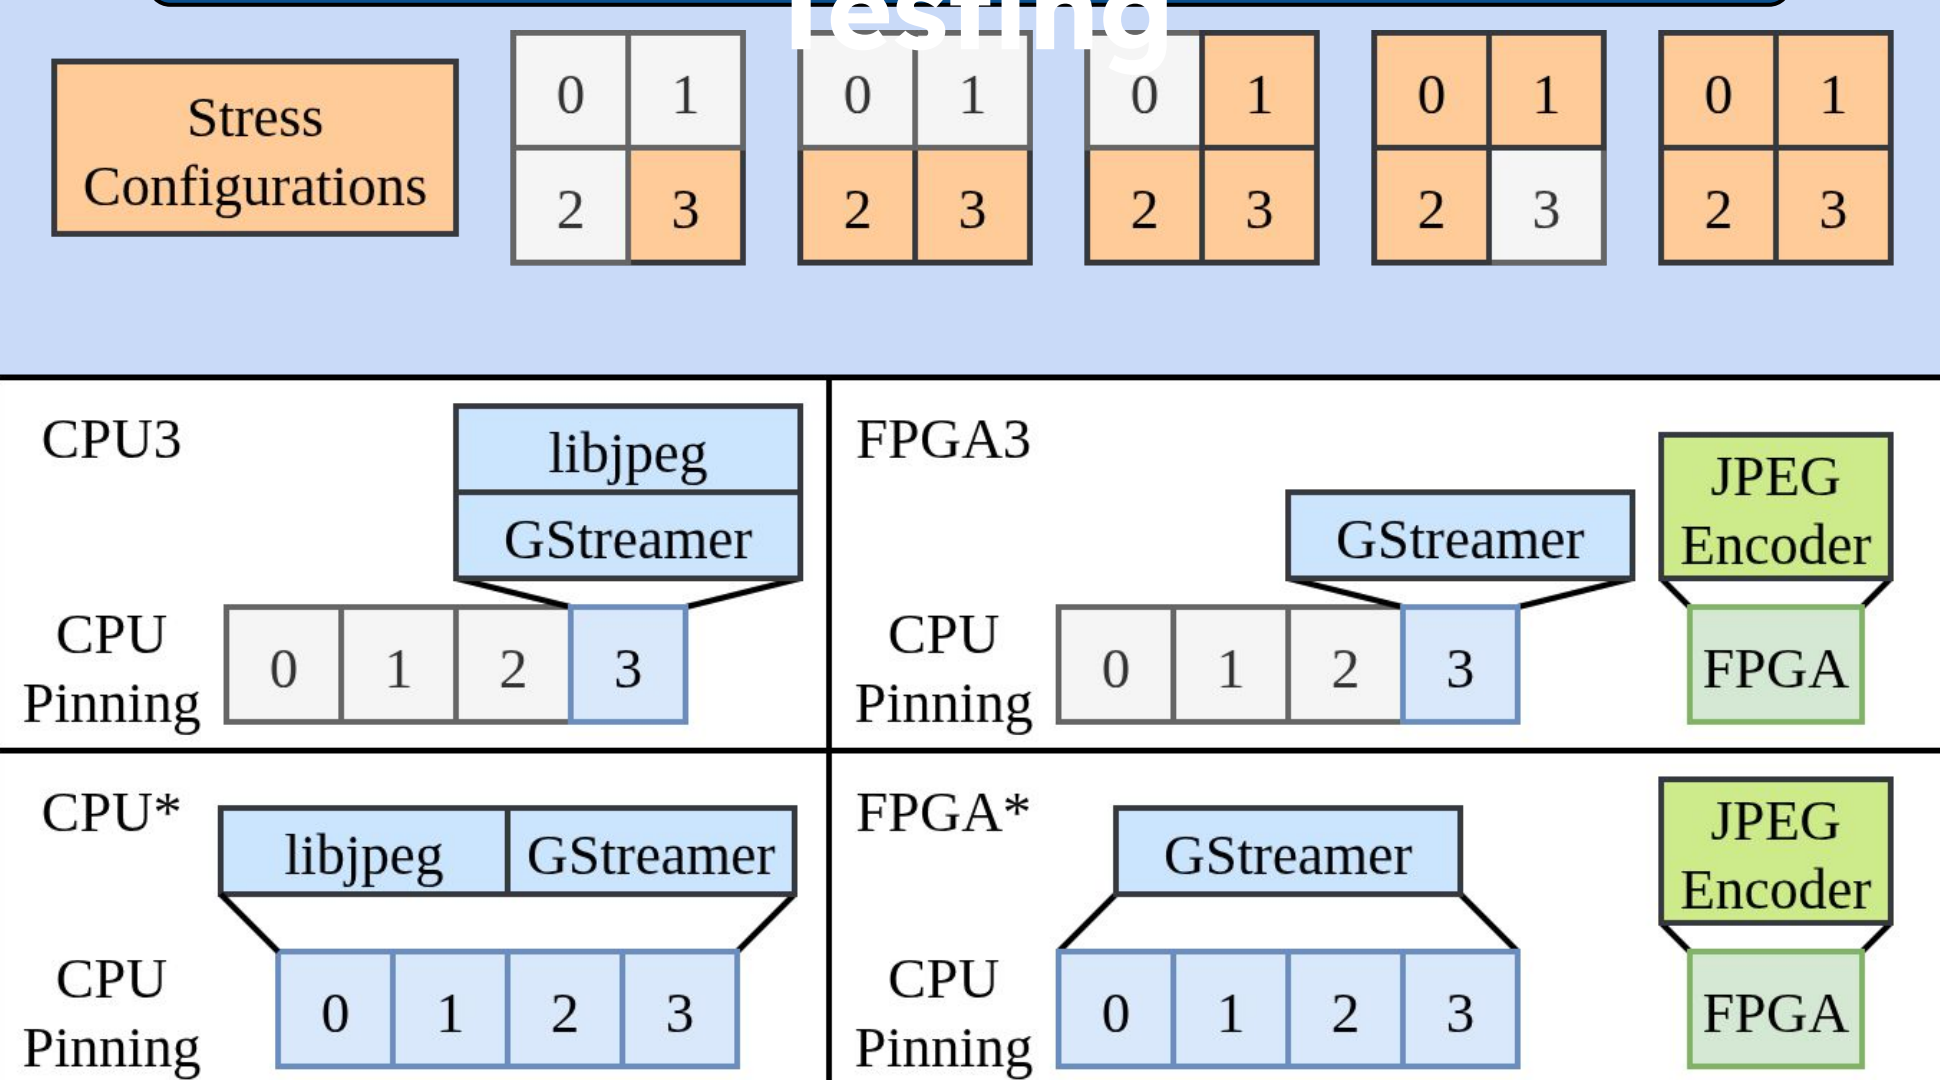
	Freq.	BRAM	URAM	DSP	FF	LUT
JPEG ¹	250MHz	6	12	85	18150	13868
Webp ²	100MHz	81	46	834	71227	68906
Webp ²	250MHz	229	10	414	92030	68755
JXL ²	260MHz	584	122	644	270565	199537
PIK ²	200MHz	614	473	2398	549987	449995

¹Our work

²Vitis Codec Libraries [4]

Experiments

CPU/Memory Stress Testing



Objective: Test the **slowdown** behavior of both the **jpegecc-based CPU encoder** and our **FPGA encoder** when the system is under load.

Discussion and Conclusions

1. Suitability for Edge
- 2.
3. Increased Predictability

References

- [1] AMD. "Kria KV260 Vision AI Starter Kit". <https://www.amd.com/en/products/system-on-modules/kria/kv260-vision-starter-kit.html>. Accessed 2025-08-20.
- [2] AMD. "Vitis HLS". <https://www.amd.com/en/products/software/adaptive-socs-and-fpgas/vitis/vitis-hls.html>. Accessed 2025-05-30.
- [4] AMD. "Vitis Libraries". https://github.com/Xilinx/Vitis_Libraries. Accessed 2025-05-30.

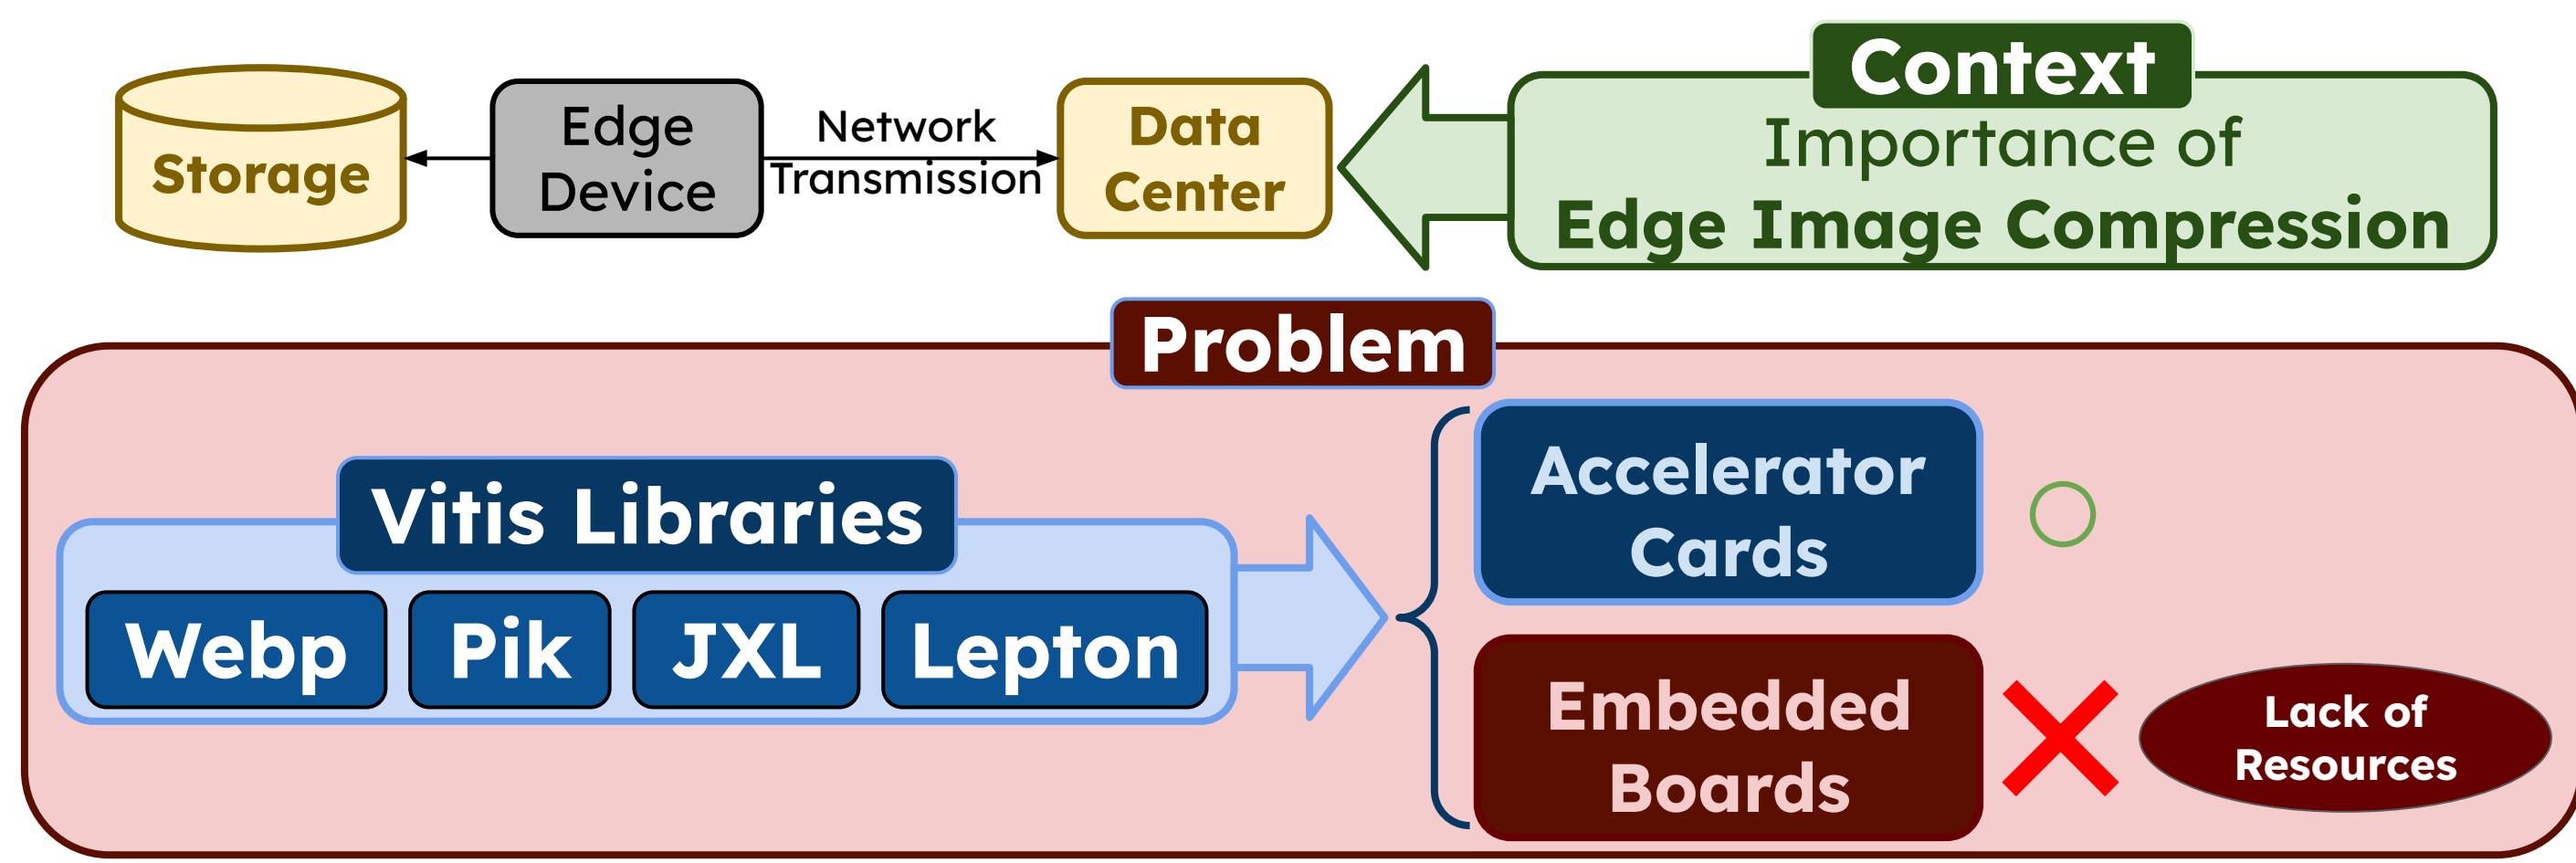
GStreamer-integrated HLS-based JPEG Encoder for Edge FPGA SoCs

Yuri Guimaraes Pereira Primo da Silva¹, Shinya Honda¹,
Sugako Otani¹², Masato Eda¹, Abraham Monrroy Cano³

¹Nagoya University, ²Renesas Electronics,
³Map IV Inc.

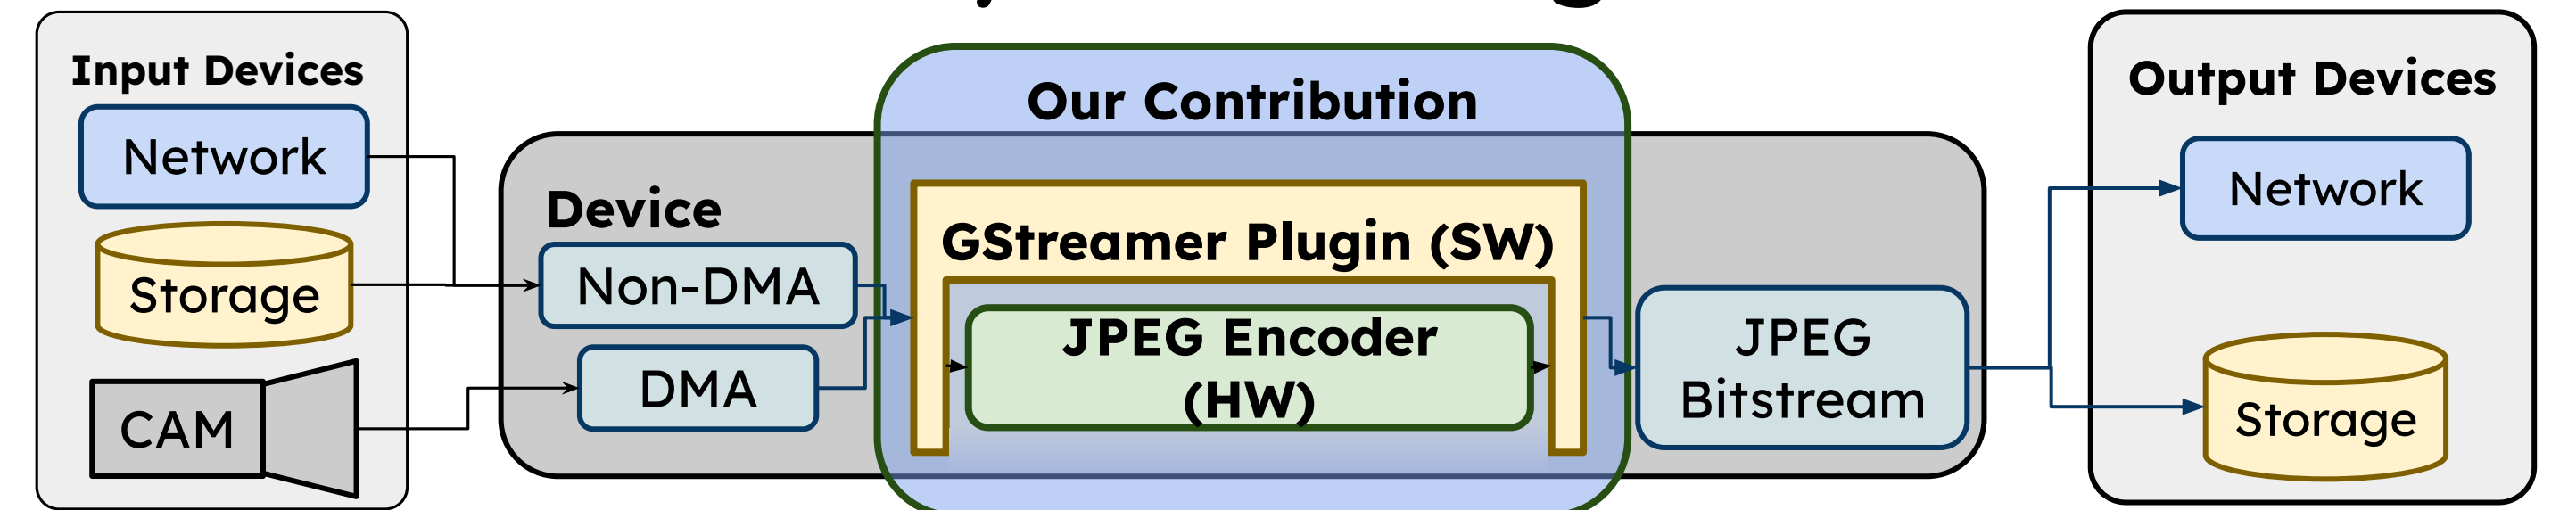
E-mail: yuri_gpps1@ert1.jp

Motivation



Approach

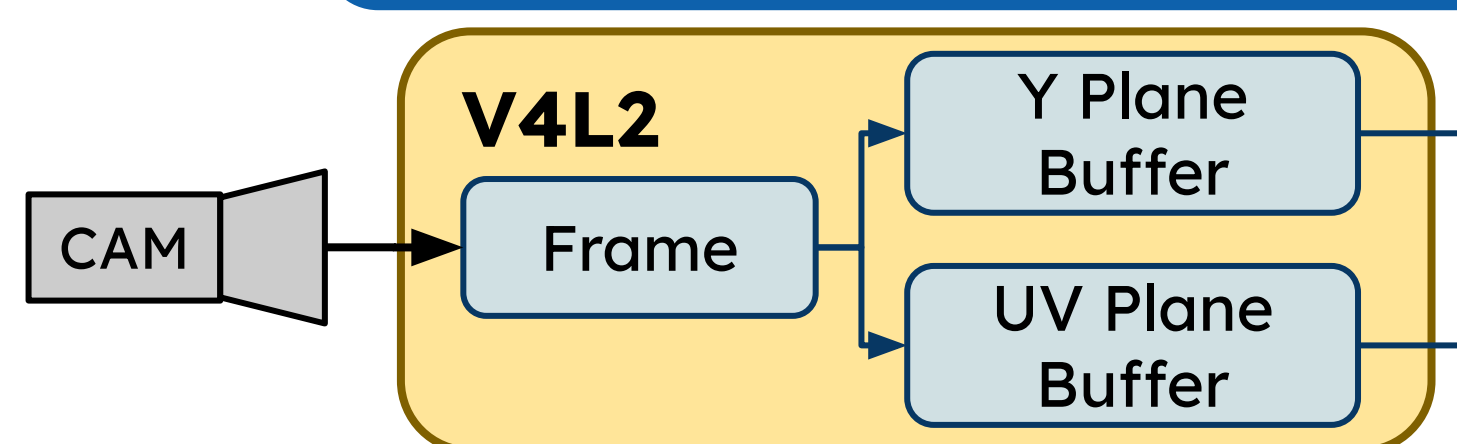
System Design



Contributions

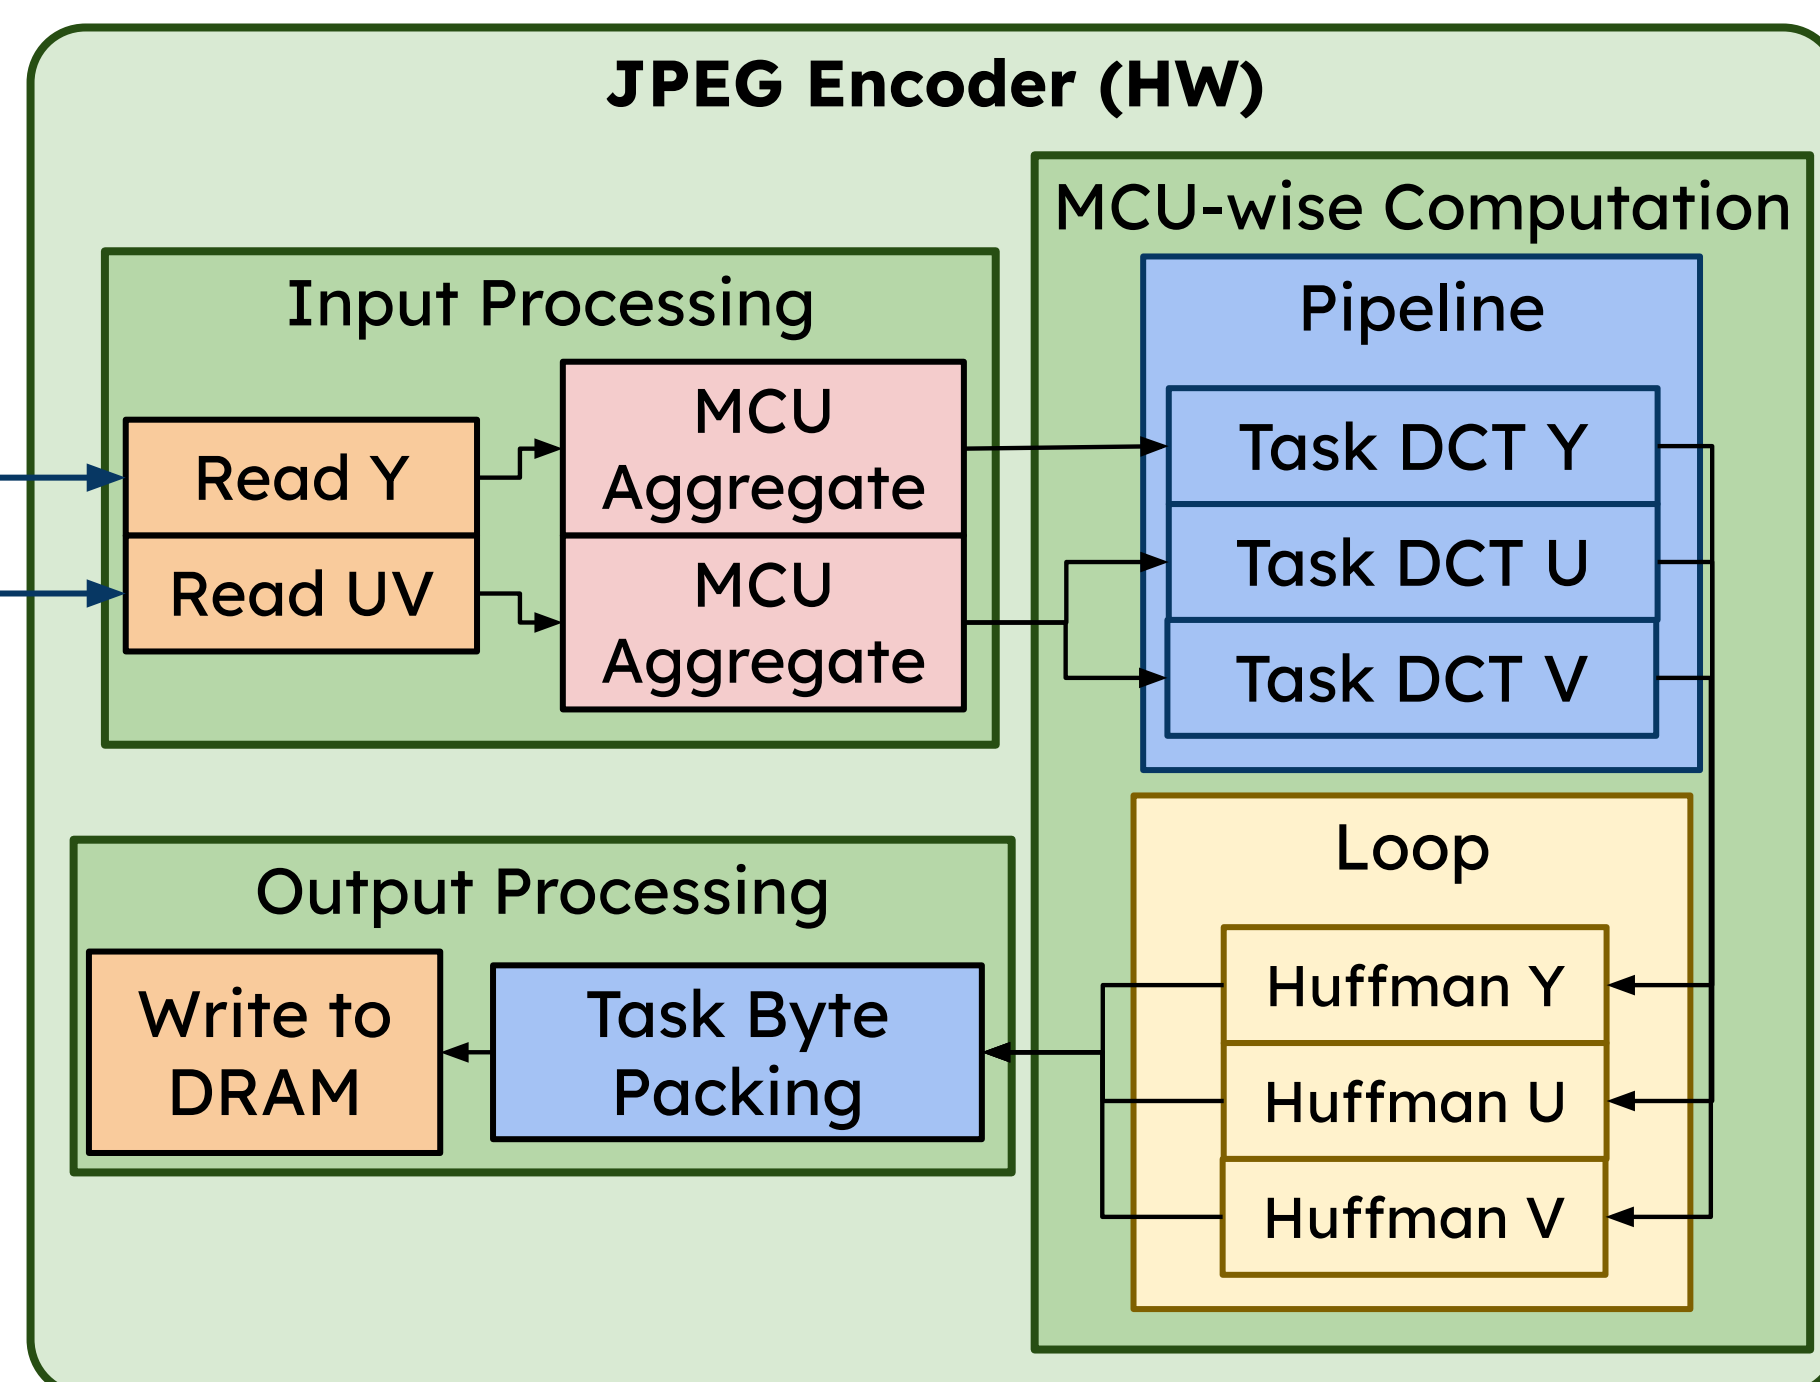
- **HW JPEG Encoder** implemented in *High-level Synthesis*.
- **GStreamer plugin** interfaces with encoder.
 - Allows for mix and matching of pipeline elements for arbitrary media processing.
 - Easy usage of DMA capabilities.

Hardware Architecture

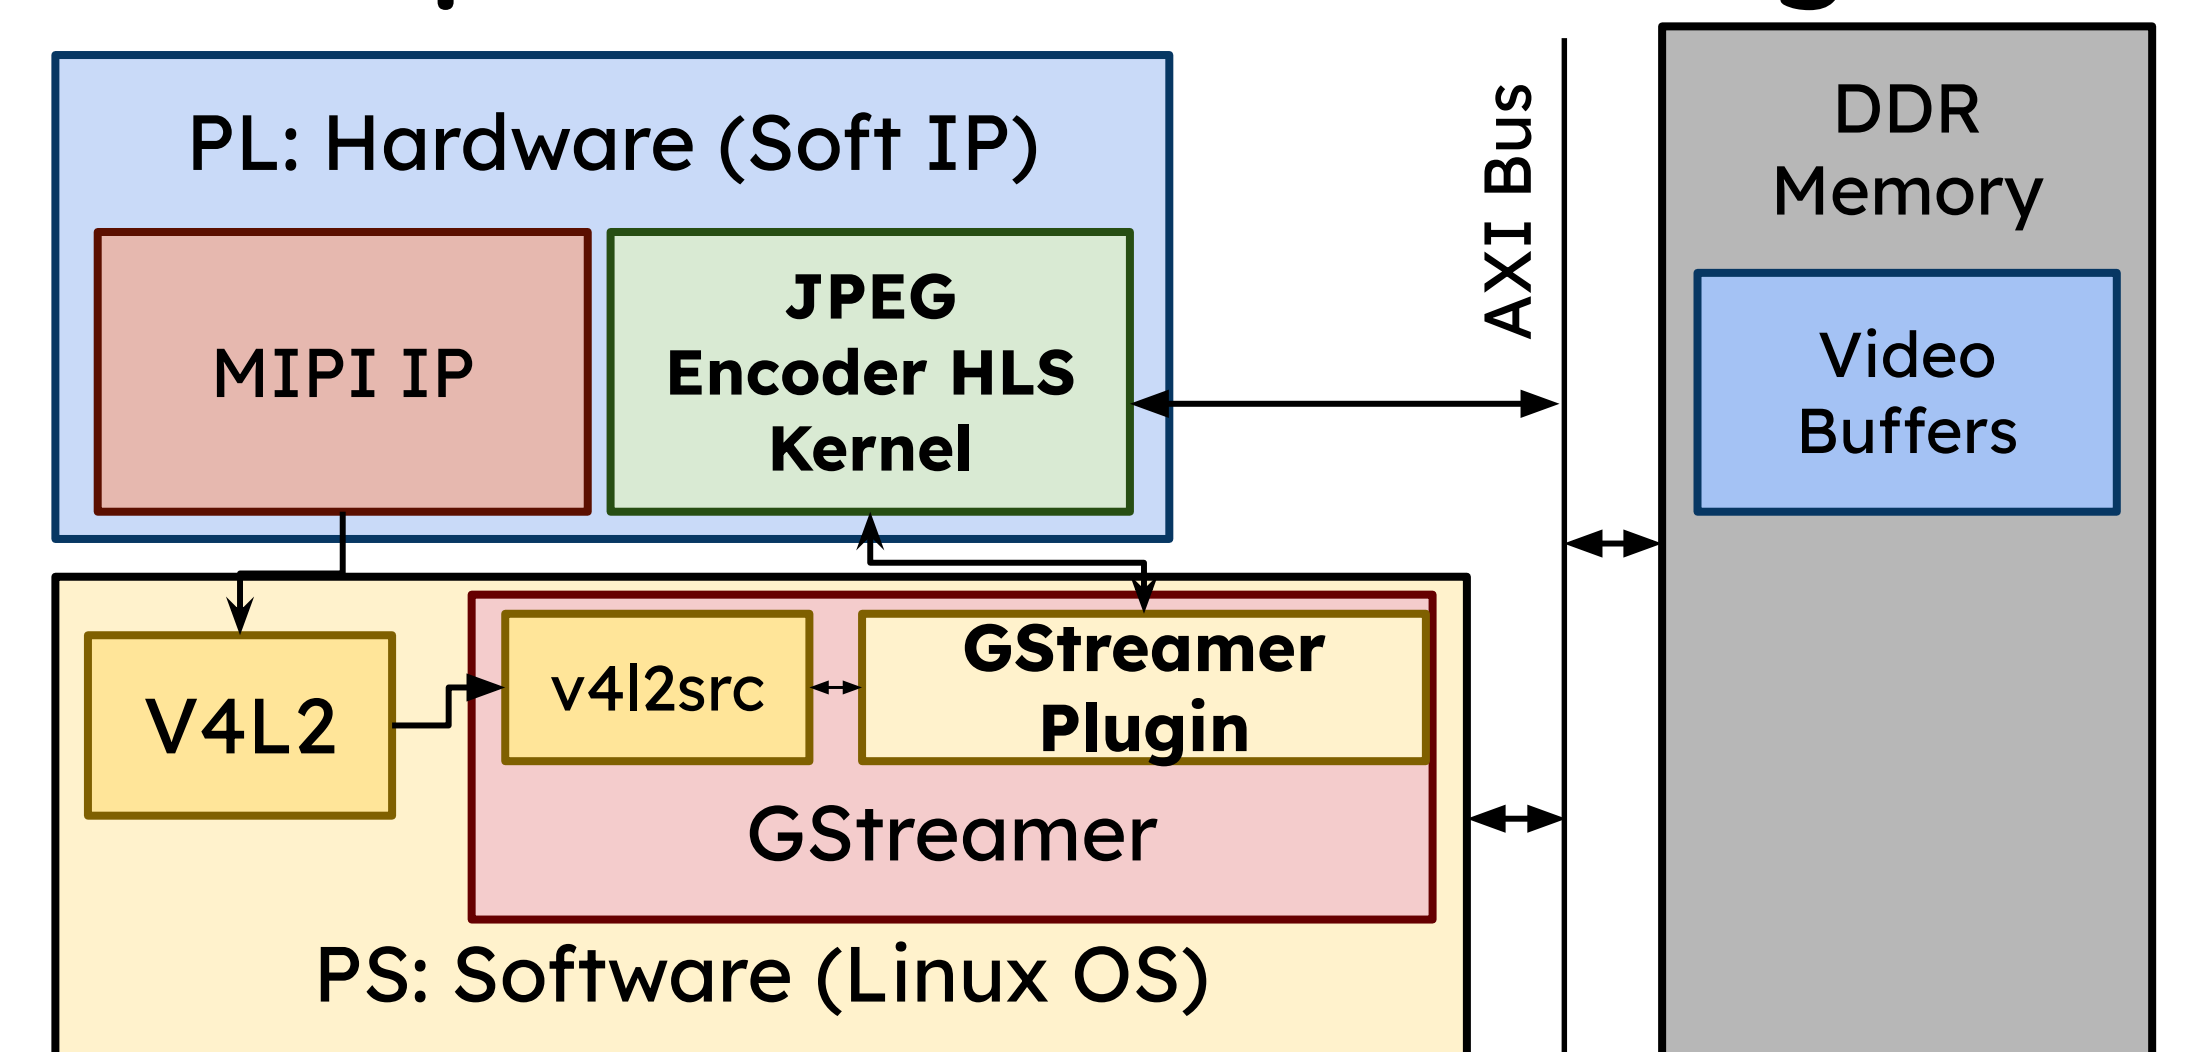


- Conventional JPEG algorithm design implemented C++ HLS.
- Streaming design:
 - Data is packed into MCU streams at input processing.
 - Each MCU in a stream goes through pipelined DCT and RLE pipelined, and sequential Huffman.
 - MCUs are packed into the expected JPEG bitstream format at the end.

- Send each plane to the encoder via AXI
 - Avoid memory copy in case of DMA

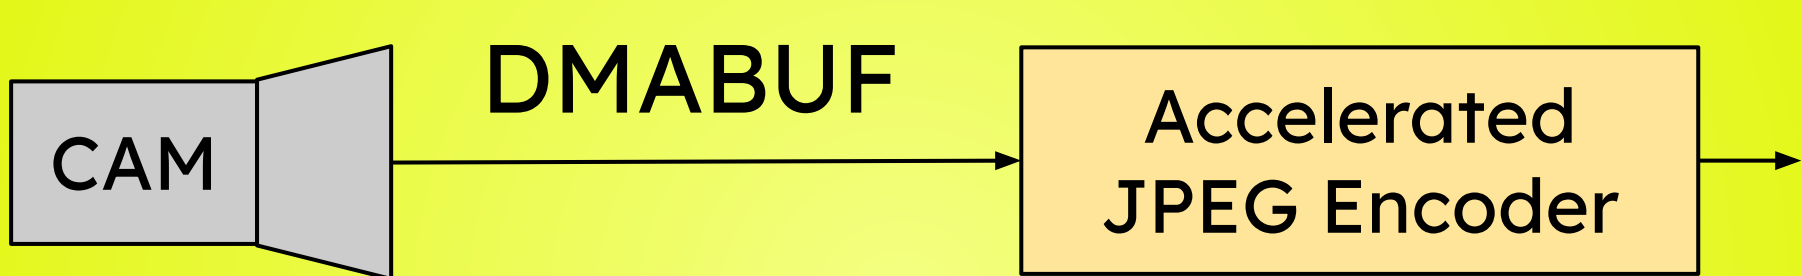


Top-level Hardware Design

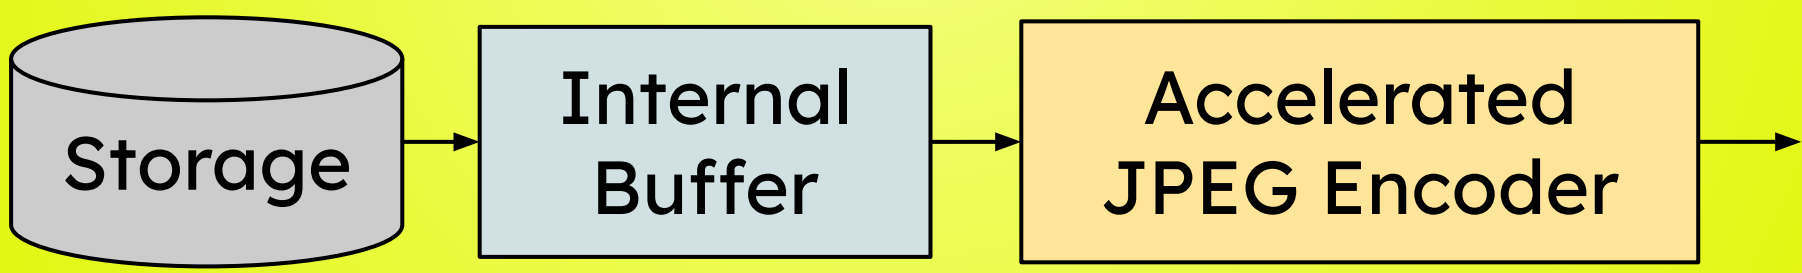


Definition: Pipelines

cam: From Camera Device (DMA-able)



file: From Storage (Requires Copy)



Performance Evaluation

FPGA vs. CPU (jpegenc)

		CPU	FPGA
cam	FPS	2.6	19.5
	CPU %	93	5
file	FPS	22.6	21.7
	CPU %	107	14

Resource Utilization

	Freq.	BRAM	URAM	DSP	FF	LUT
JPEG ¹	250MHz	6	12	85	18150	13868
Webp ²	100MHz	81	46	834	71227	68906
Webp ²	250MHz	229	10	414	92030	68755
JXL ²	260MHz	584	122	644	270565	199537
PIK ²	200MHz	614	473	2398	549987	449995

Webp Encoder²

- **High Resource Usage**
 - Requires lowering **target frequency** to 100Hz

JPEG Encoder¹

- **7~25% of Webp Resources**
 - Can fit **four** JPEG encoders in place of a **single** Webp Encoder

¹Our work

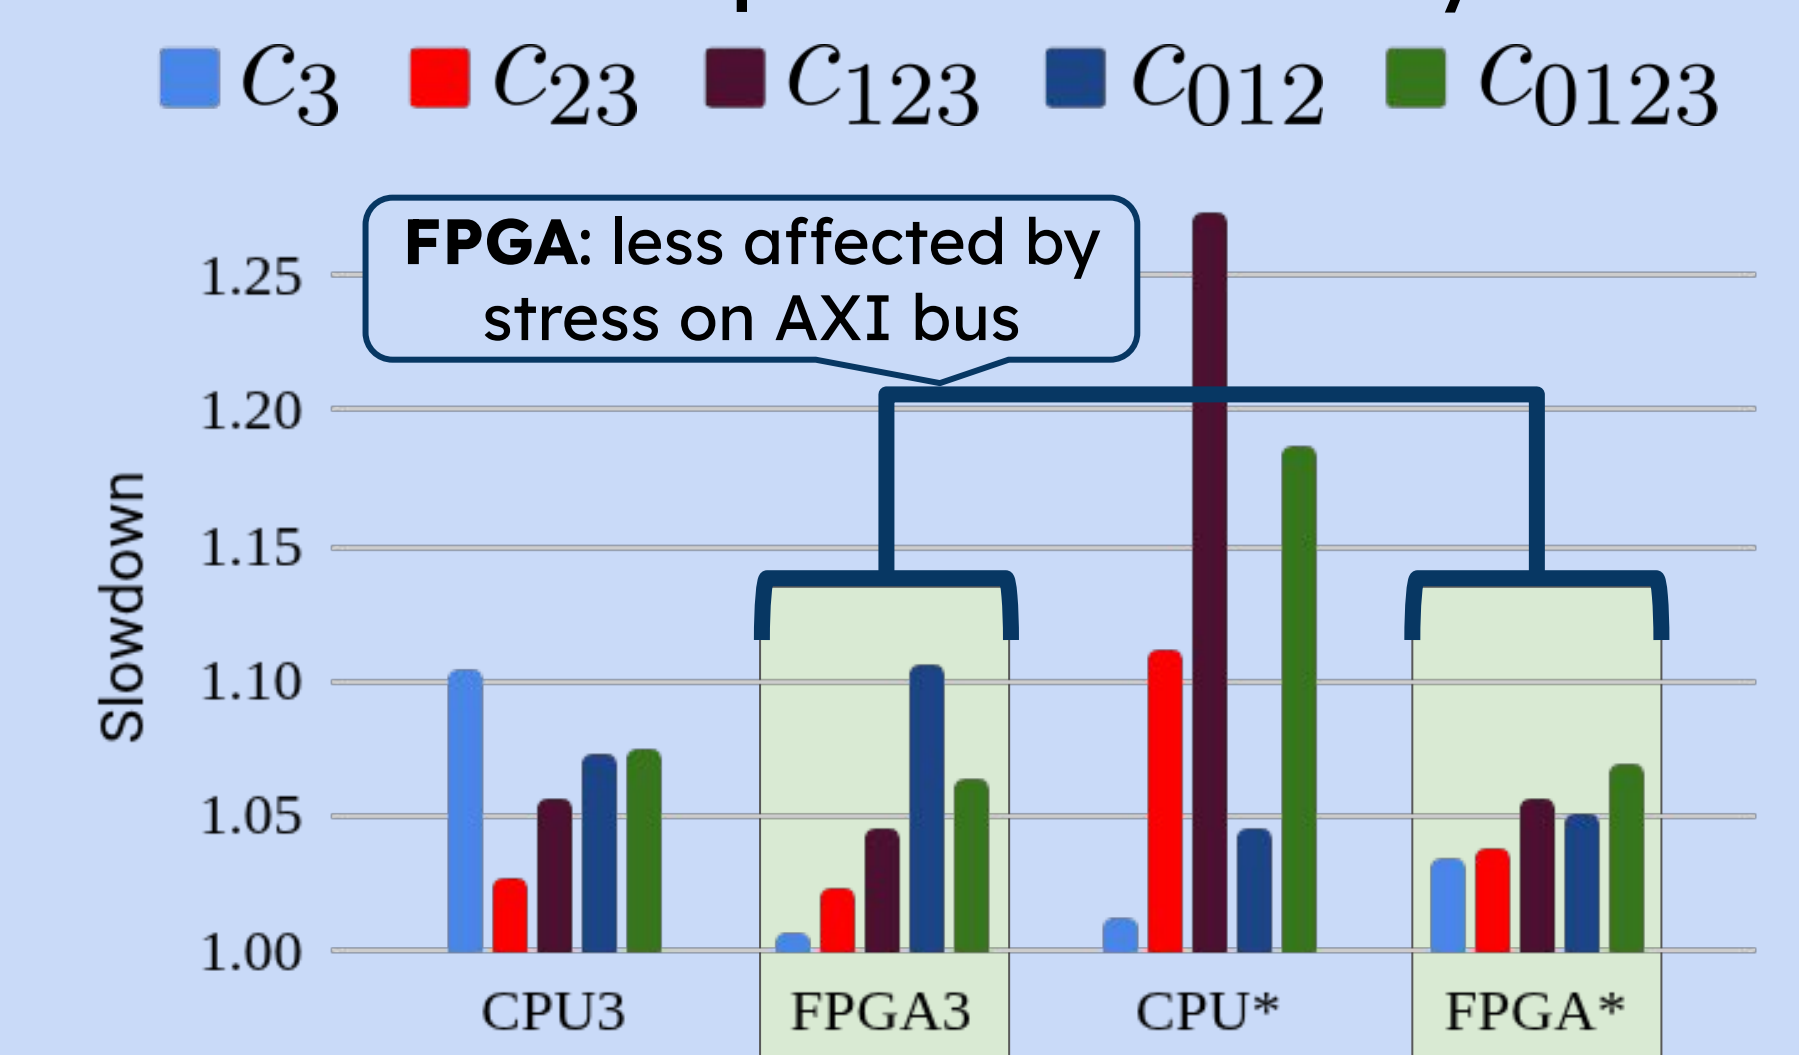
²Vitis Codec Libraries

Experiments

Memory Stress Testing

Stress Configurations	0	1	0	1	0	1	0	1	0	1	
	2	3	2	3	2	3	2	3	2	3	
$C_3 \quad C_{23} \quad C_{123} \quad C_{012} \quad C_{0123}$											
CPU3	libjpeg				GStreamer						
CPU Pinning	0	1	2	3							
FPGA3					GStreamer				JPEG Encoder		
CPU Pinning	0	1	2	3					FPGA		
CPU*	libjpeg		GStreamer		GStreamer				JPEG Encoder		
CPU Pinning	0	1	2	3	0				1	2	3
FPGA*					GStreamer				JPEG Encoder		
CPU Pinning					0				1	2	3

file: **Slowdown Comparison under Memory Stress**



Discussion and Conclusions

1. Suitability for Edge

- Low resource usage
- Realtime performance

2. Usability

- GStreamer integration
 - Easy composition of media pipelines

3. Increased Predictability

- Resilience against memory stress

Future Work

- Streamlined flow for GStreamer integration with arbitrary HLS kernels

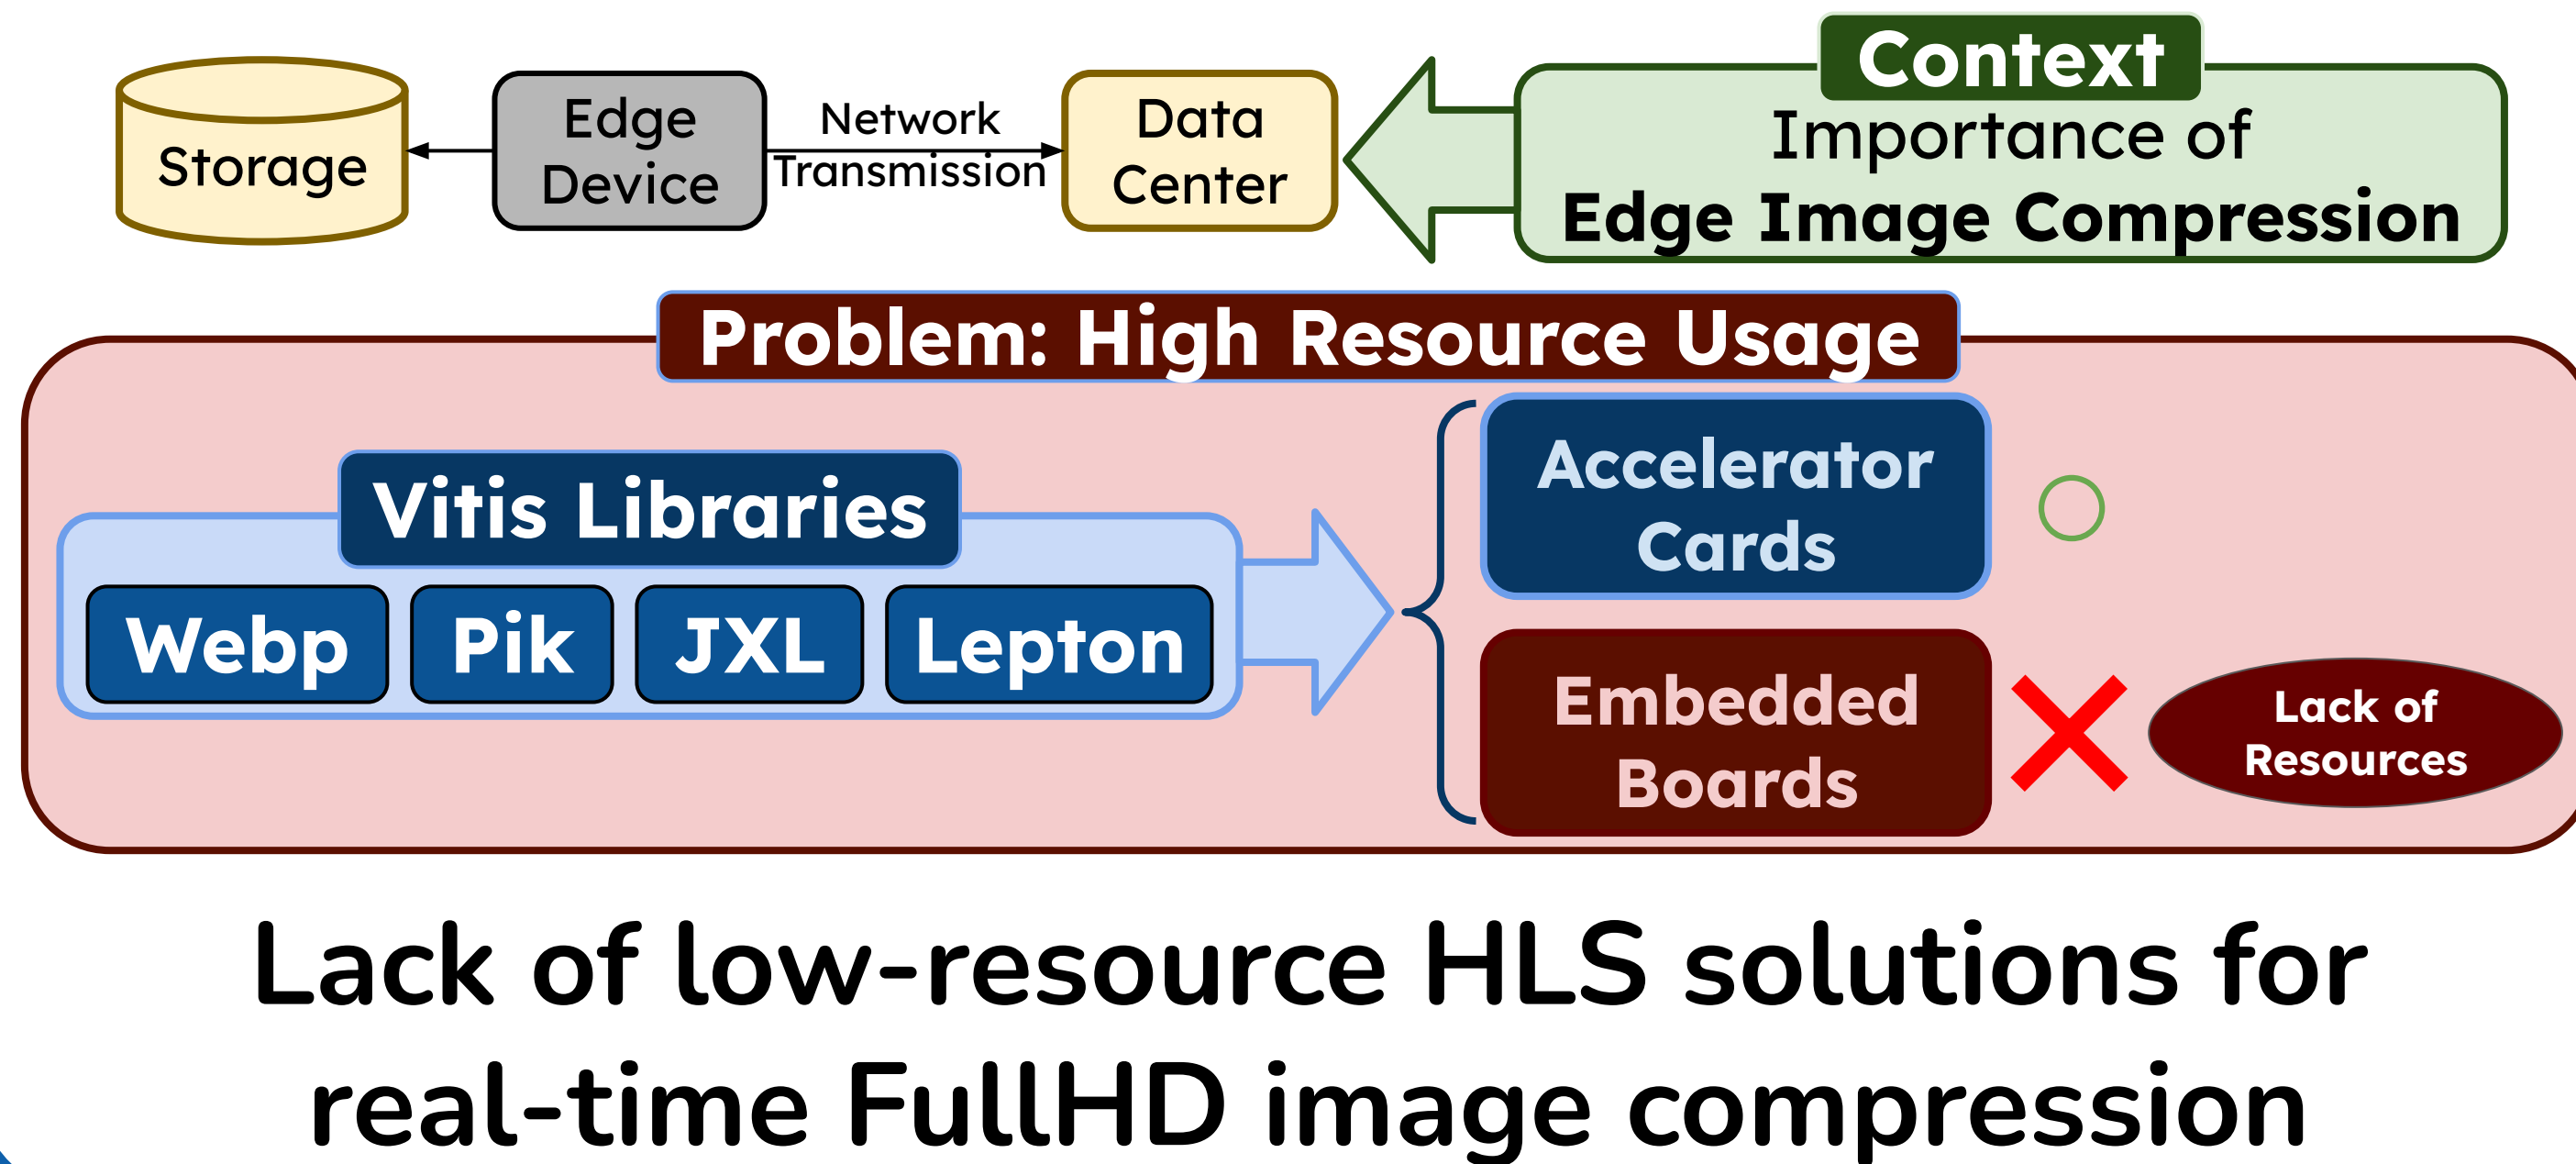
GStreamer-integrated HLS-based JPEG Encoder for Edge FPGA SoCs

Yuri Guimaraes Pereira Primo da Silva¹, Shinya Honda¹,
Sugako Otani¹², Masato Eda¹, Abraham Monrroy Cano³

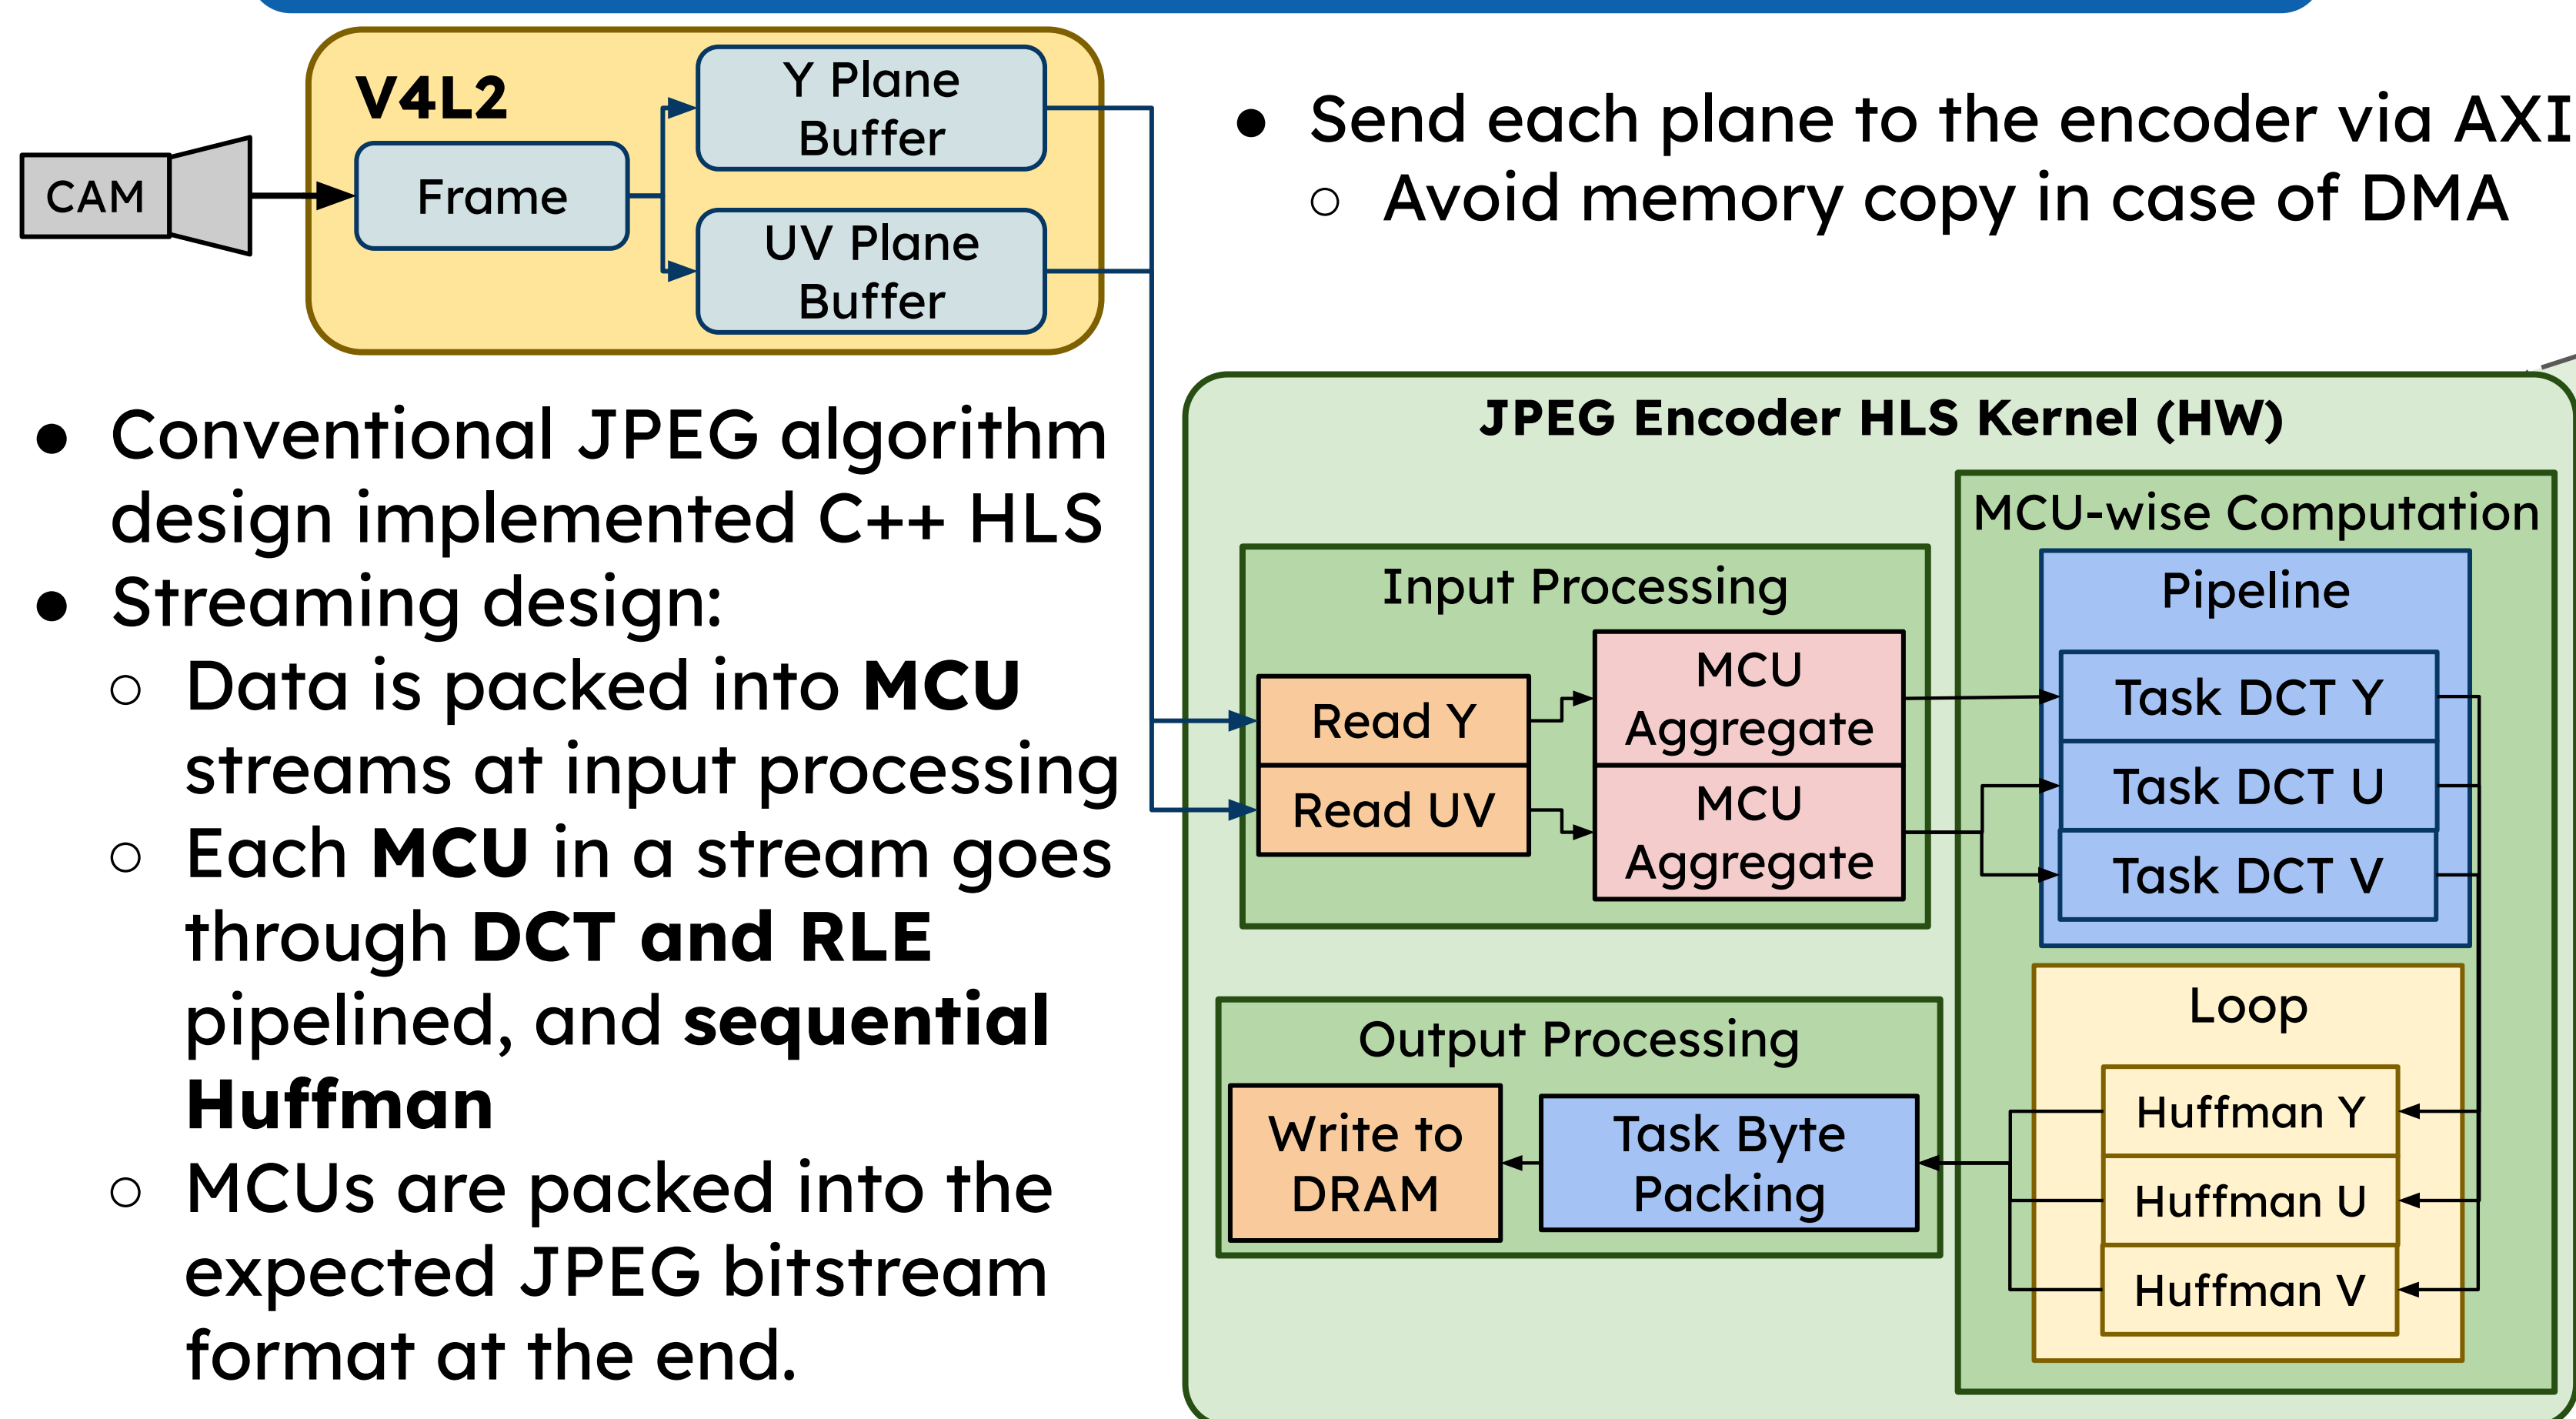
¹Nagoya University, ²Renesas Electronics,
³Map IV Inc.

E-mail: yuri_gpps1@ert1.jp

Motivation

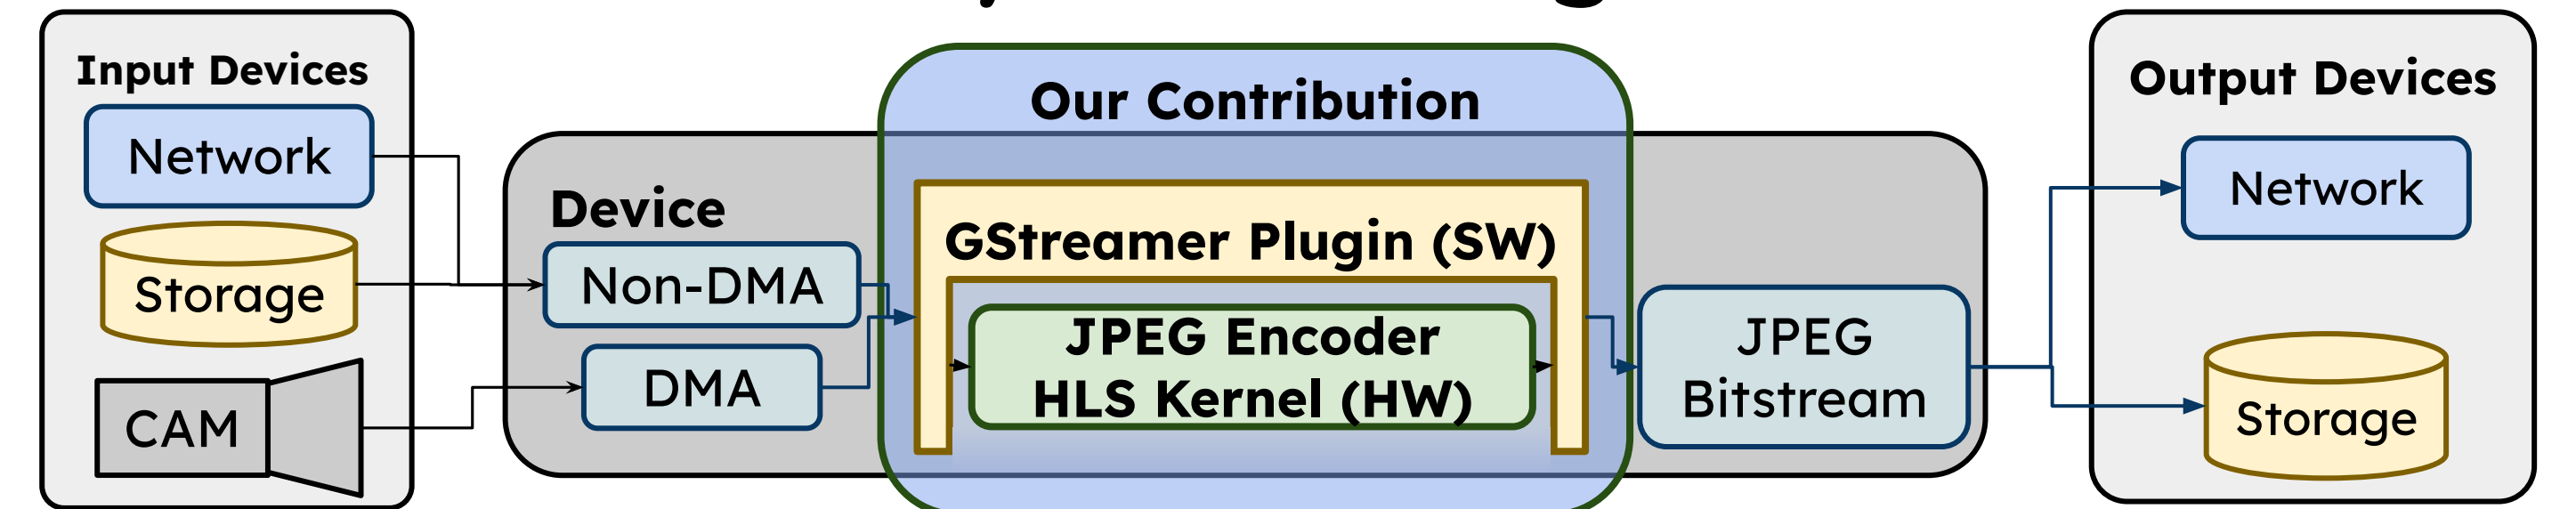


Hardware Architecture



Approach

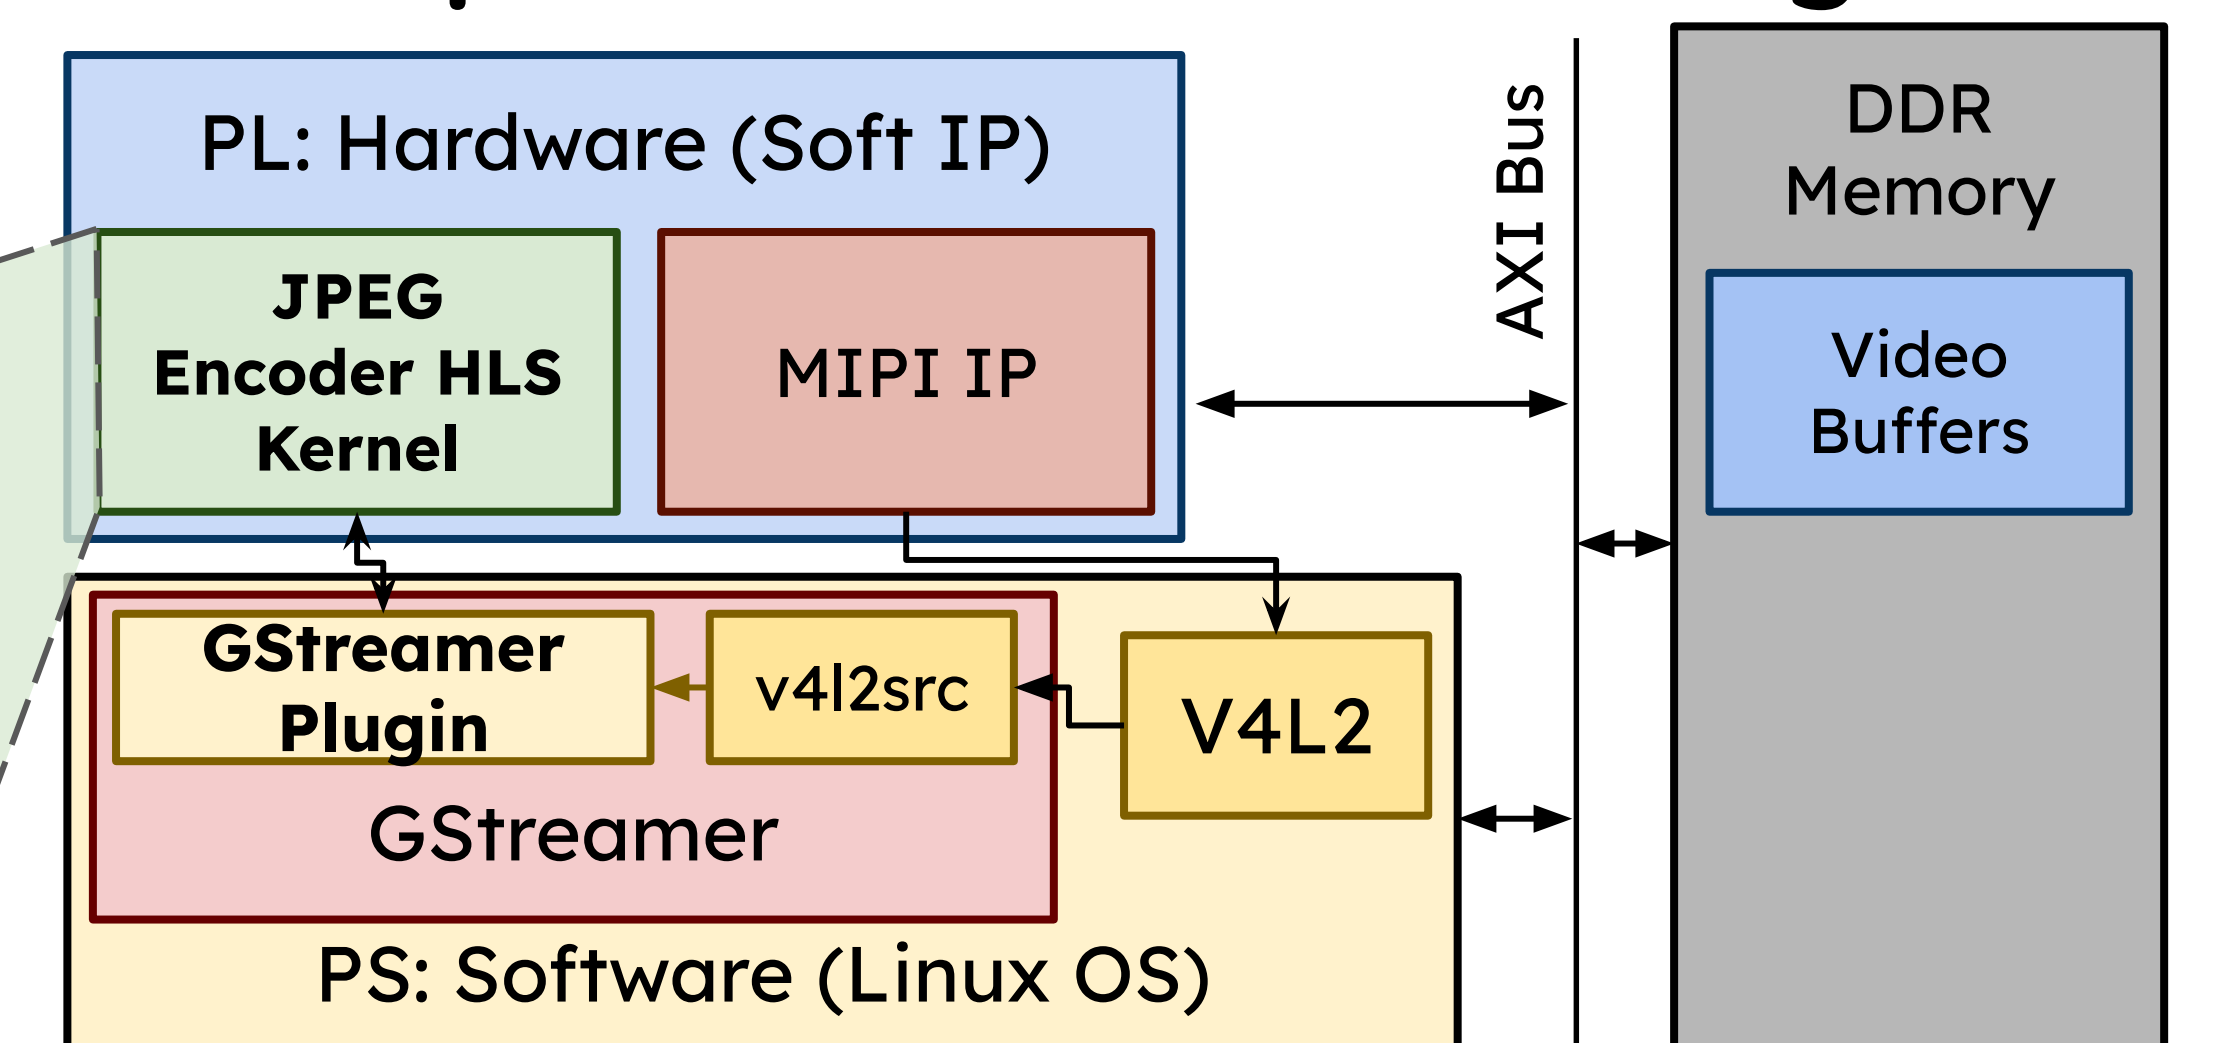
System Design



Contributions

- HW JPEG Encoder** implemented in *High-level Synthesis*.
- GStreamer plugin** interfaces with encoder.
 - Allows for mix and matching of pipeline elements for arbitrary media processing.
 - Easy usage of DMA capabilities.

Top-level Hardware Design

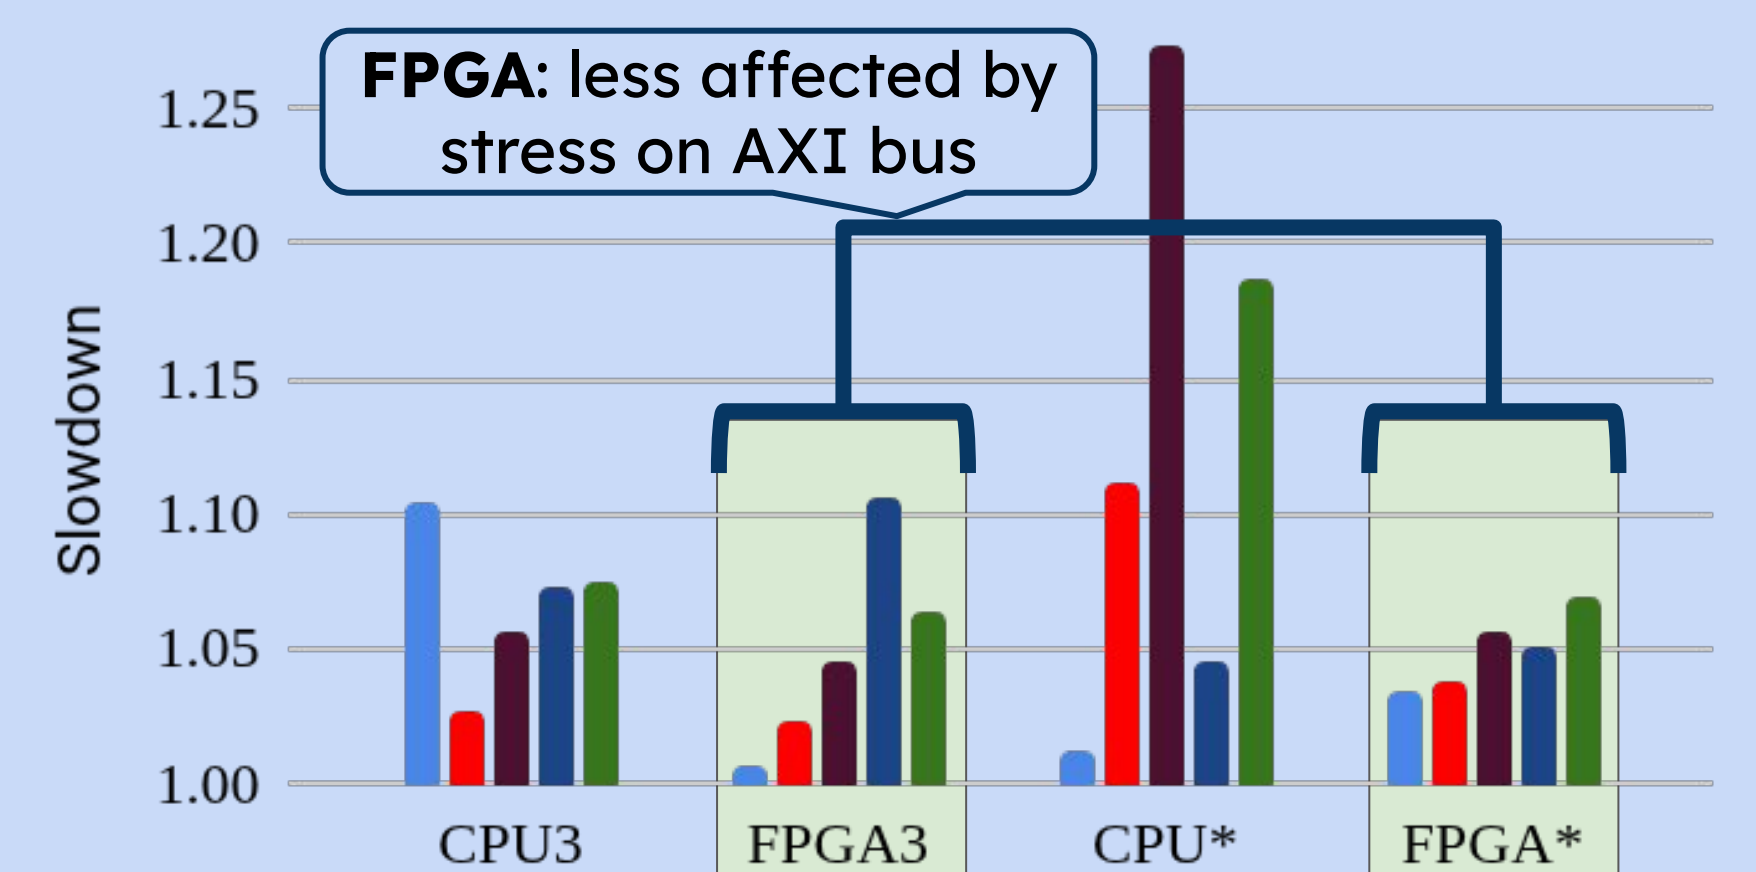
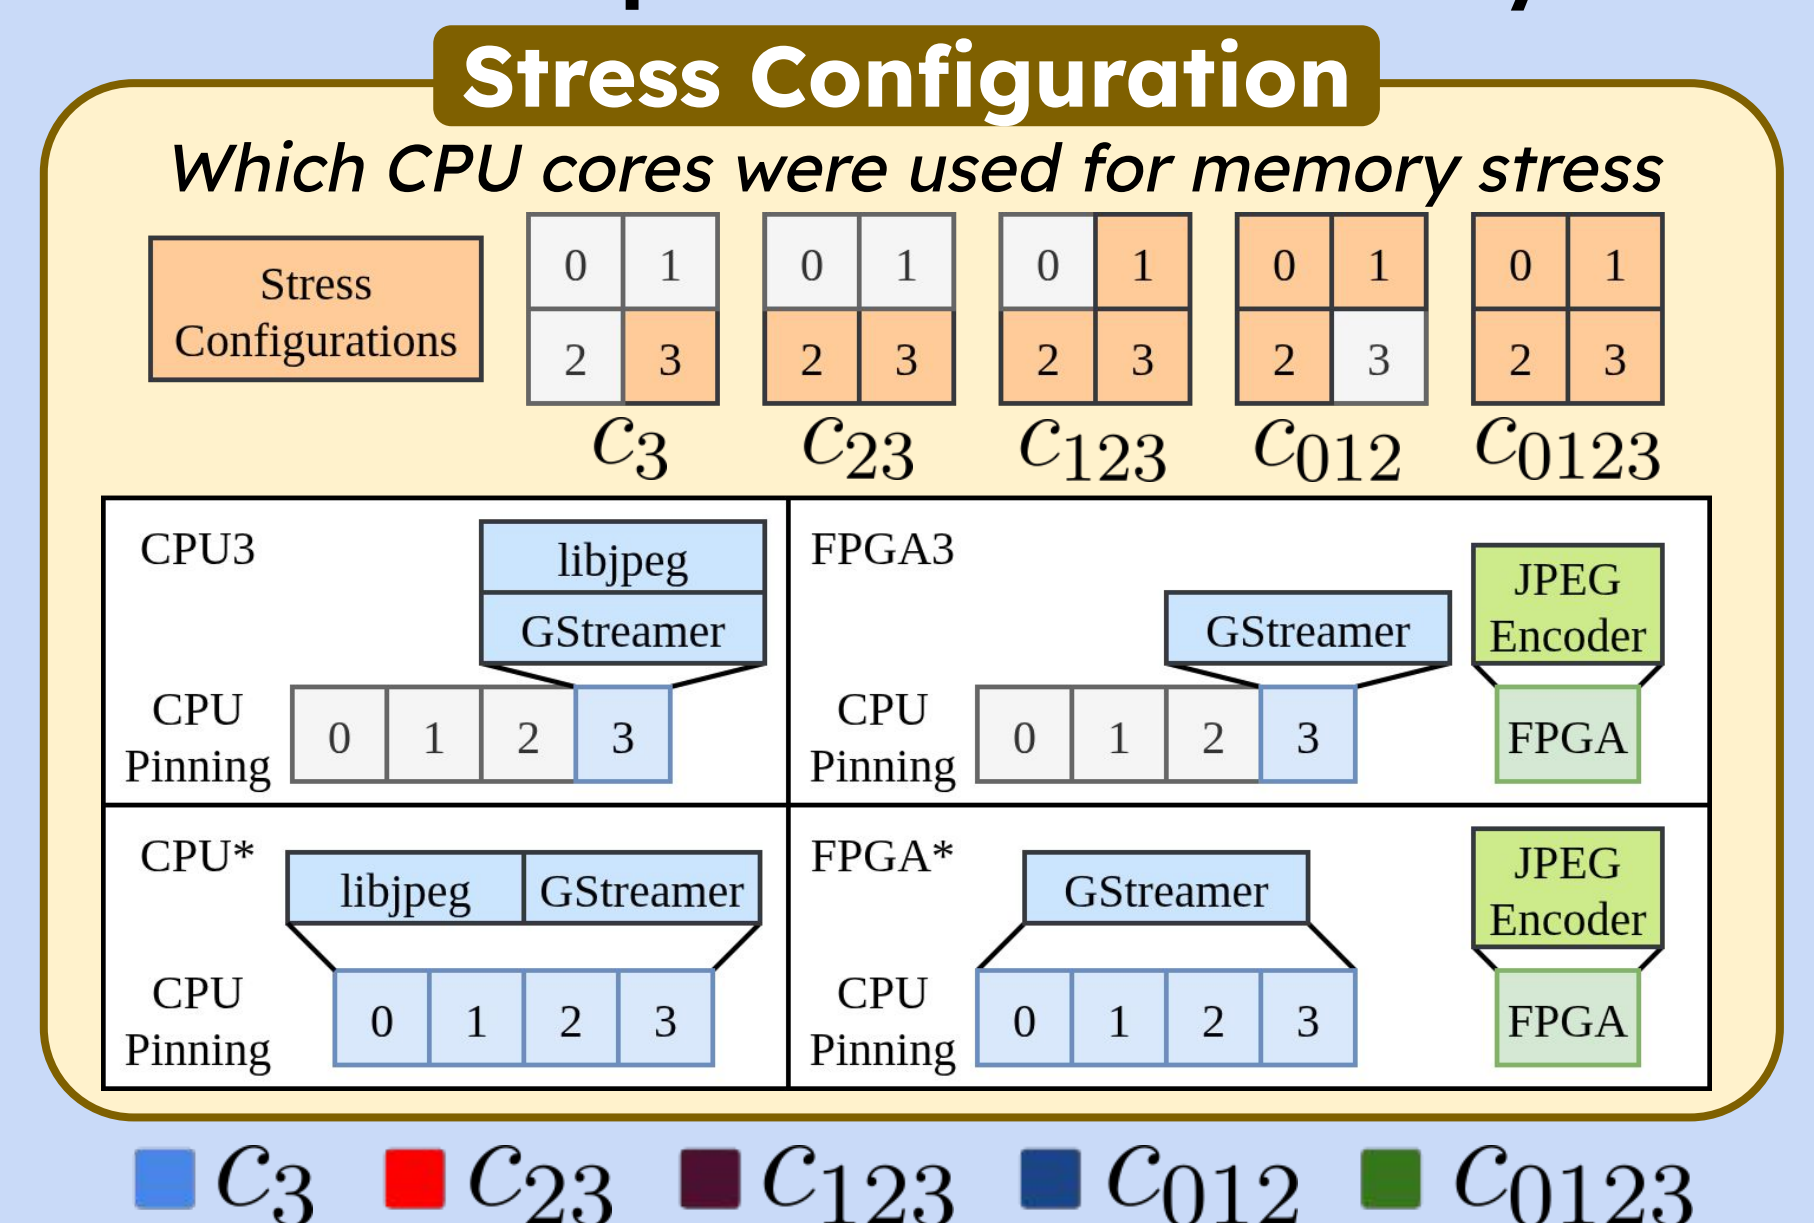


- Camera accessed through V4L2
- Accesses must go through AXI bus

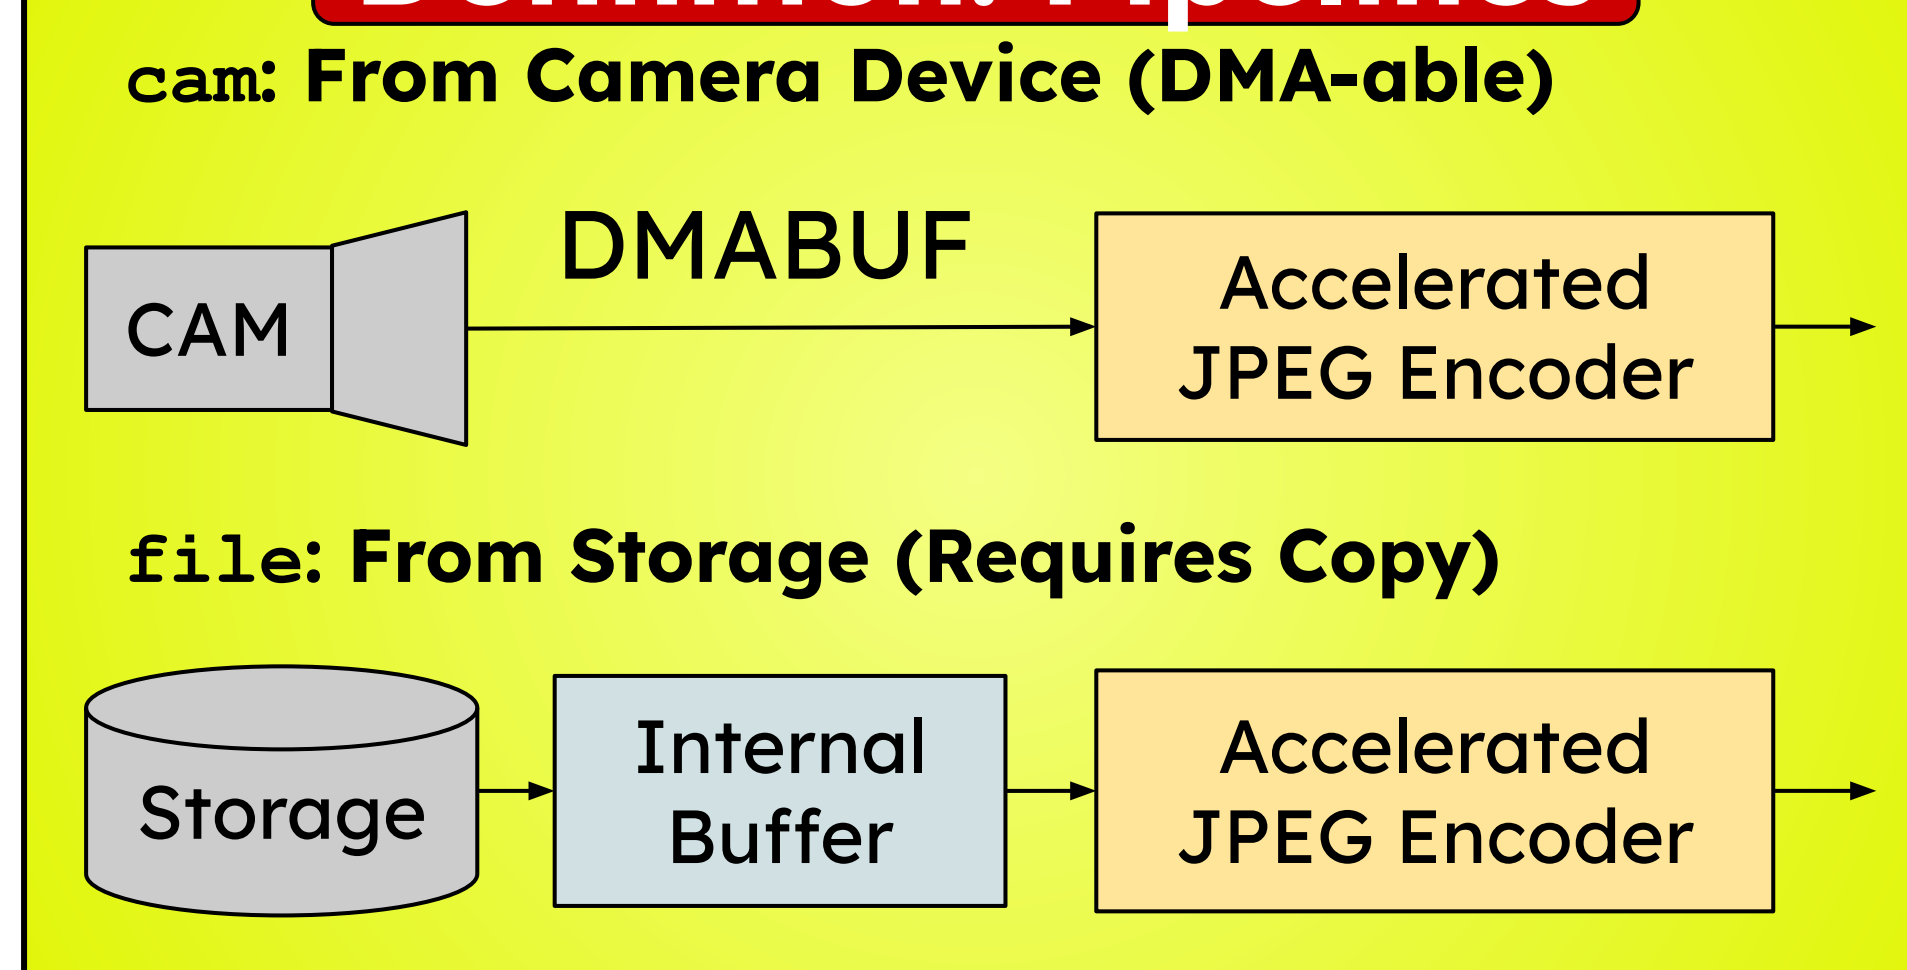
Experiments

Memory Stress Testing

Slowdown Comparison under Memory Stress



Definition: Pipelines



Performance

FPGA vs. CPU (jpegenc)

		CPU	FPGA	FPGA/CPU
cam	FPS	2.6	19.5	7.5
	CPU %	93	5	0.054
file	FPS	22.6	21.7	.96
	CPU %	107	14	0.13

Resource Utilization

JPEG (our) vs. Vitis Codec Libraries

	Freq.	BRAM	URAM	DSP	FF	LUT
JPEG ¹	250MHz	6	12	85	18150	13868
Webp ²	100MHz	81	46	834	71227	68906
Webp ²	250MHz	229	10	414	92030	68755
JXL ²	260MHz	584	122	644	270565	199537
PIK ²	200MHz	614	473	2398	549987	449995

¹Our work

²Vitis Codec Libraries

Webp Encoder²

- High Resource Usage**
 - Requires lowering target frequency to 100Hz

JPEG Encoder¹

- 7~25% of Webp Resources**
 - Can fit **four** JPEG encoders in place of a single Webp Encoder

Discussion and Conclusions

- Suitability for Edge
 - Low resource usage
 - Realtime performance
- Usability
 - GStreamer integration
 - Easy composition of media pipelines
- Increased Predictability
 - Resilience against memory stress

Future Work

- Streamlined flow for GStreamer integration with arbitrary HLS kernels

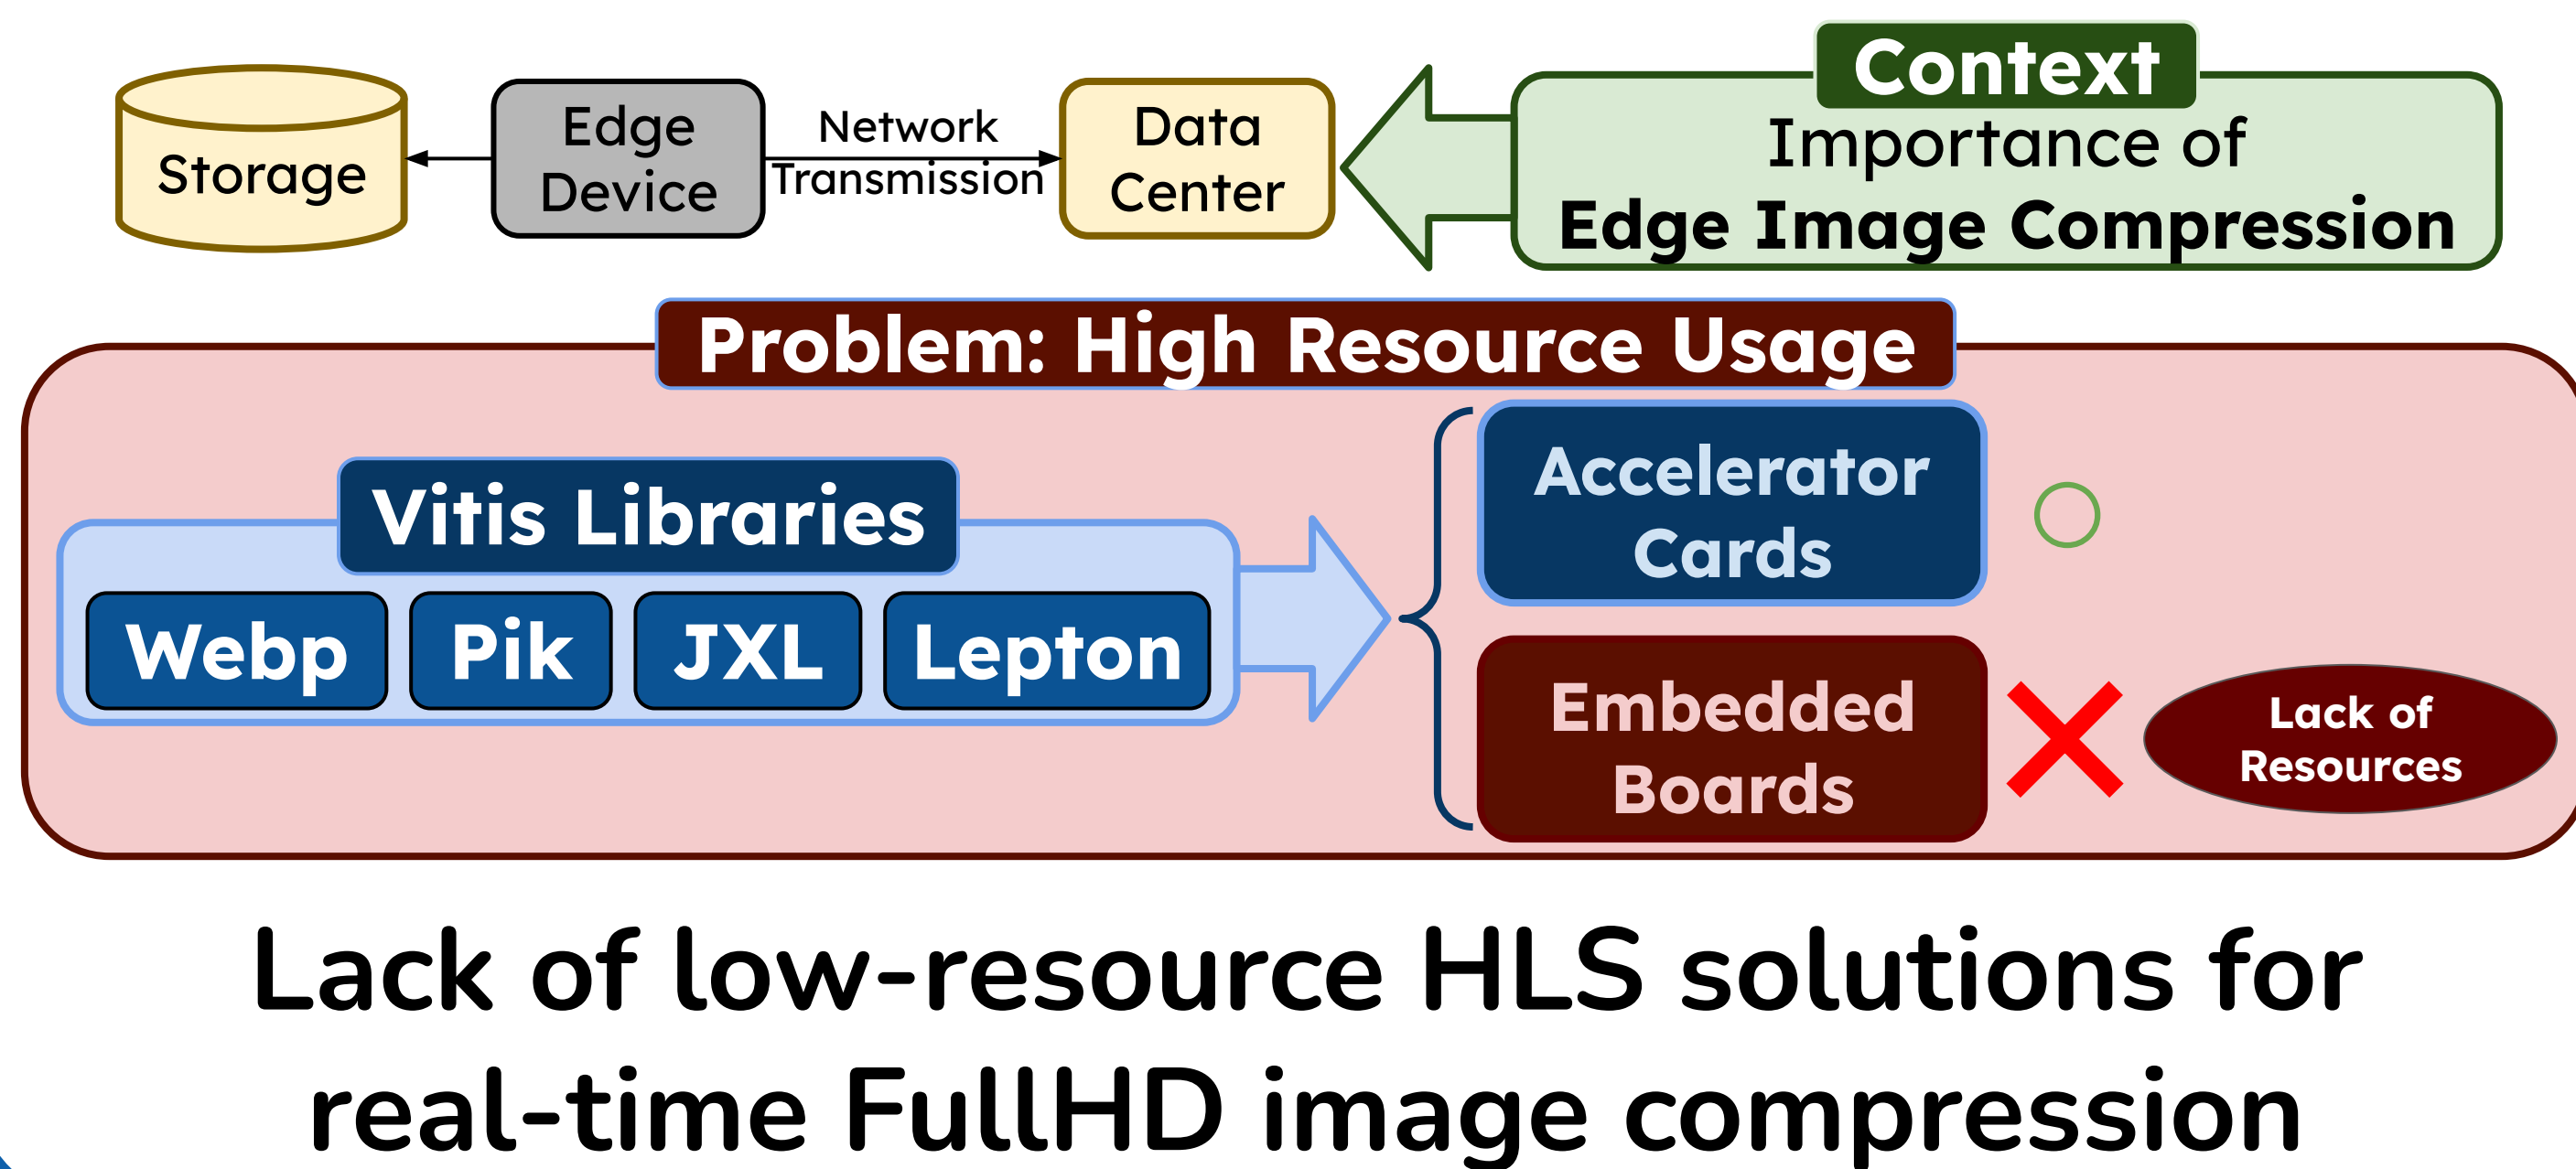
GStreamer-integrated HLS-based JPEG Encoder for Edge FPGA SoCs

Yuri Guimaraes Pereira Primo da Silva¹, Shinya Honda¹,
Sugako Otani¹², Masato Eda¹, Abraham Monrroy Cano³

¹Nagoya University, ²Renesas Electronics,
³Map IV Inc.

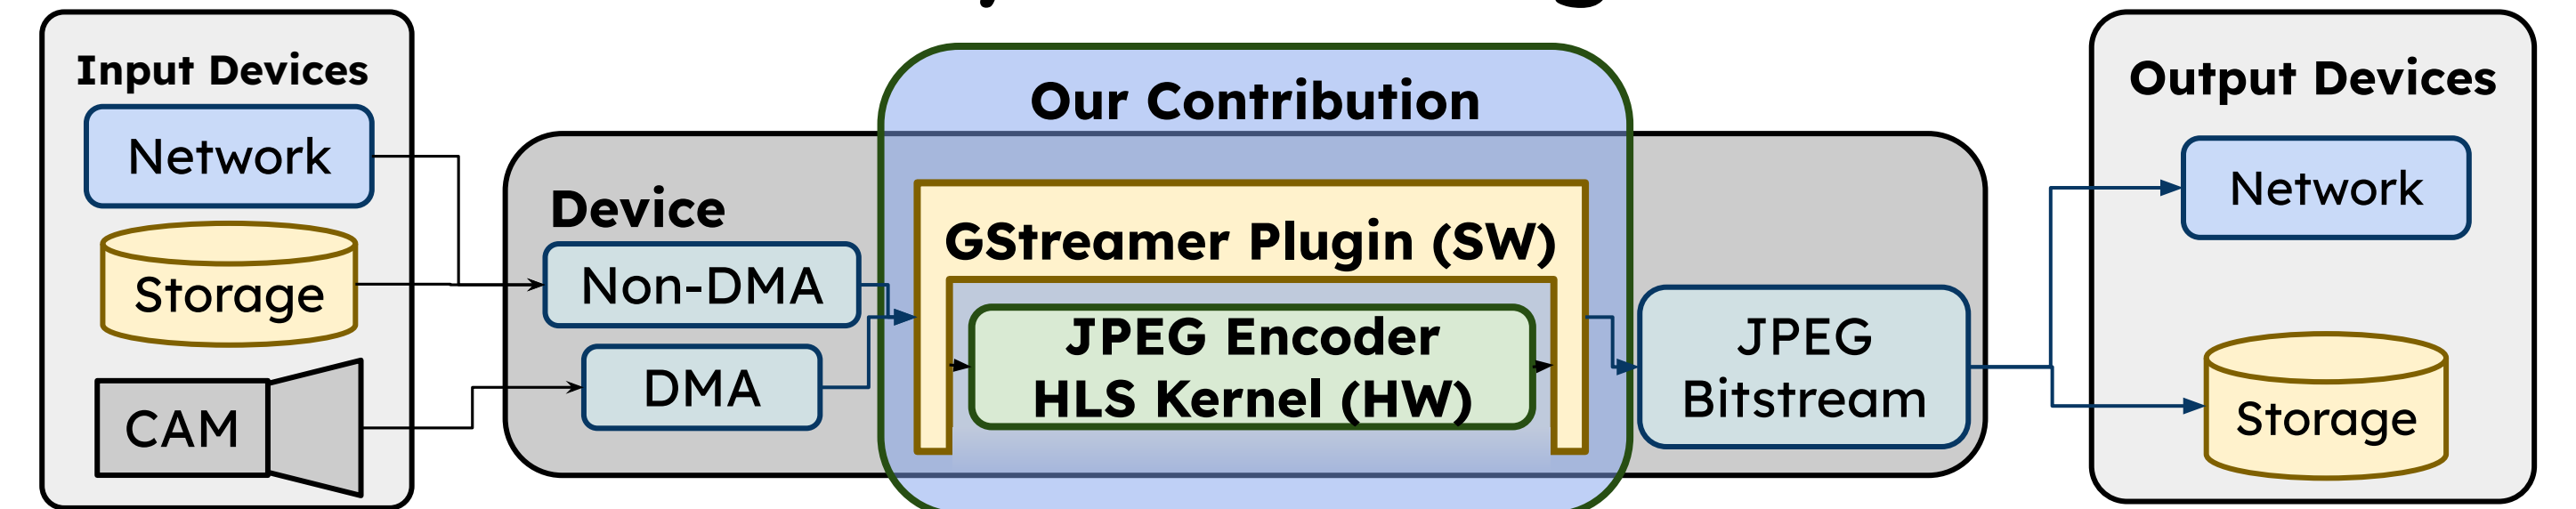
E-mail: yuri_gpps1@ert1.jp

Motivation



Approach

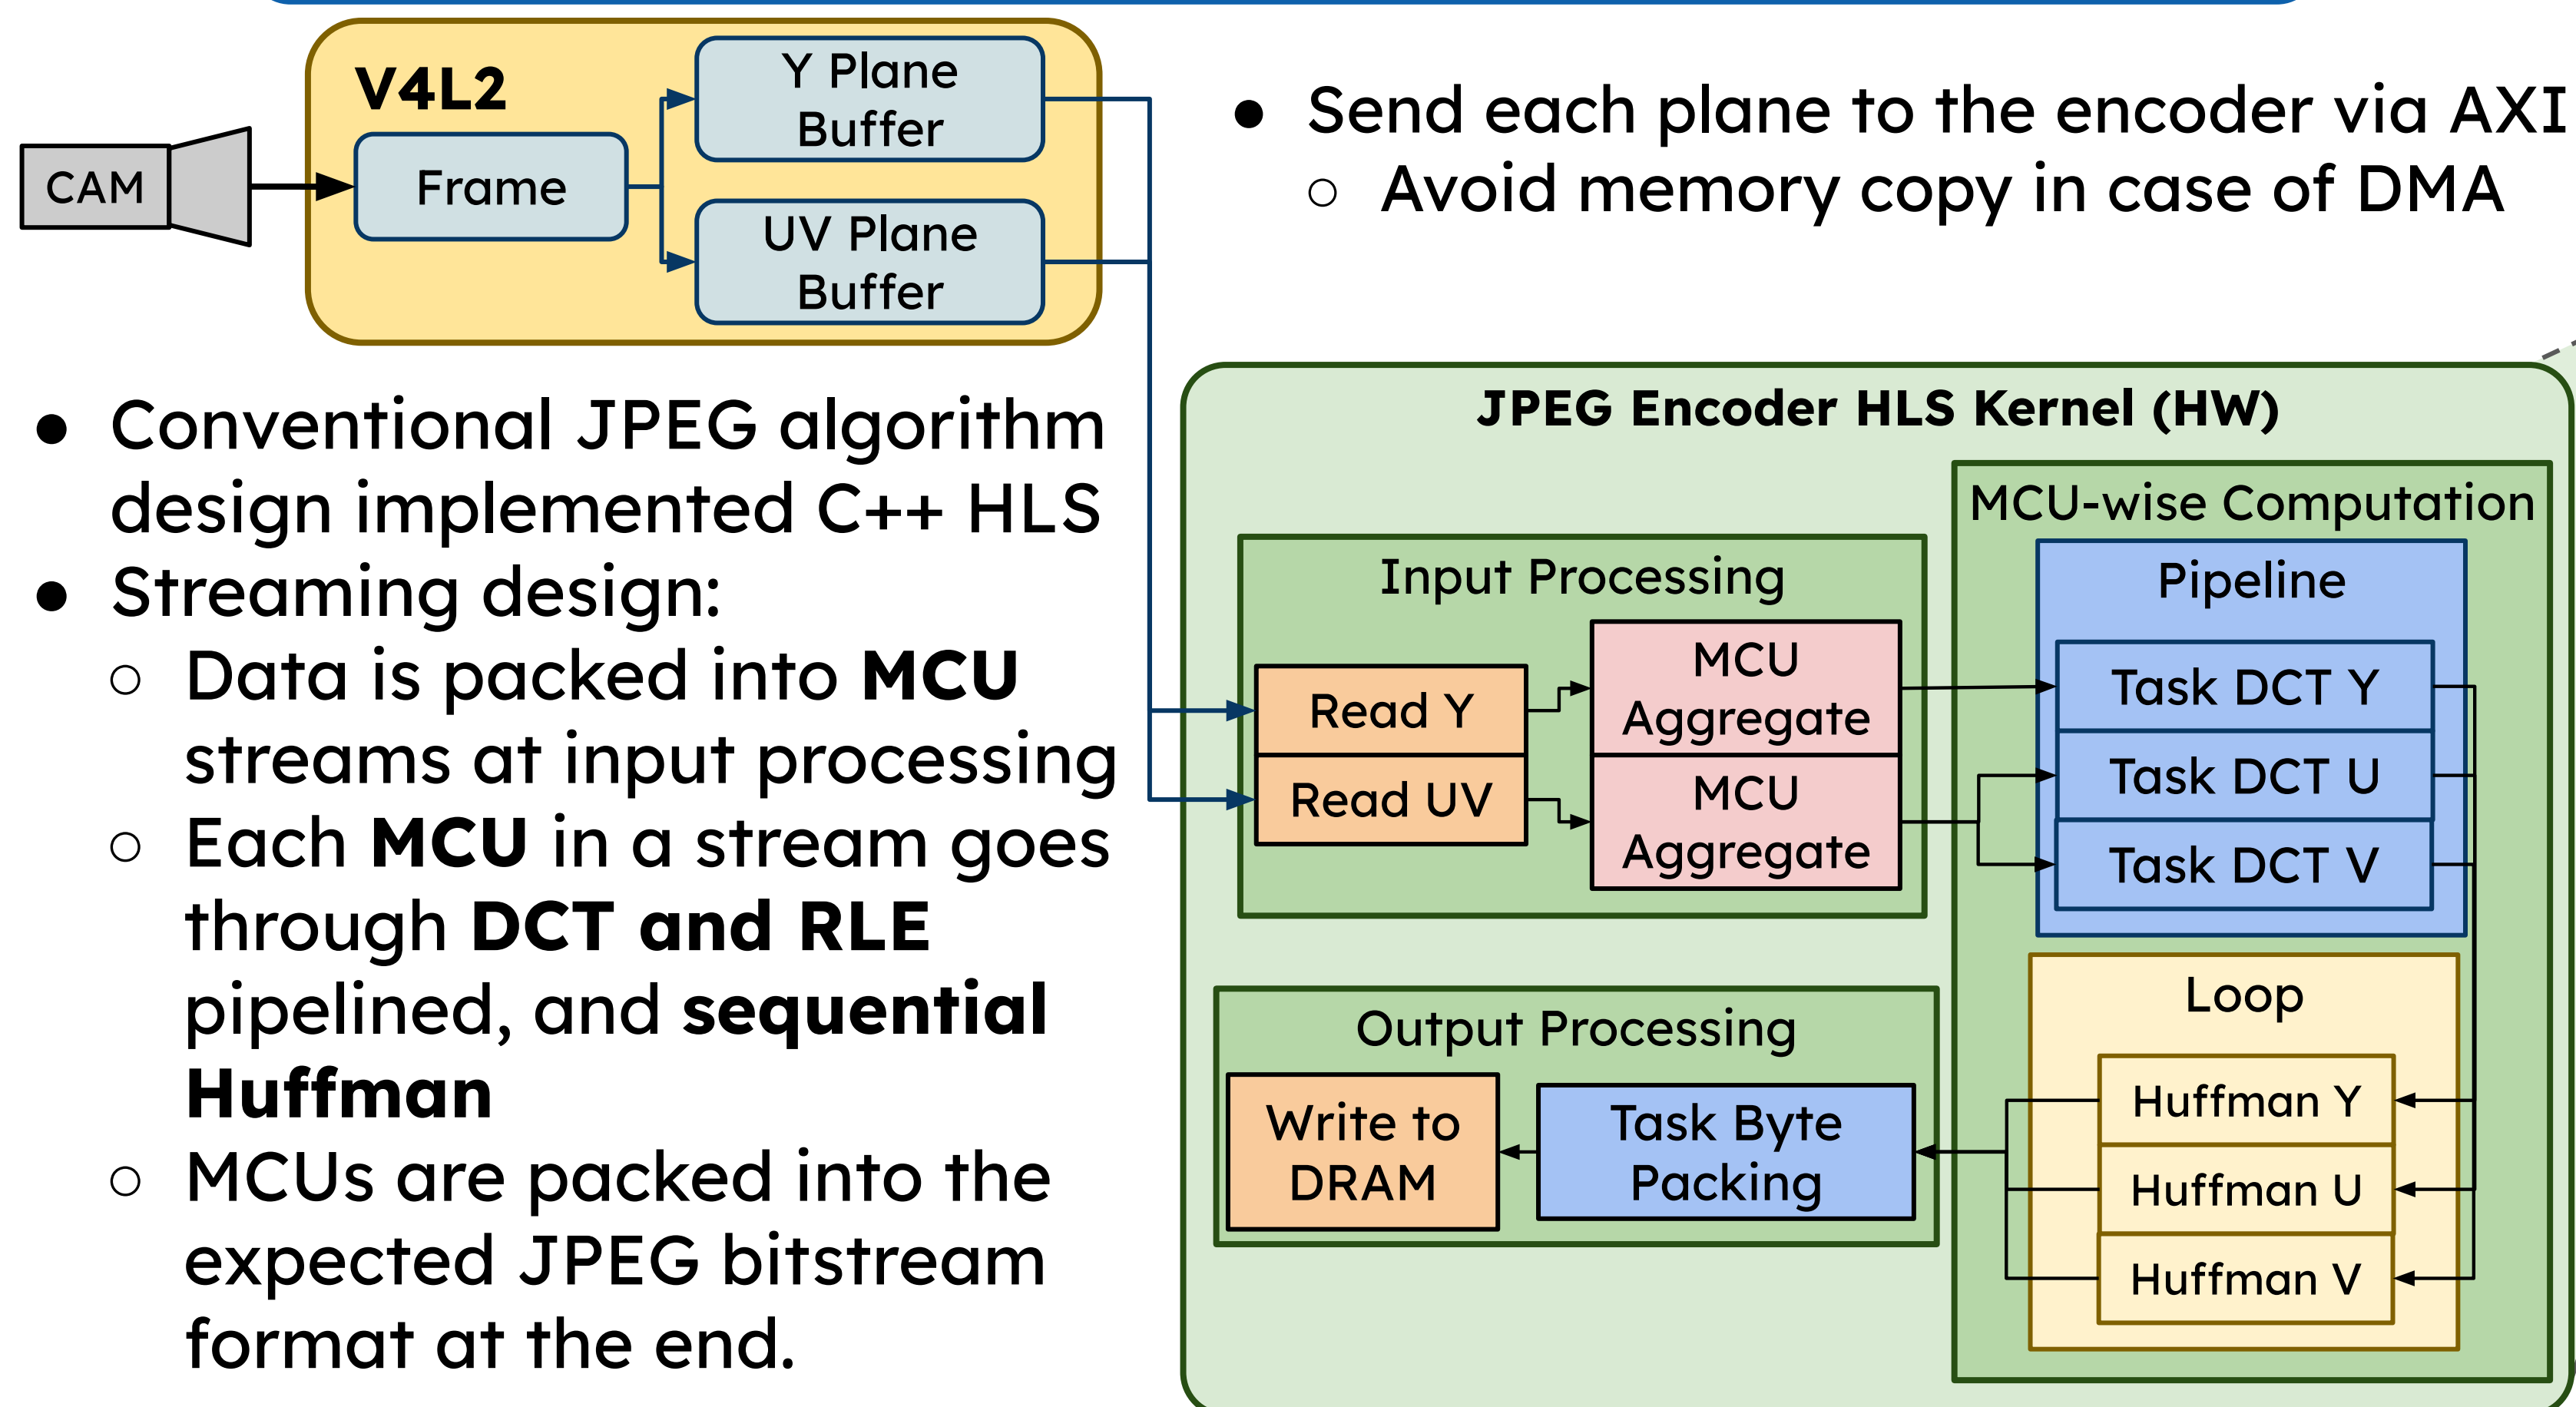
System Design



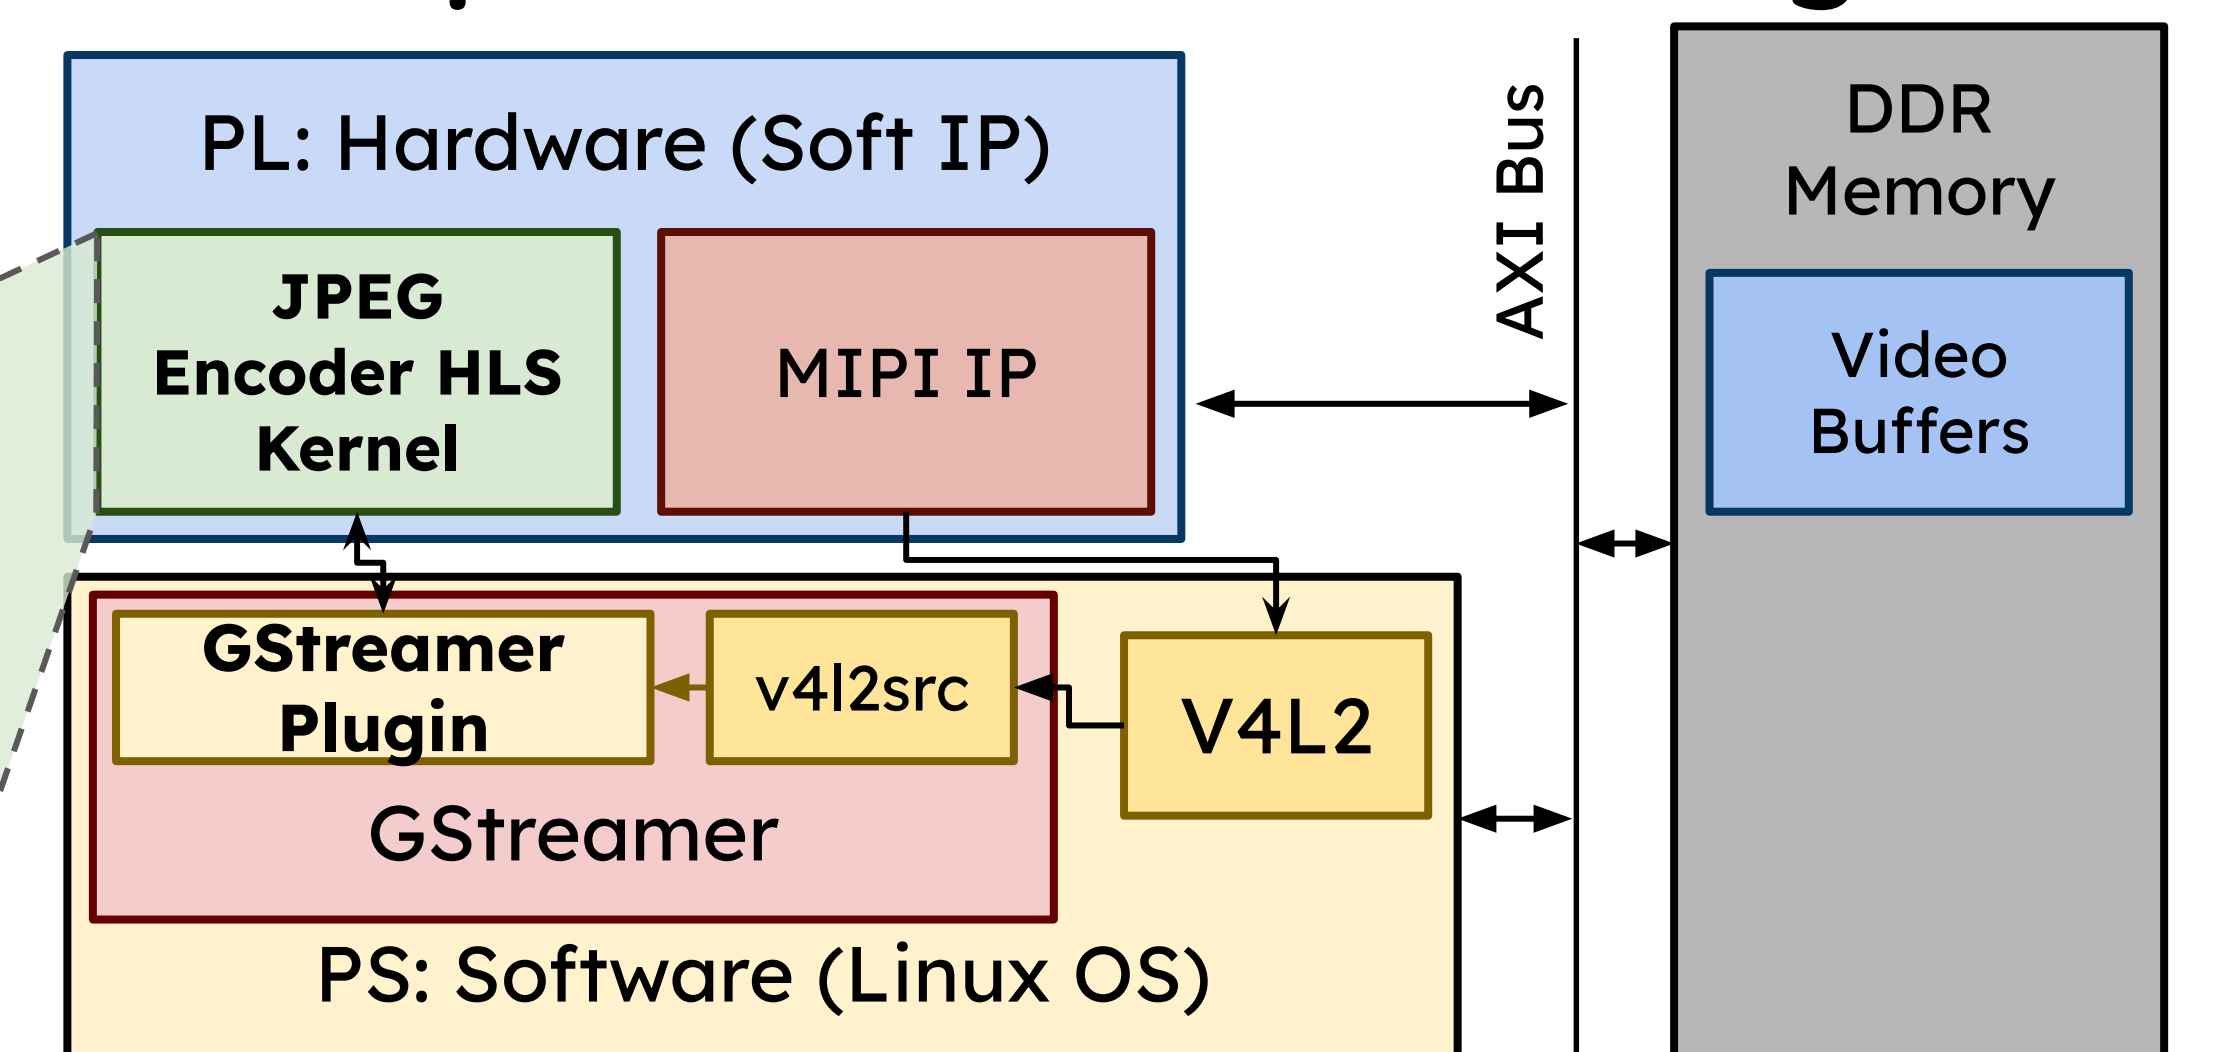
Contributions

- **HW JPEG Encoder** implemented in *High-level Synthesis*.
- **GStreamer plugin** interfaces with encoder.
 - Allows for mix and matching of pipeline elements for arbitrary media processing.
 - Easy usage of DMA capabilities.

Hardware Architecture



Top-level Hardware Design

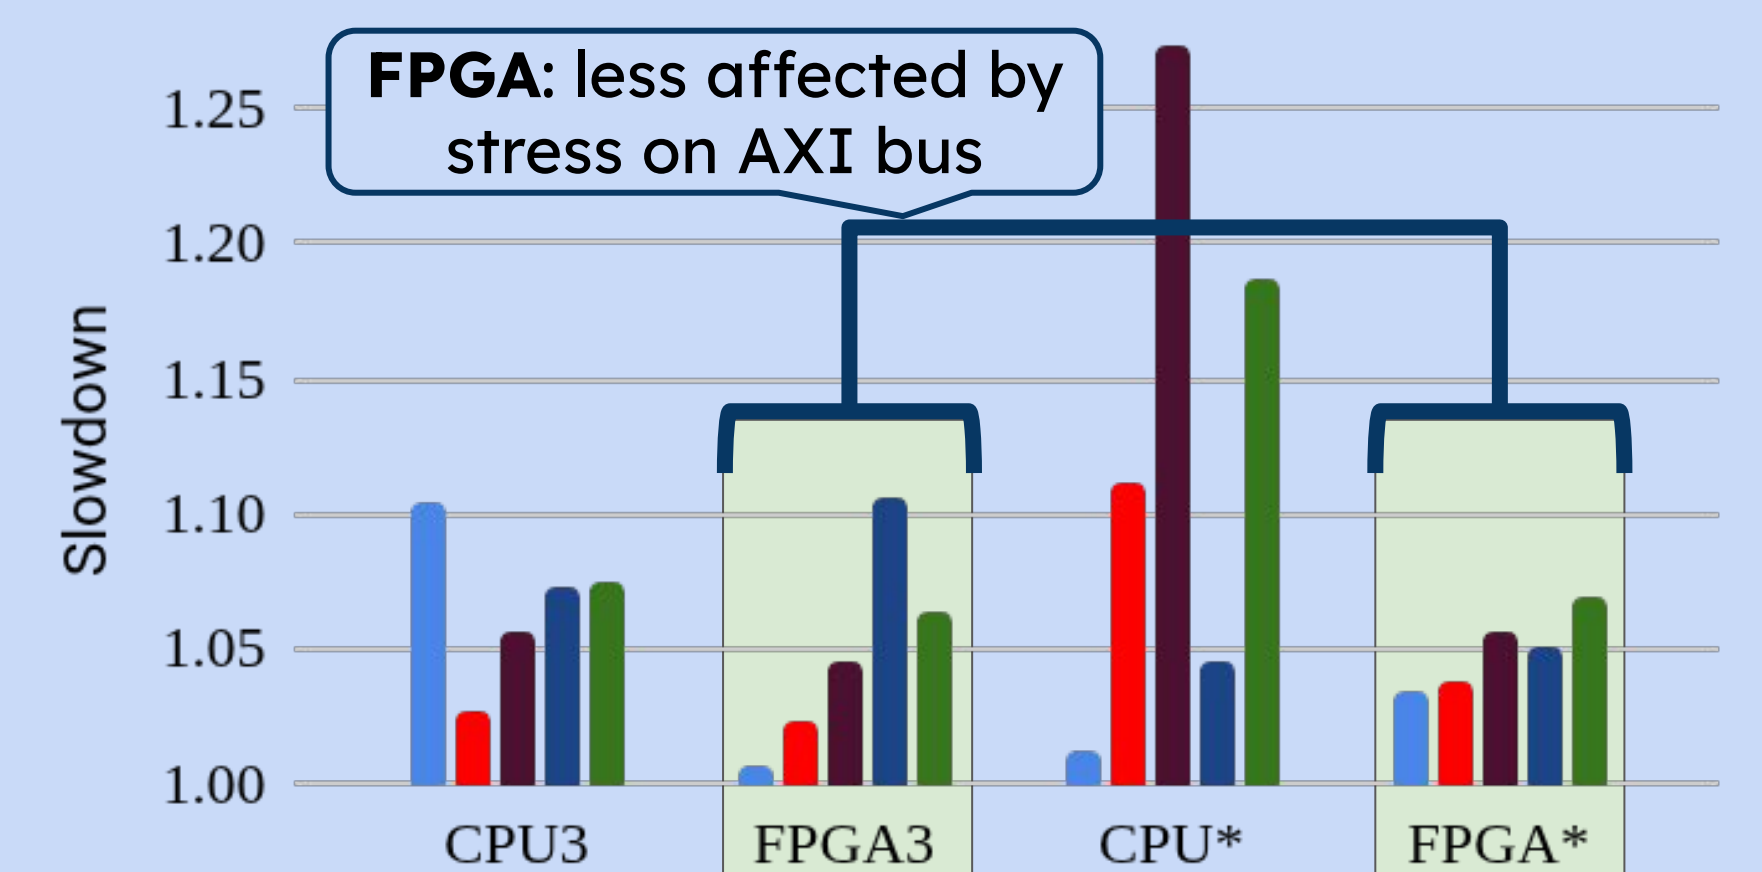
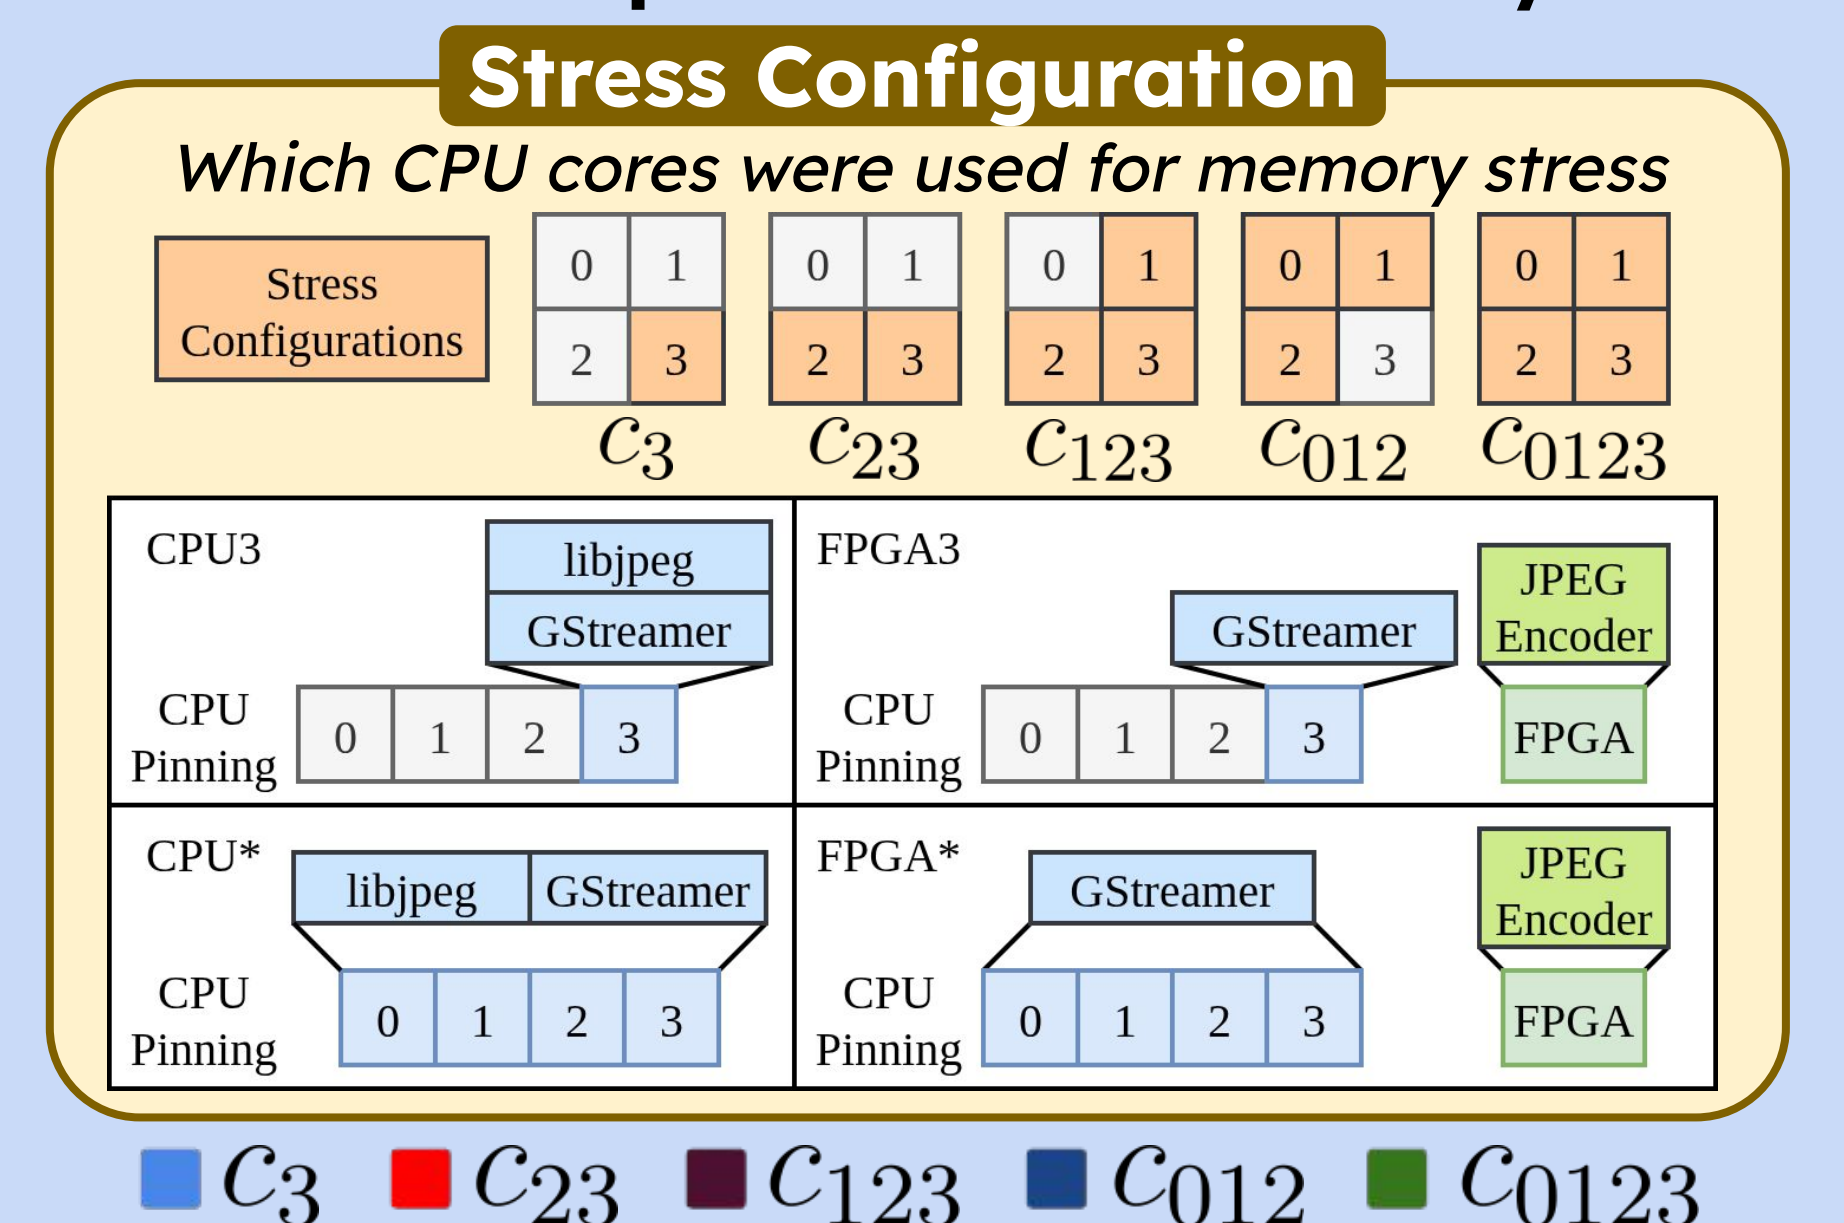


- Camera accessed through V4L2
 - Accesses must go through AXI bus

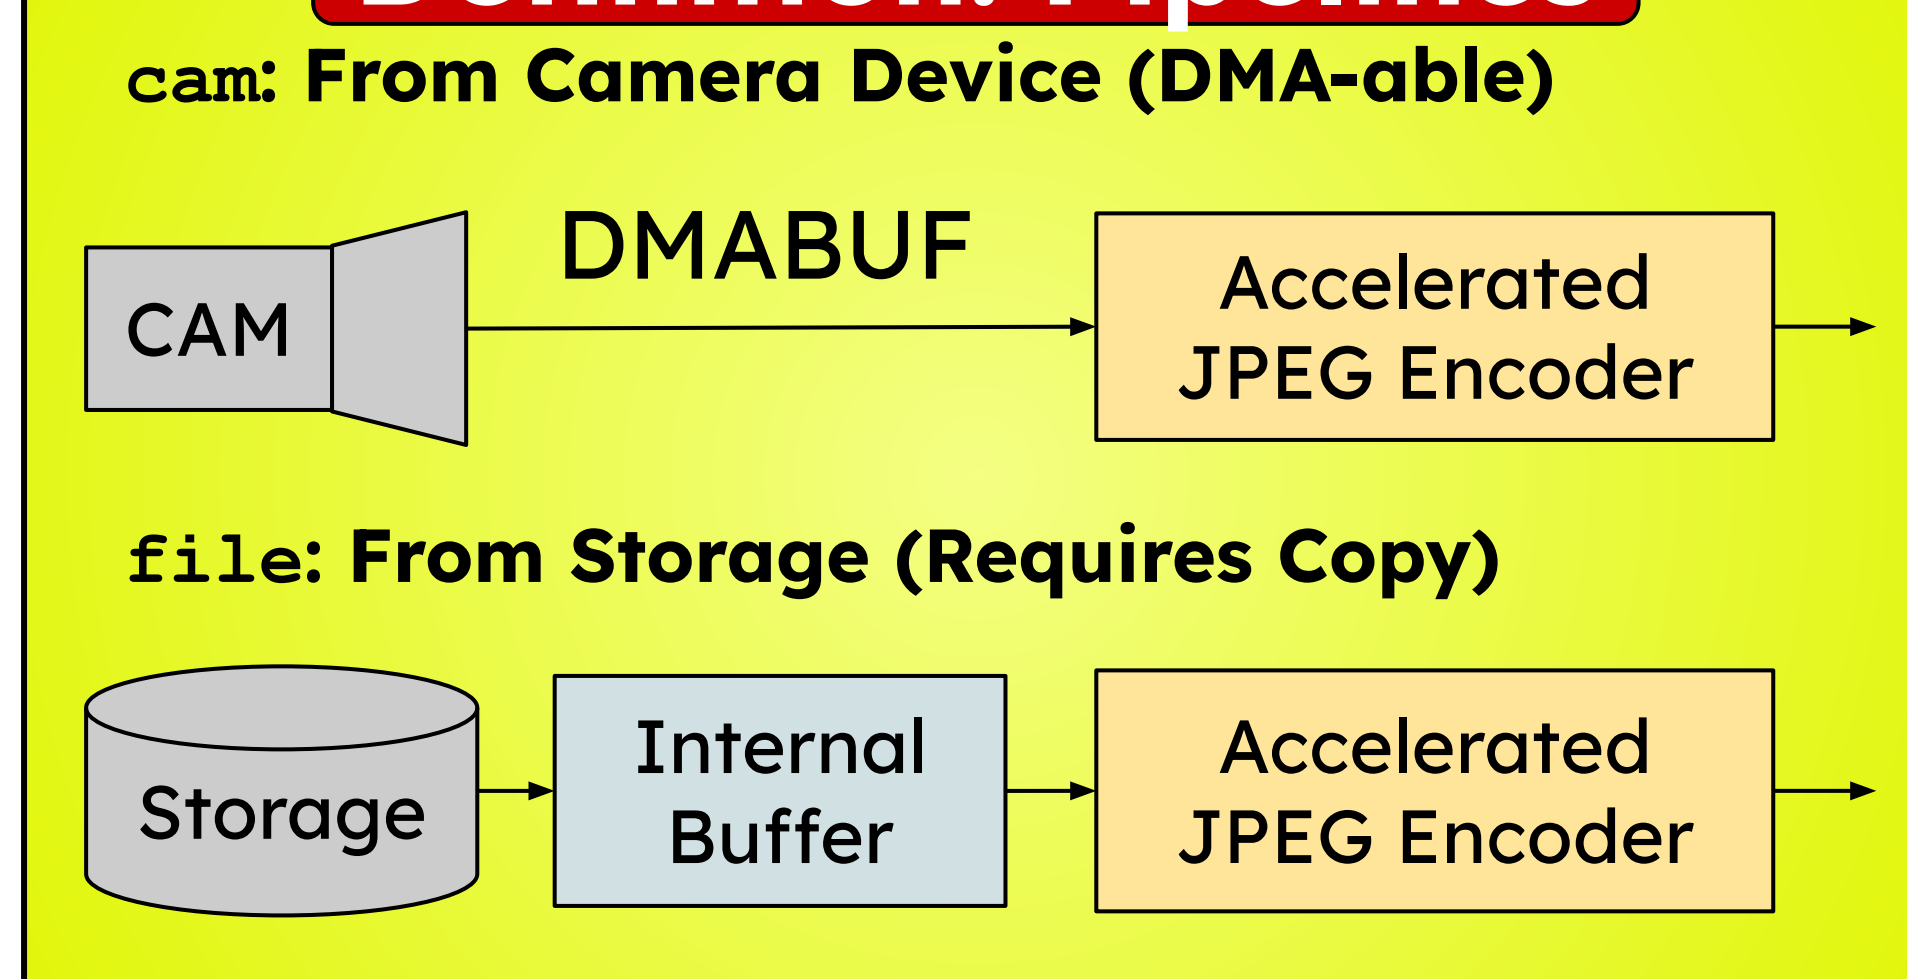
Experiments

Memory Stress Testing

Slowdown Comparison under Memory Stress



Definition: Pipelines



Performance

FPGA vs. CPU (jpegenc)

		CPU	FPGA	FPGA/CPU
cam	FPS	2.6	19.5	7.5
	CPU %	93	5	0.054
file	FPS	22.6	21.7	.96
	CPU %	107	14	0.13

Resource Utilization

JPEG (our) vs. Vitis Codec Libraries

	Freq.	BRAM	URAM	DSP	FF	LUT
JPEG ¹	250MHz	6	12	85	18150	13868
Webp ²	100MHz	81	46	834	71227	68906
Webp ²	250MHz	229	10	414	92030	68755
JXL ²	260MHz	584	122	644	270565	199537
PIK ²	200MHz	614	473	2398	549987	449995

¹Our work

²Vitis Codec Libraries

Webp Encoder²

- **High Resource Usage**
 - Requires lowering target frequency to 100Hz

JPEG Encoder¹

- **7~25% of Webp Resources**
 - Can fit **four** JPEG encoders in place of a **single** Webp Encoder

Discussion and Conclusions

1. Suitability for Edge
 - Low resource usage
 - Realtime performance
2. Usability
 - GStreamer integration
 - Easy composition of media pipelines
3. Increased Predictability
 - Resilience against memory stress

Future Work

- Streamlined flow for GStreamer integration with arbitrary HLS kernels

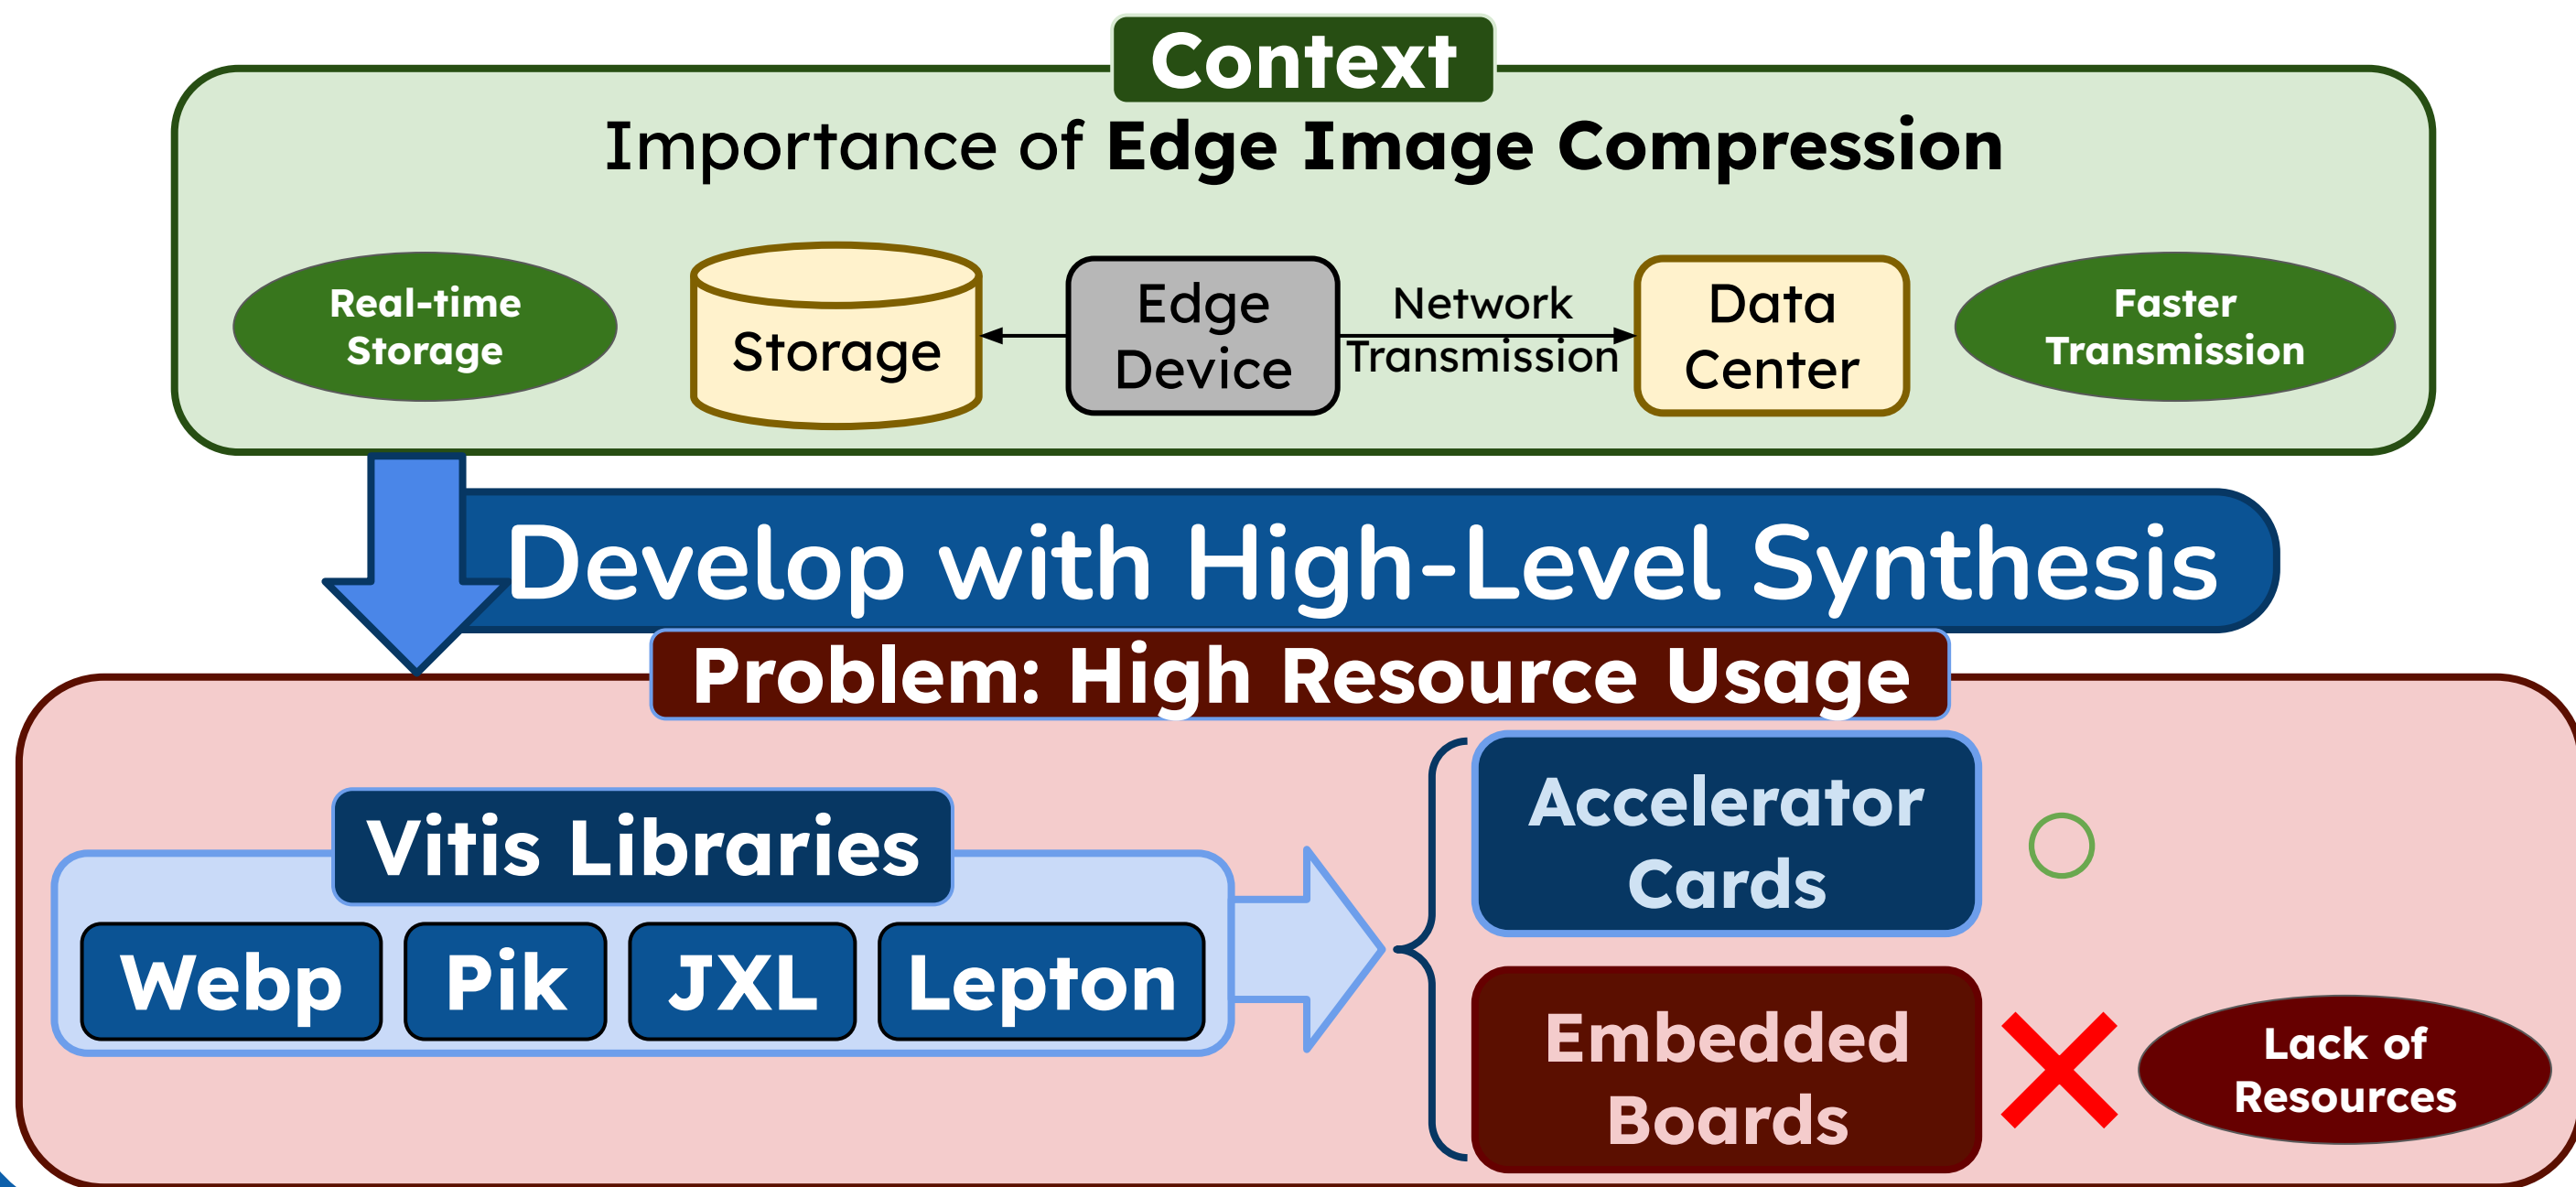
GStreamer-integrated HLS-based JPEG Encoder for Edge FPGA SoCs

Yuri Guimaraes Pereira Primo da Silva¹, Shinya Honda¹,
Sugako Otani¹², Masato Eda¹, Abraham Monrroy Cano³

¹Nagoya University, ²Renesas Electronics,
³Map IV Inc.

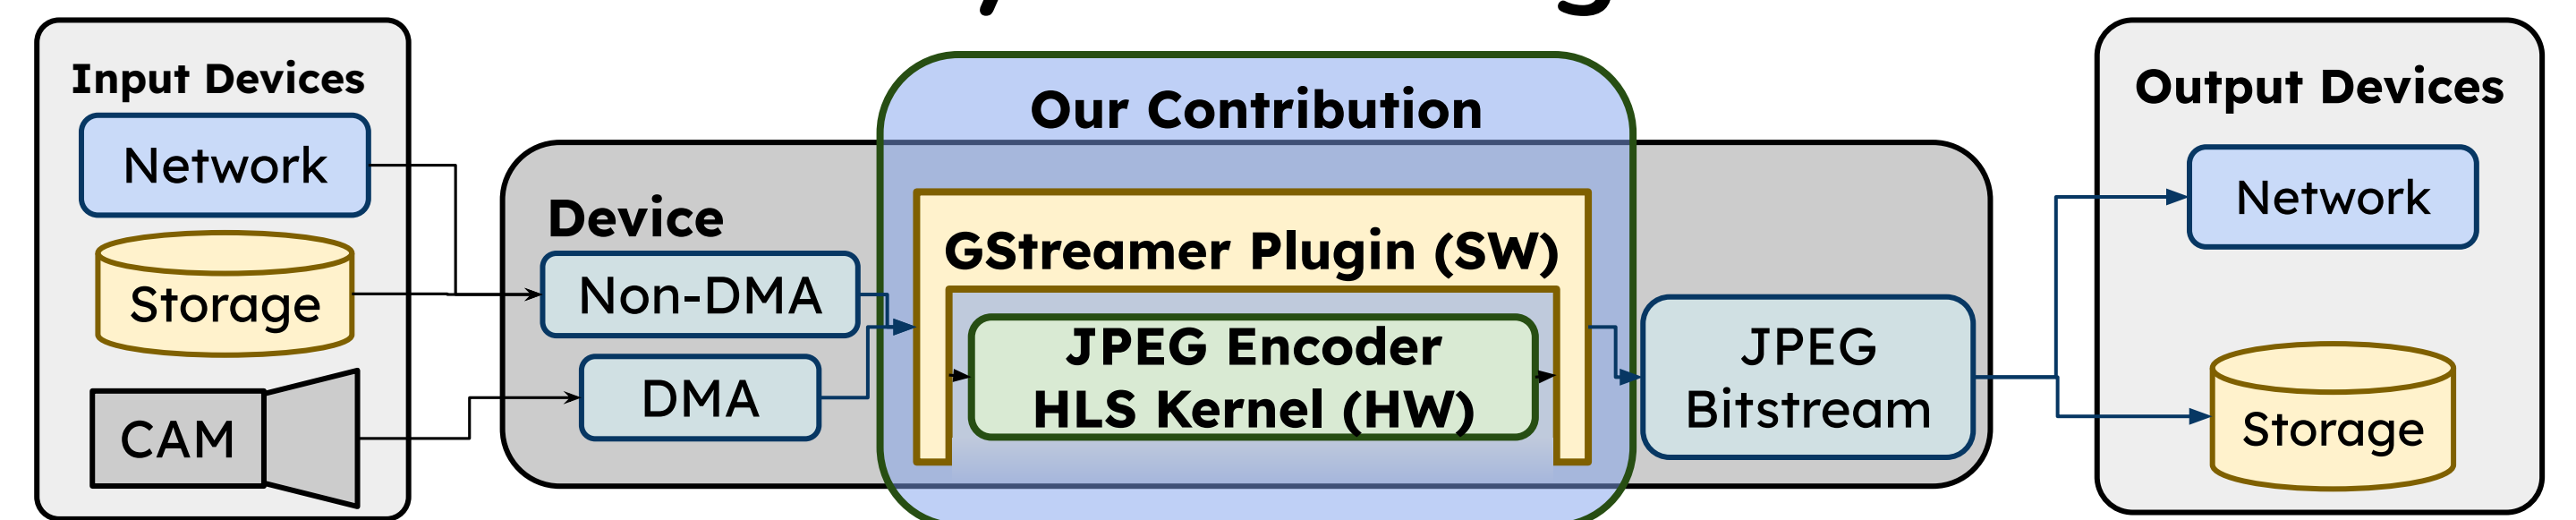
E-mail: yuri_gpps1@ert1.jp

Motivation



Approach

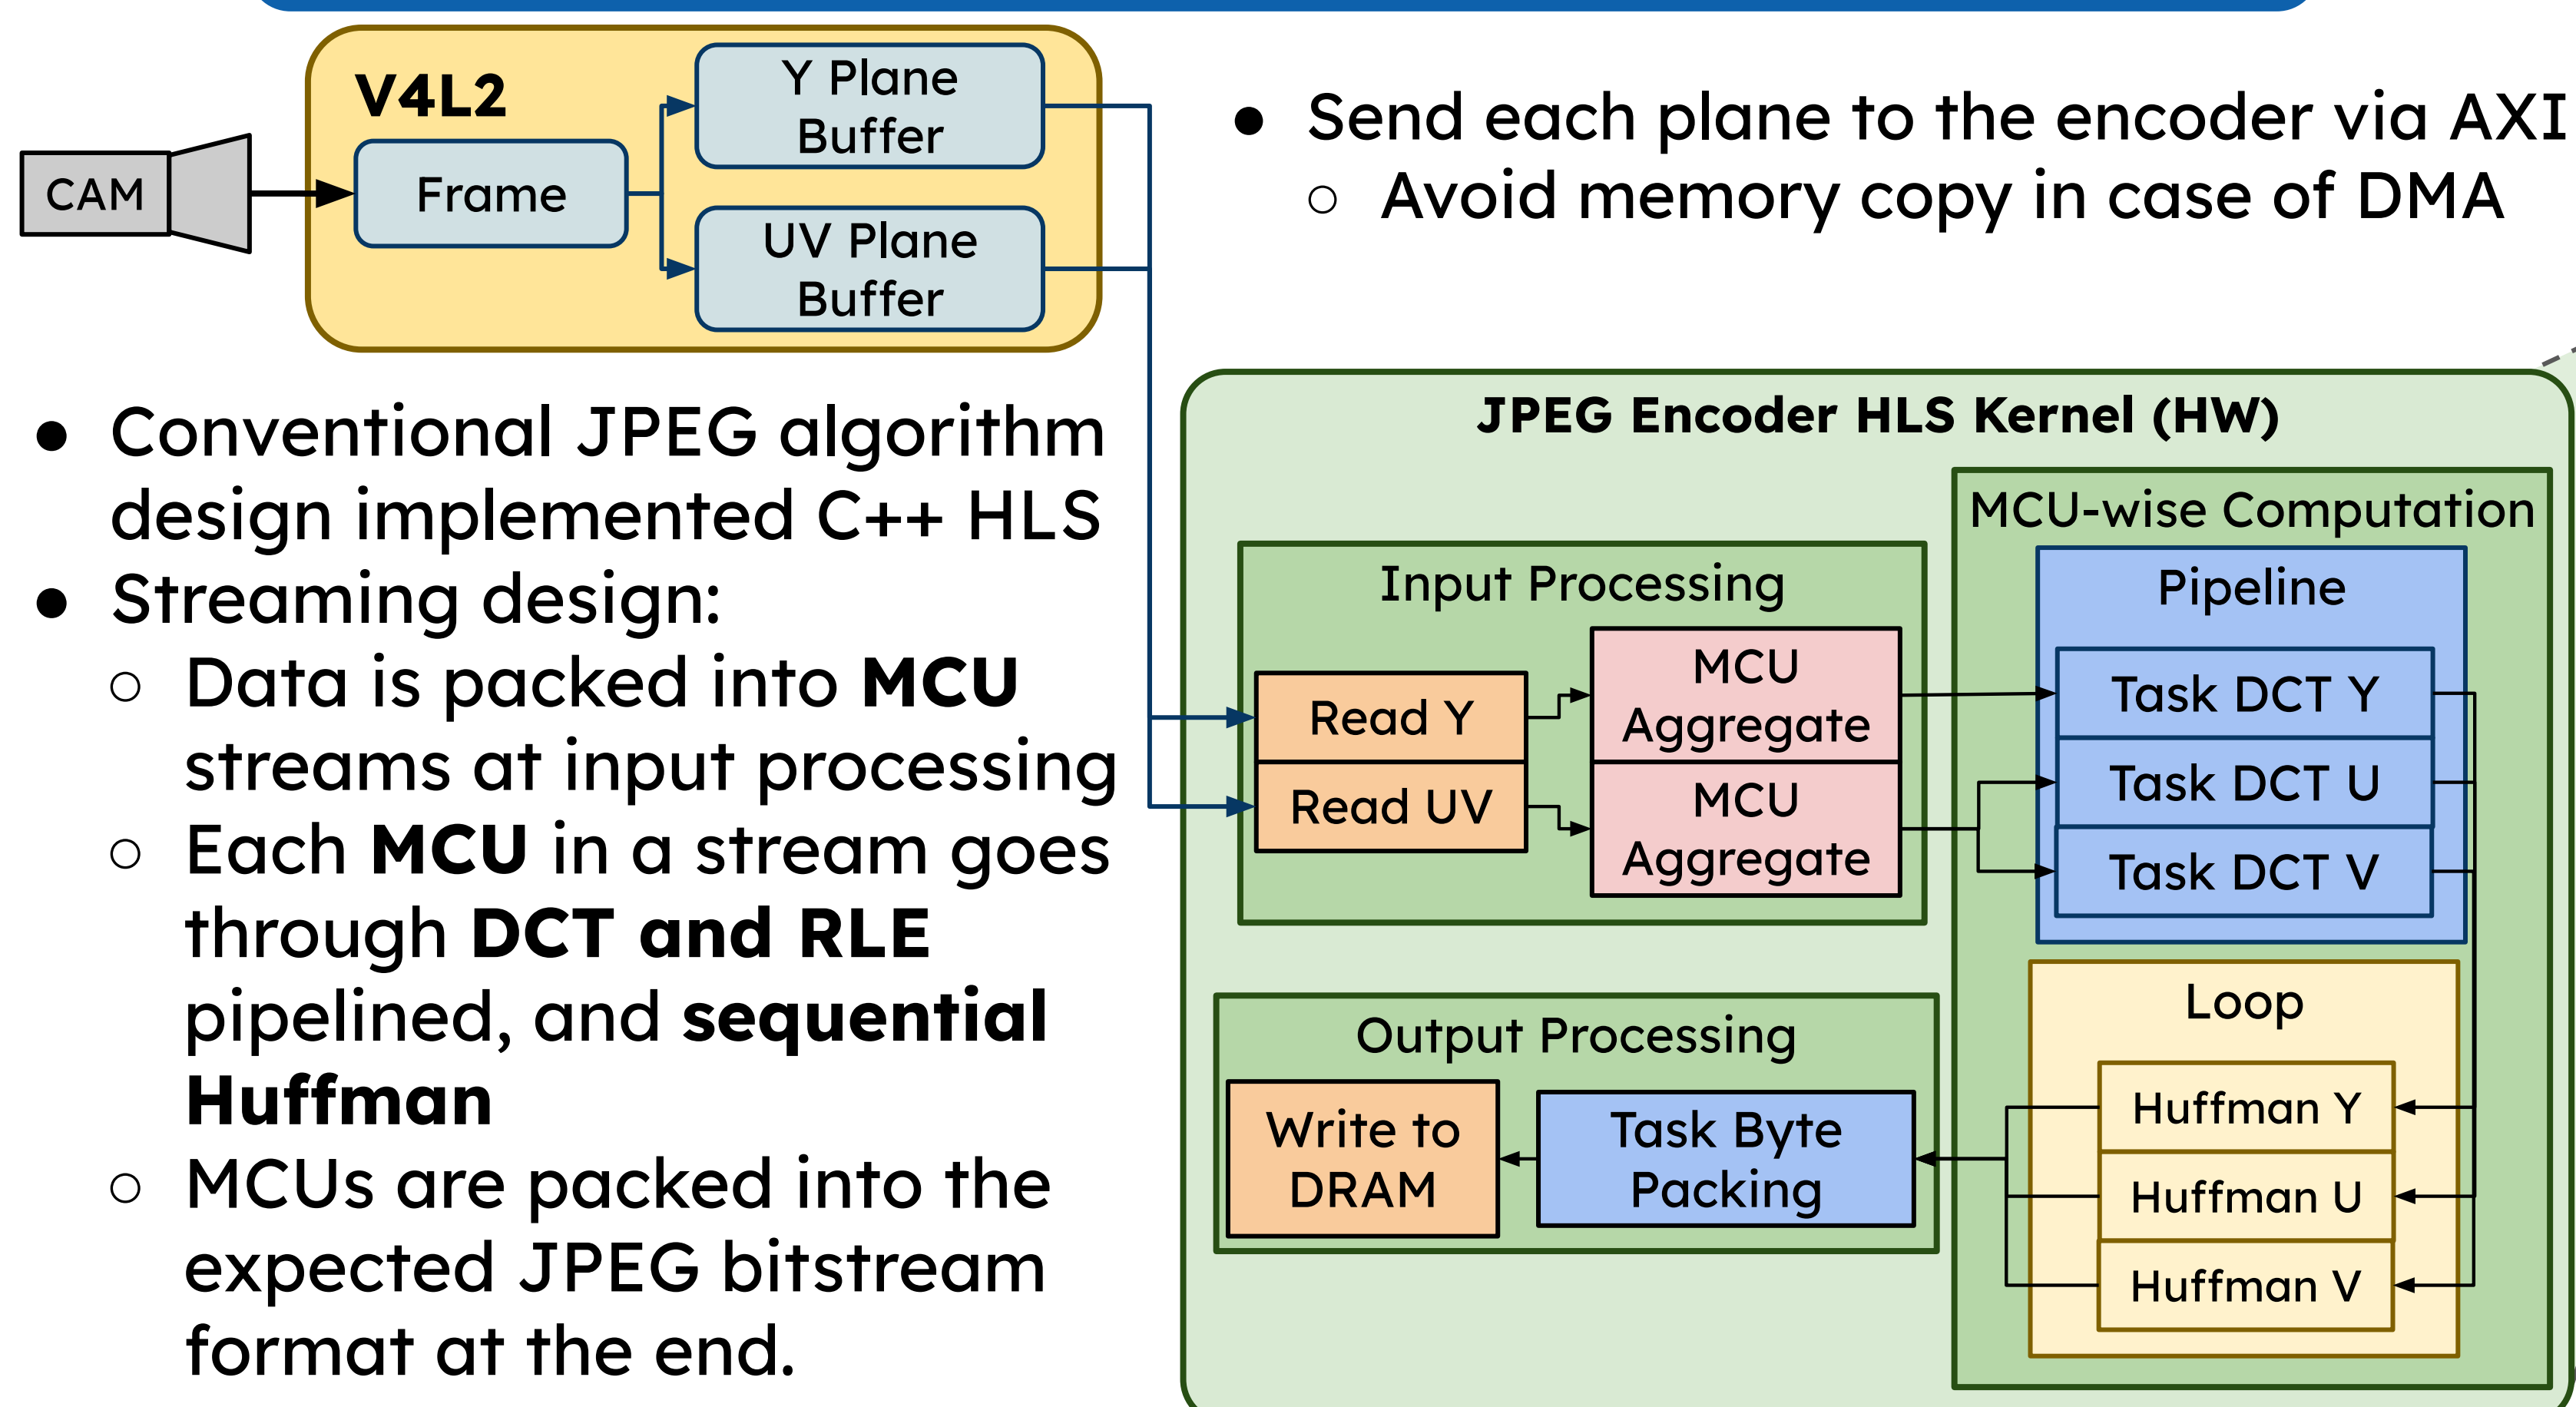
System Design



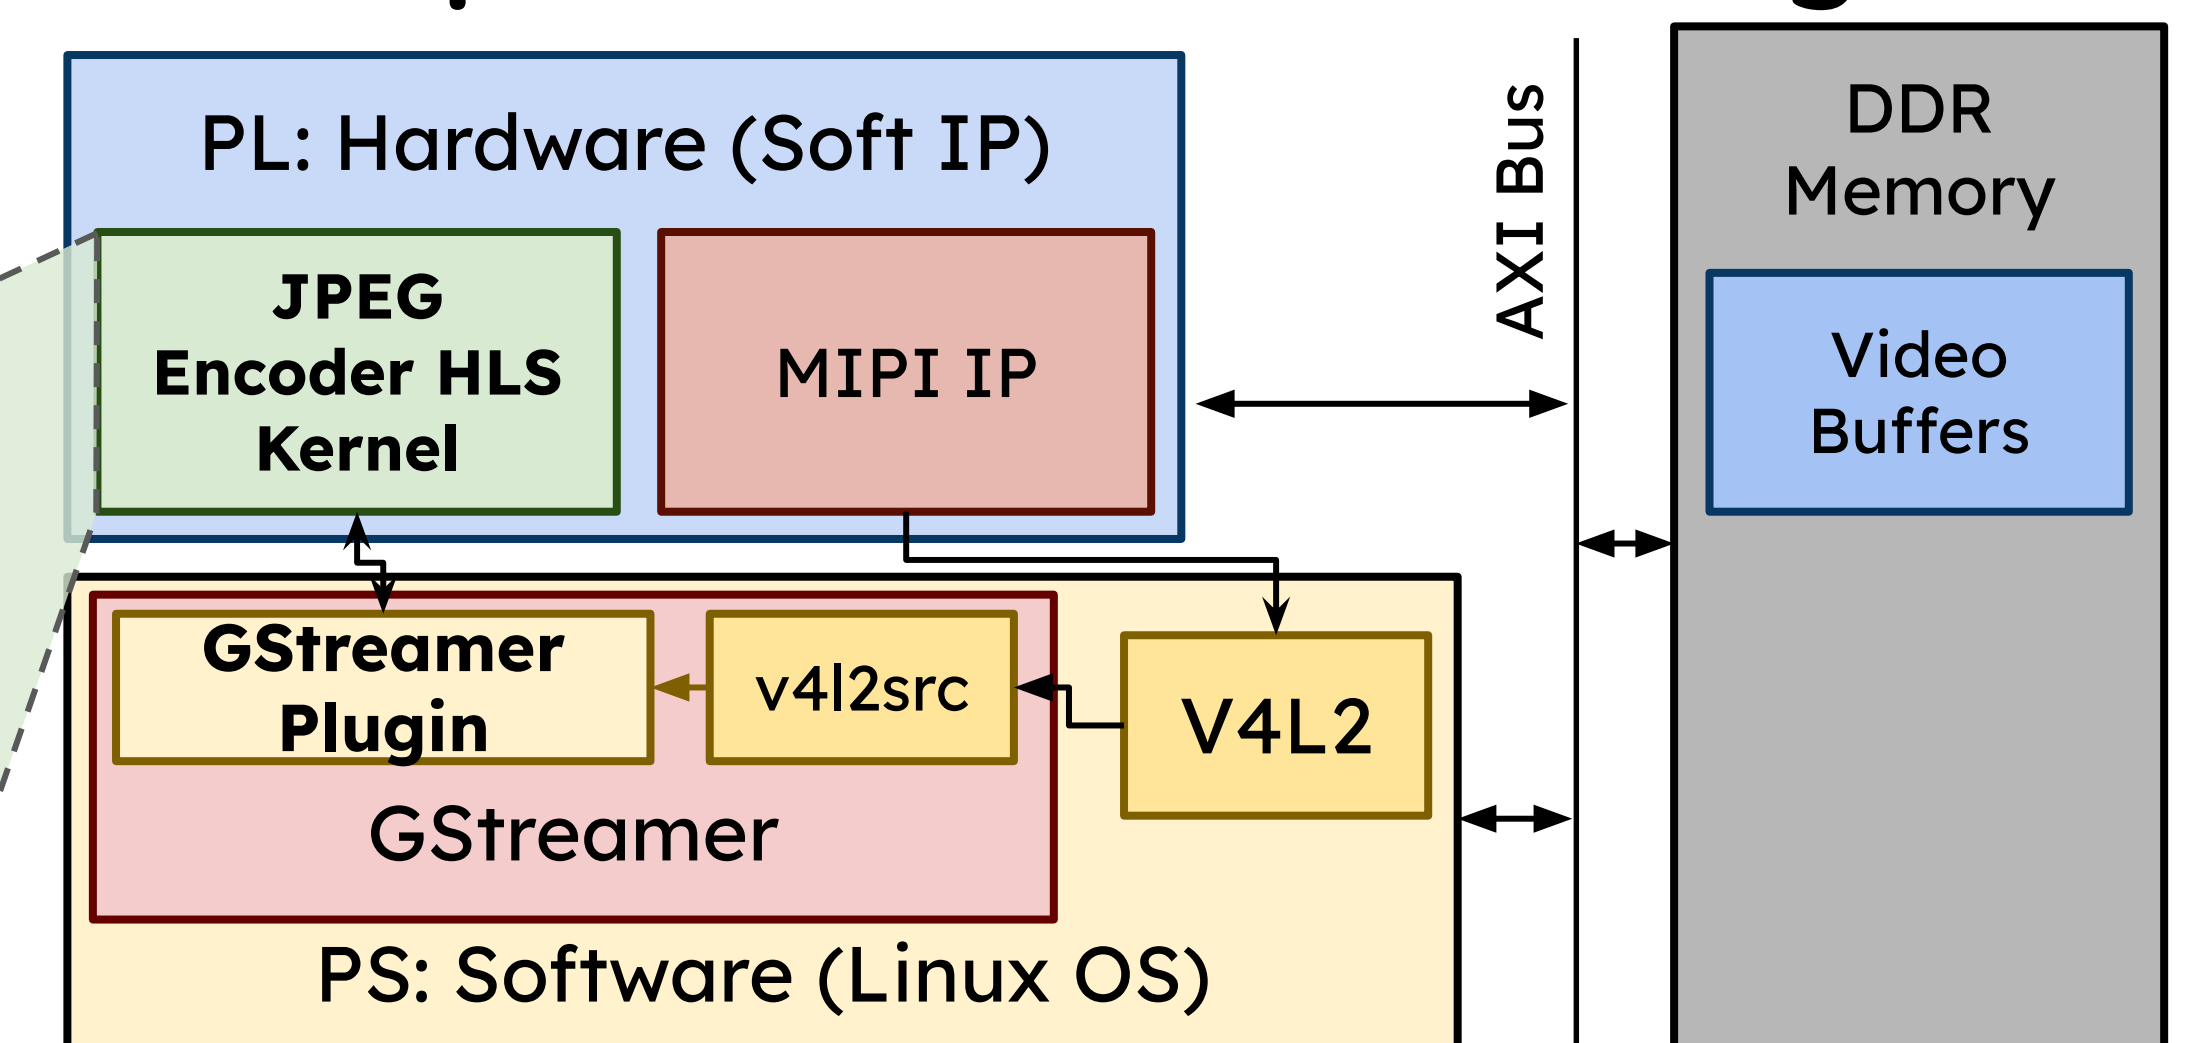
Contributions

- **HW JPEG Encoder** implemented in *High-level Synthesis*.
- **GStreamer plugin** interfaces with encoder.
 - Allows for mix and matching of pipeline elements for arbitrary media processing.
 - Easy usage of DMA capabilities.

Hardware Architecture



Top-level Hardware Design

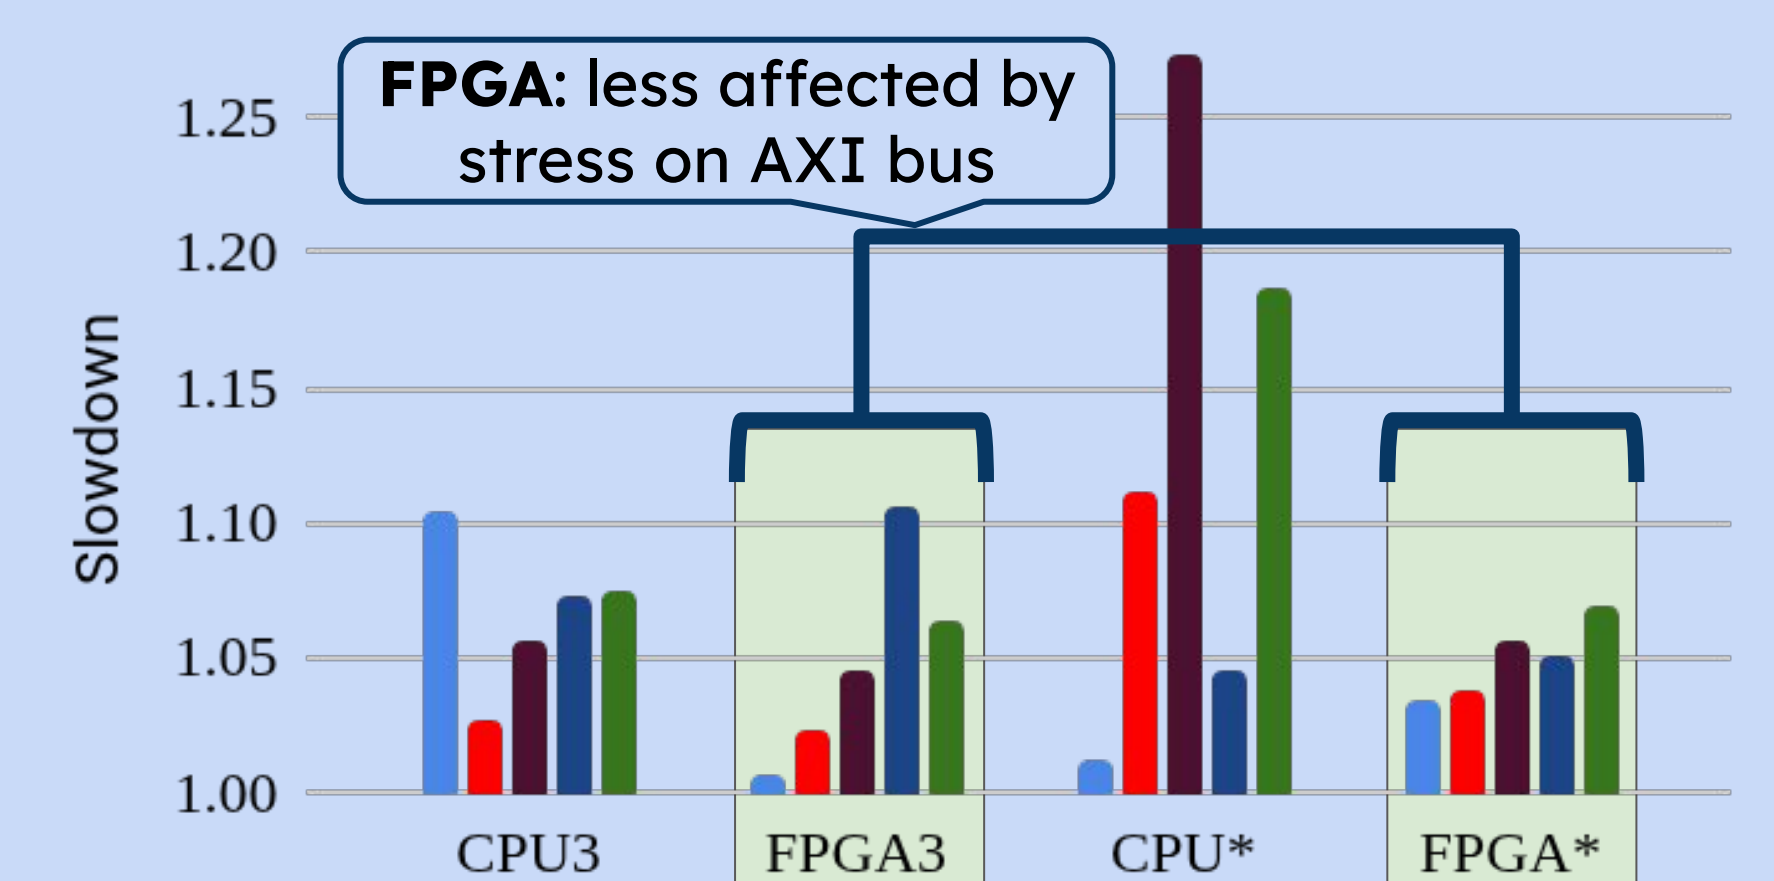
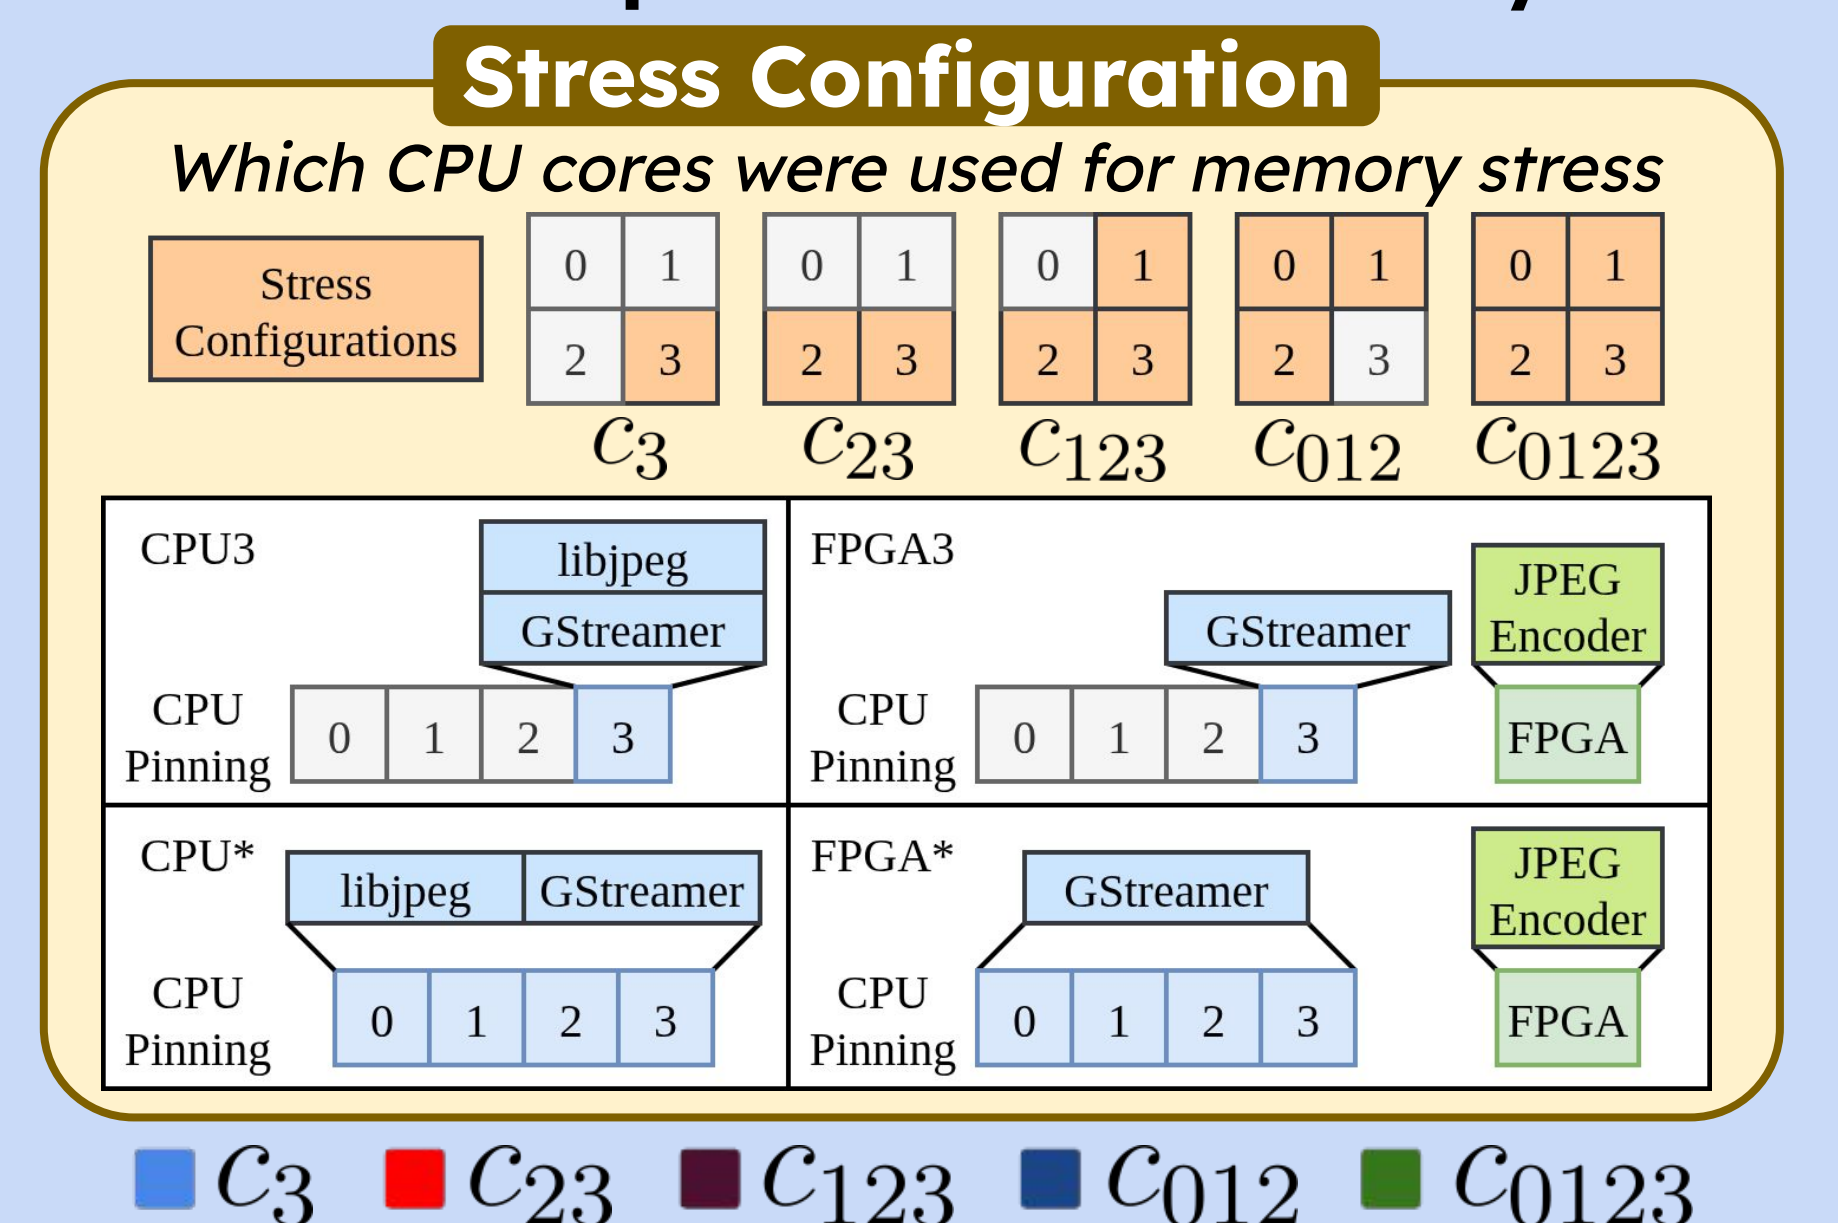


- Camera accessed through V4L2
 - Accesses must go through AXI bus

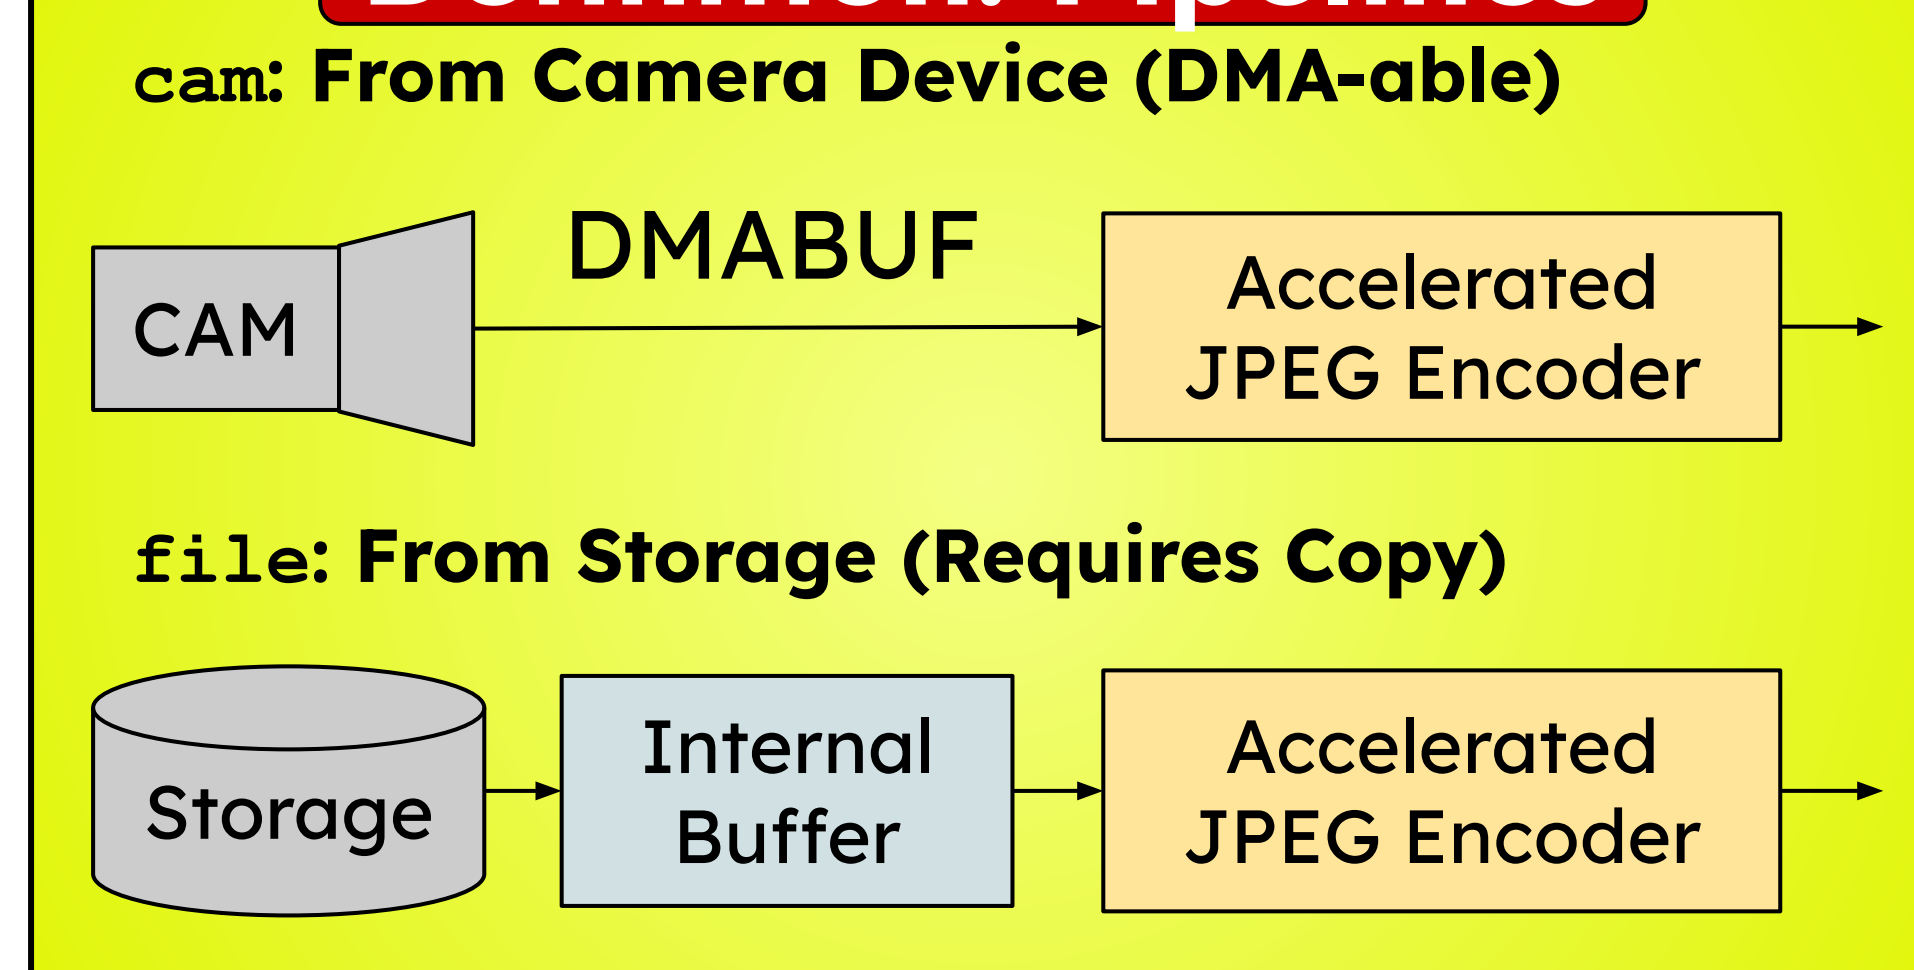
Experiments

Memory Stress Testing

Slowdown Comparison under Memory Stress



Definition: Pipelines



Performance

FPGA vs. CPU (jpegenc)

		CPU	FPGA		FPGA/CPU	
			Single Core	2 CU	1 CU	2 CU
cam	FPS	2.6	19.5	30.0	7.5	15
	CPU %	93	5	10	0.054	0.108
	FPS	22.6	21.7	30.0	.96	1.38
	CPU %	107	14	28	0.13	0.26

Resource Utilization

JPEG (our) vs. Vitis Codec Libraries

	Freq.	BRAM	URAM	DSP	FF	LUT
JPEG ¹	250MHz	6	12	85	18150	13868
Webp ²	100MHz	81	46	834	71227	68906
Webp ²	250MHz	229	10	414	92030	68755
JXL ²	260MHz	584	122	644	270565	199537
PIK ²	200MHz	614	473	2398	549987	449995

¹Our work

²Vitis Codec Libraries

Webp Encoder²

- **High Resource Usage**
 - Requires lowering target frequency to 100Hz

JPEG Encoder¹

- **7~25% of Webp Resources**
 - Can fit **four** JPEG encoders in place of a **single** Webp Encoder

Conclusions

1. Streamlined flow for GStreamer integration with arbitrary HLS kernels
2. Resilience against memory stress
3. Increased Predictability

Future Work

- Streamlined flow for GStreamer integration with arbitrary HLS kernels

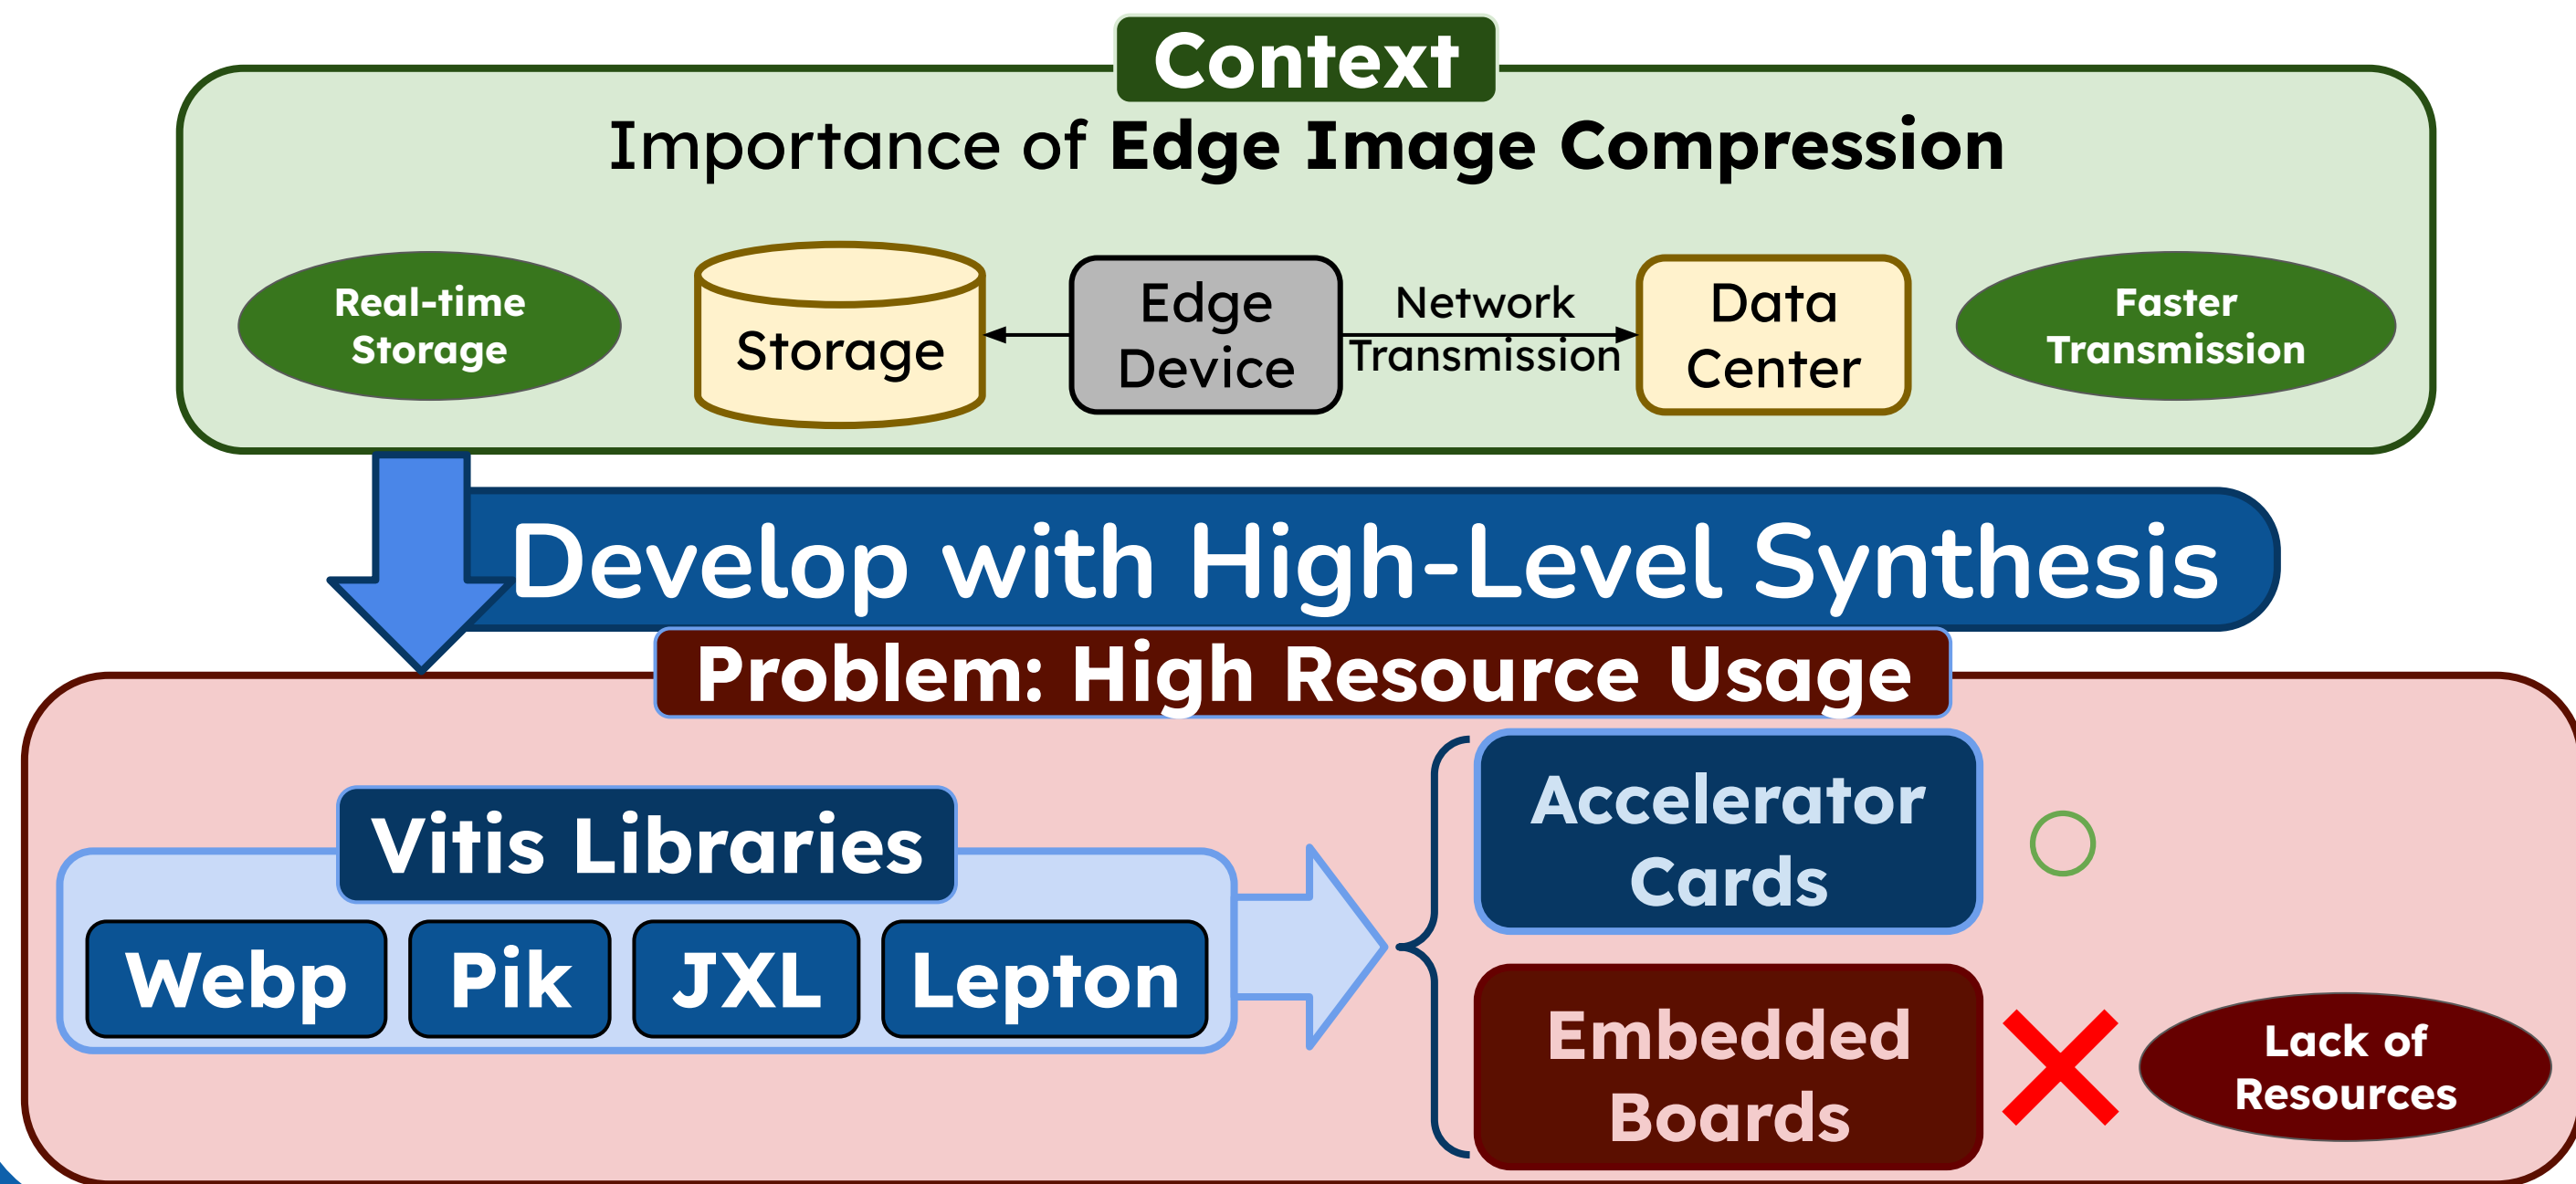
GStreamer-integrated HLS-based JPEG Encoder for Edge FPGA SoCs

Yuri Guimaraes Pereira Primo da Silva¹, Shinya Honda¹,
Sugako Otani¹², Masato Eda¹, Abraham Monrroy Cano³

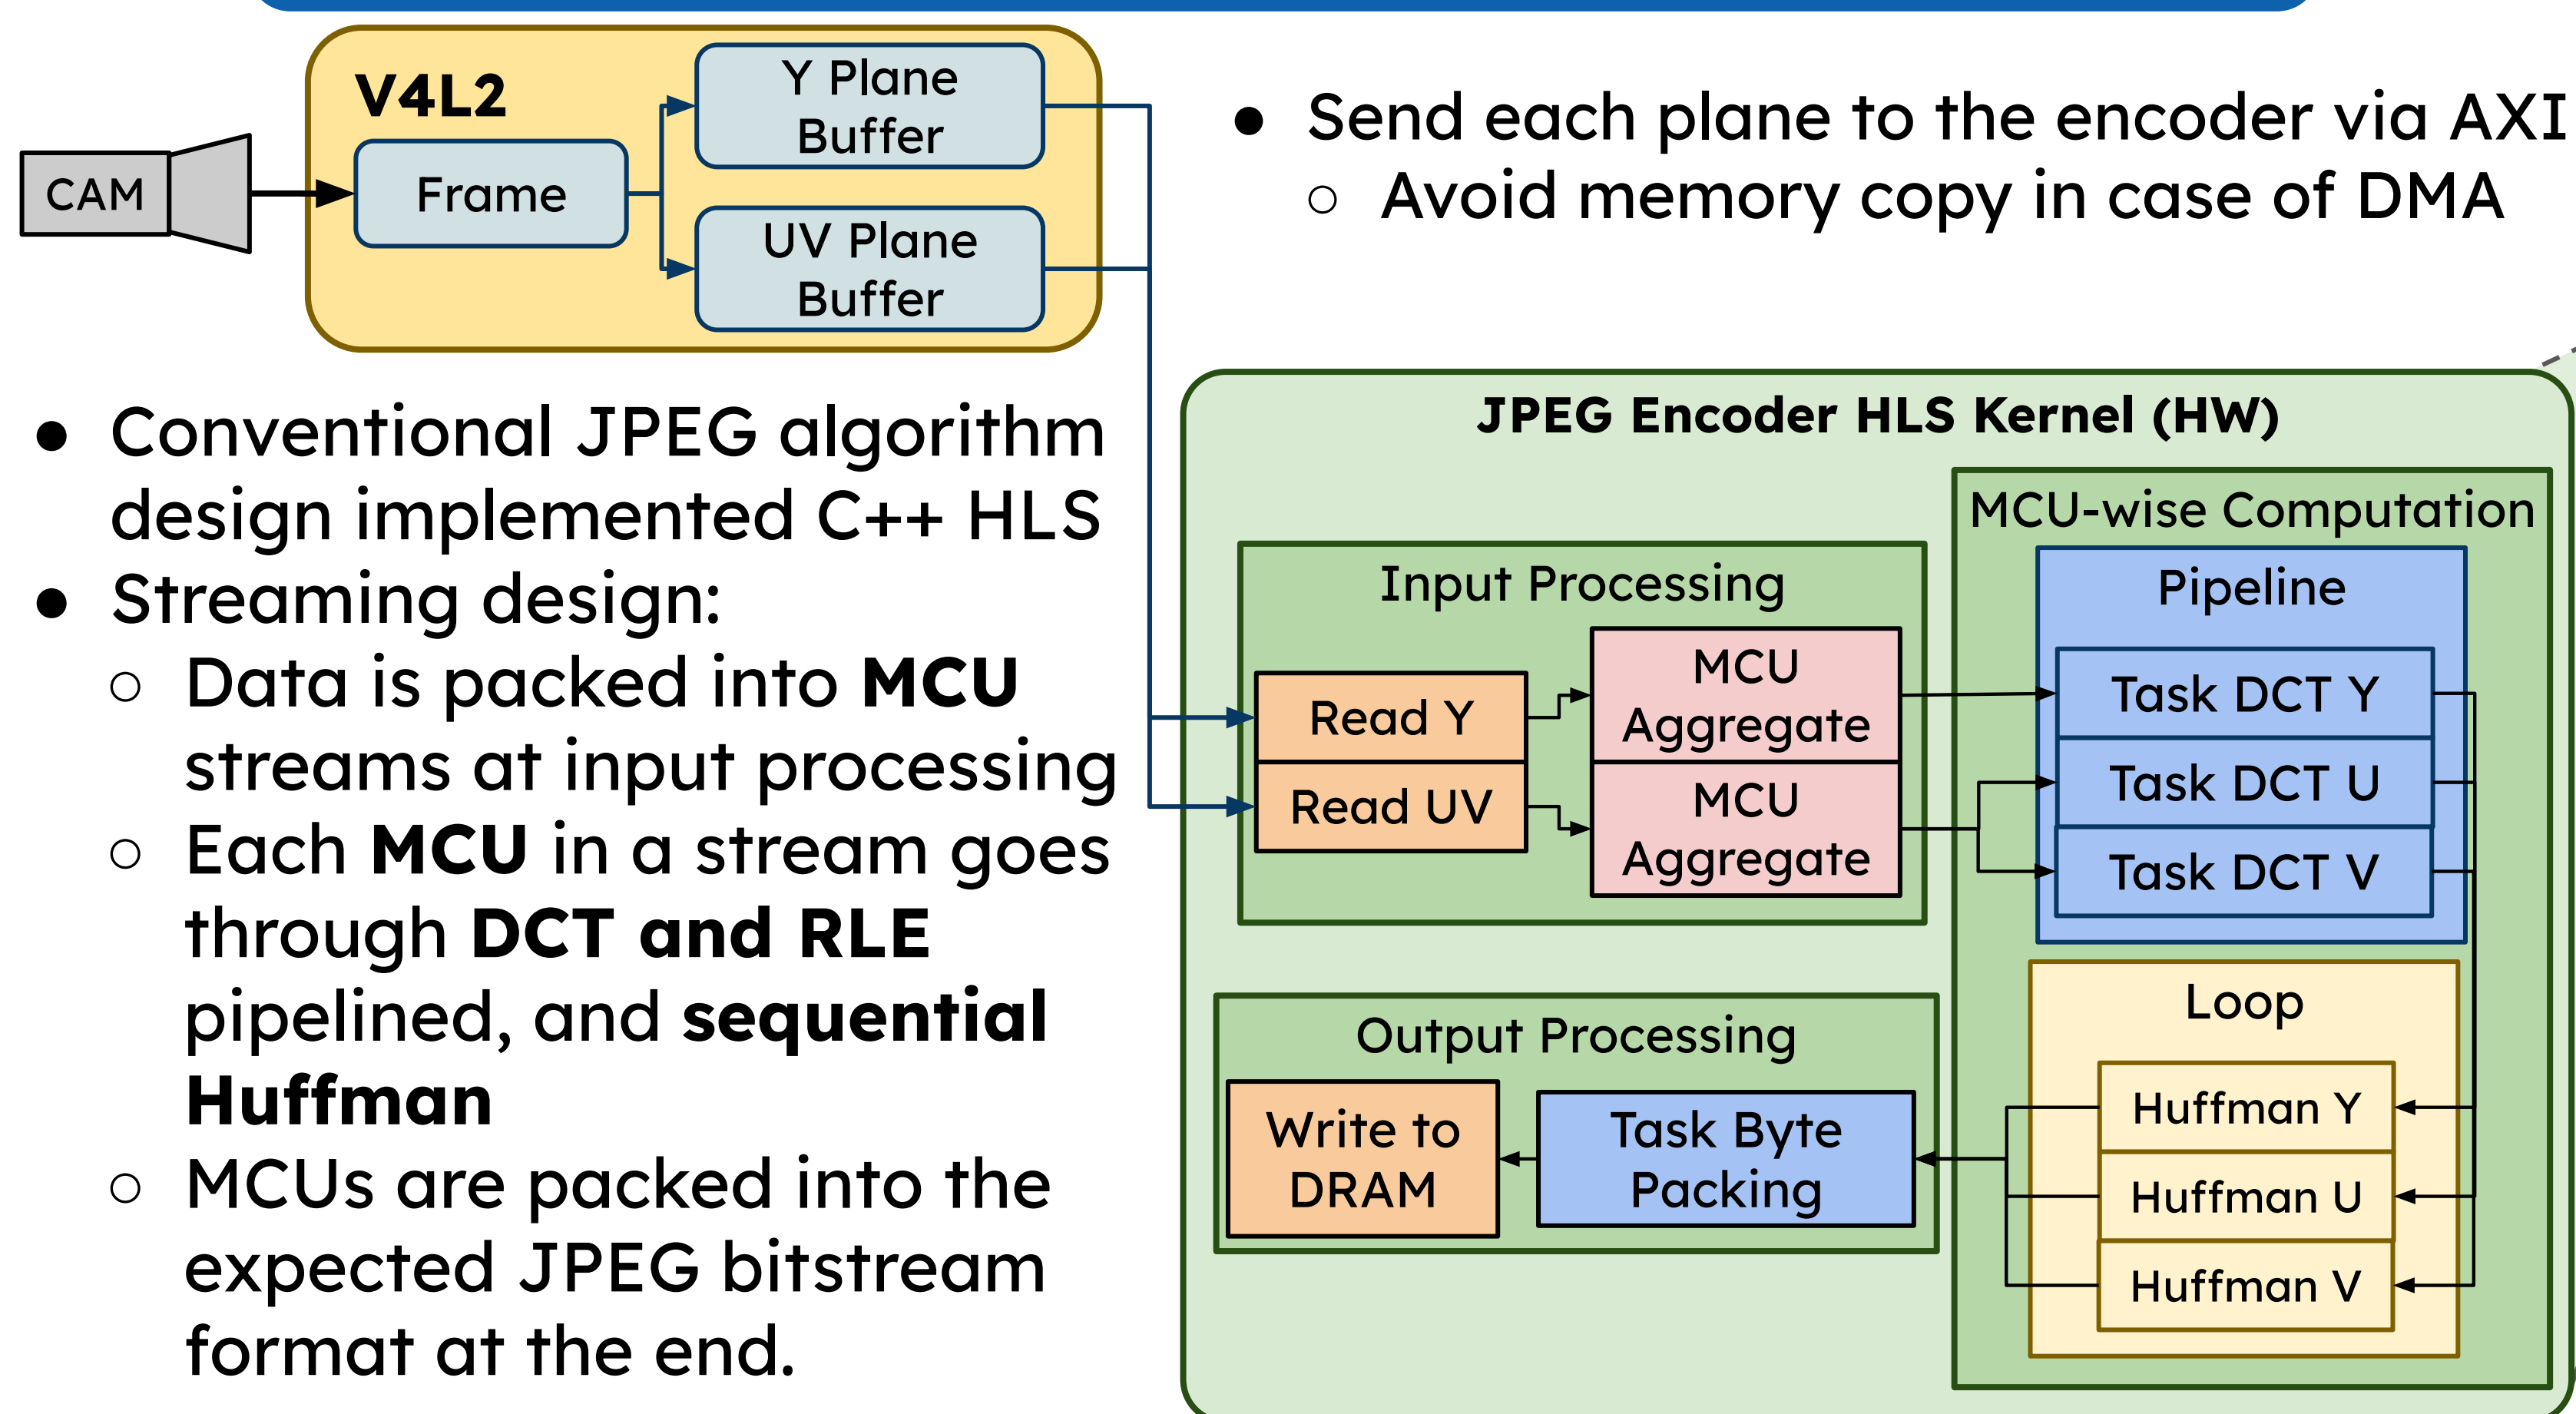
¹Nagoya University, ²Renesas Electronics,
³Map IV Inc.

E-mail: yuri_gpps1@ert1.jp

Motivation

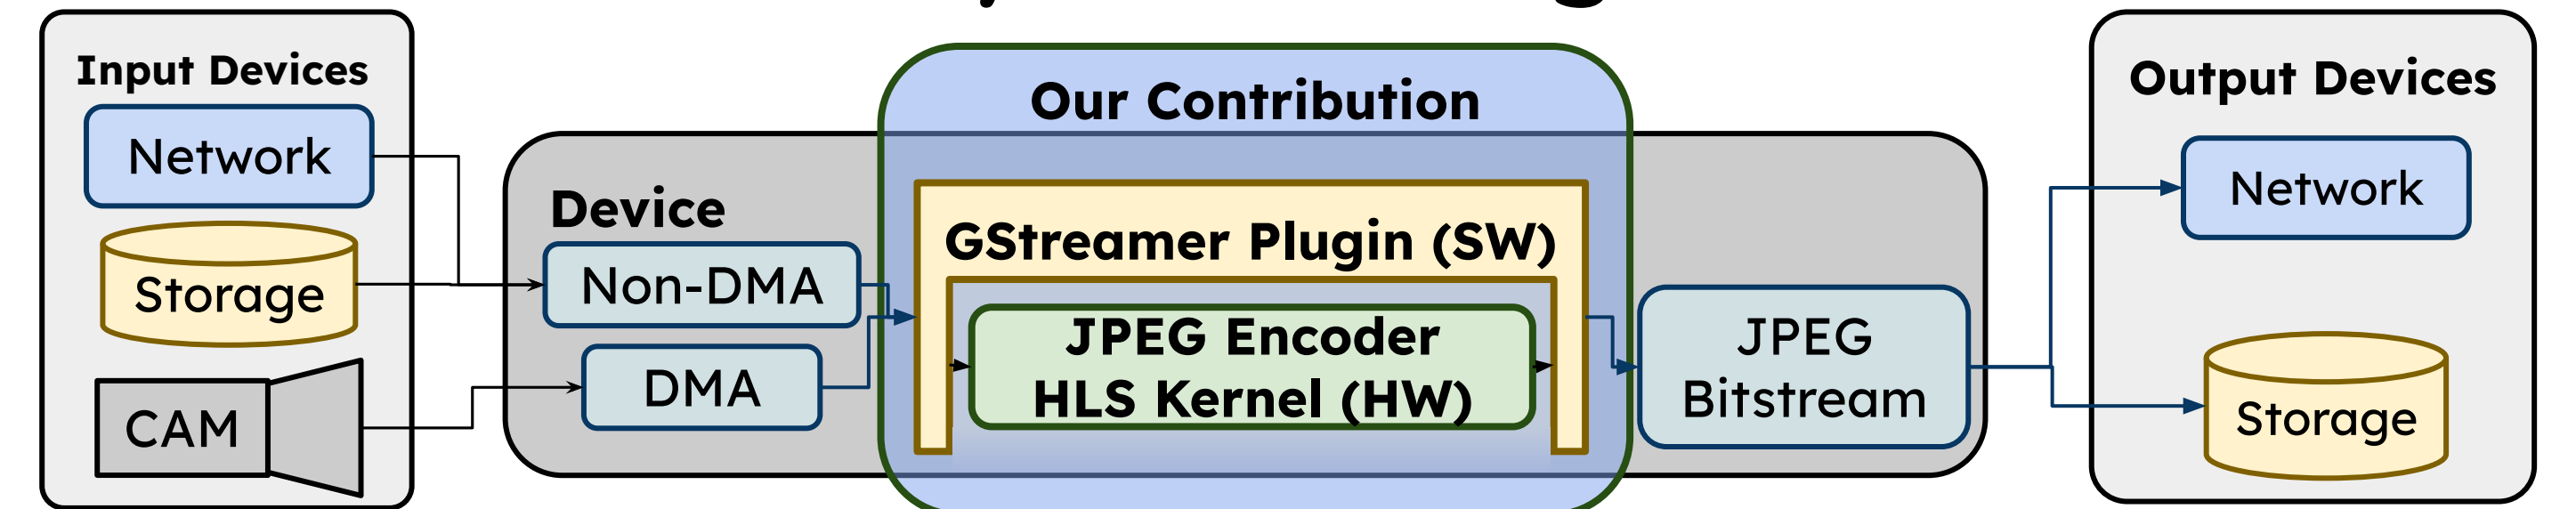


Hardware Architecture



Approach

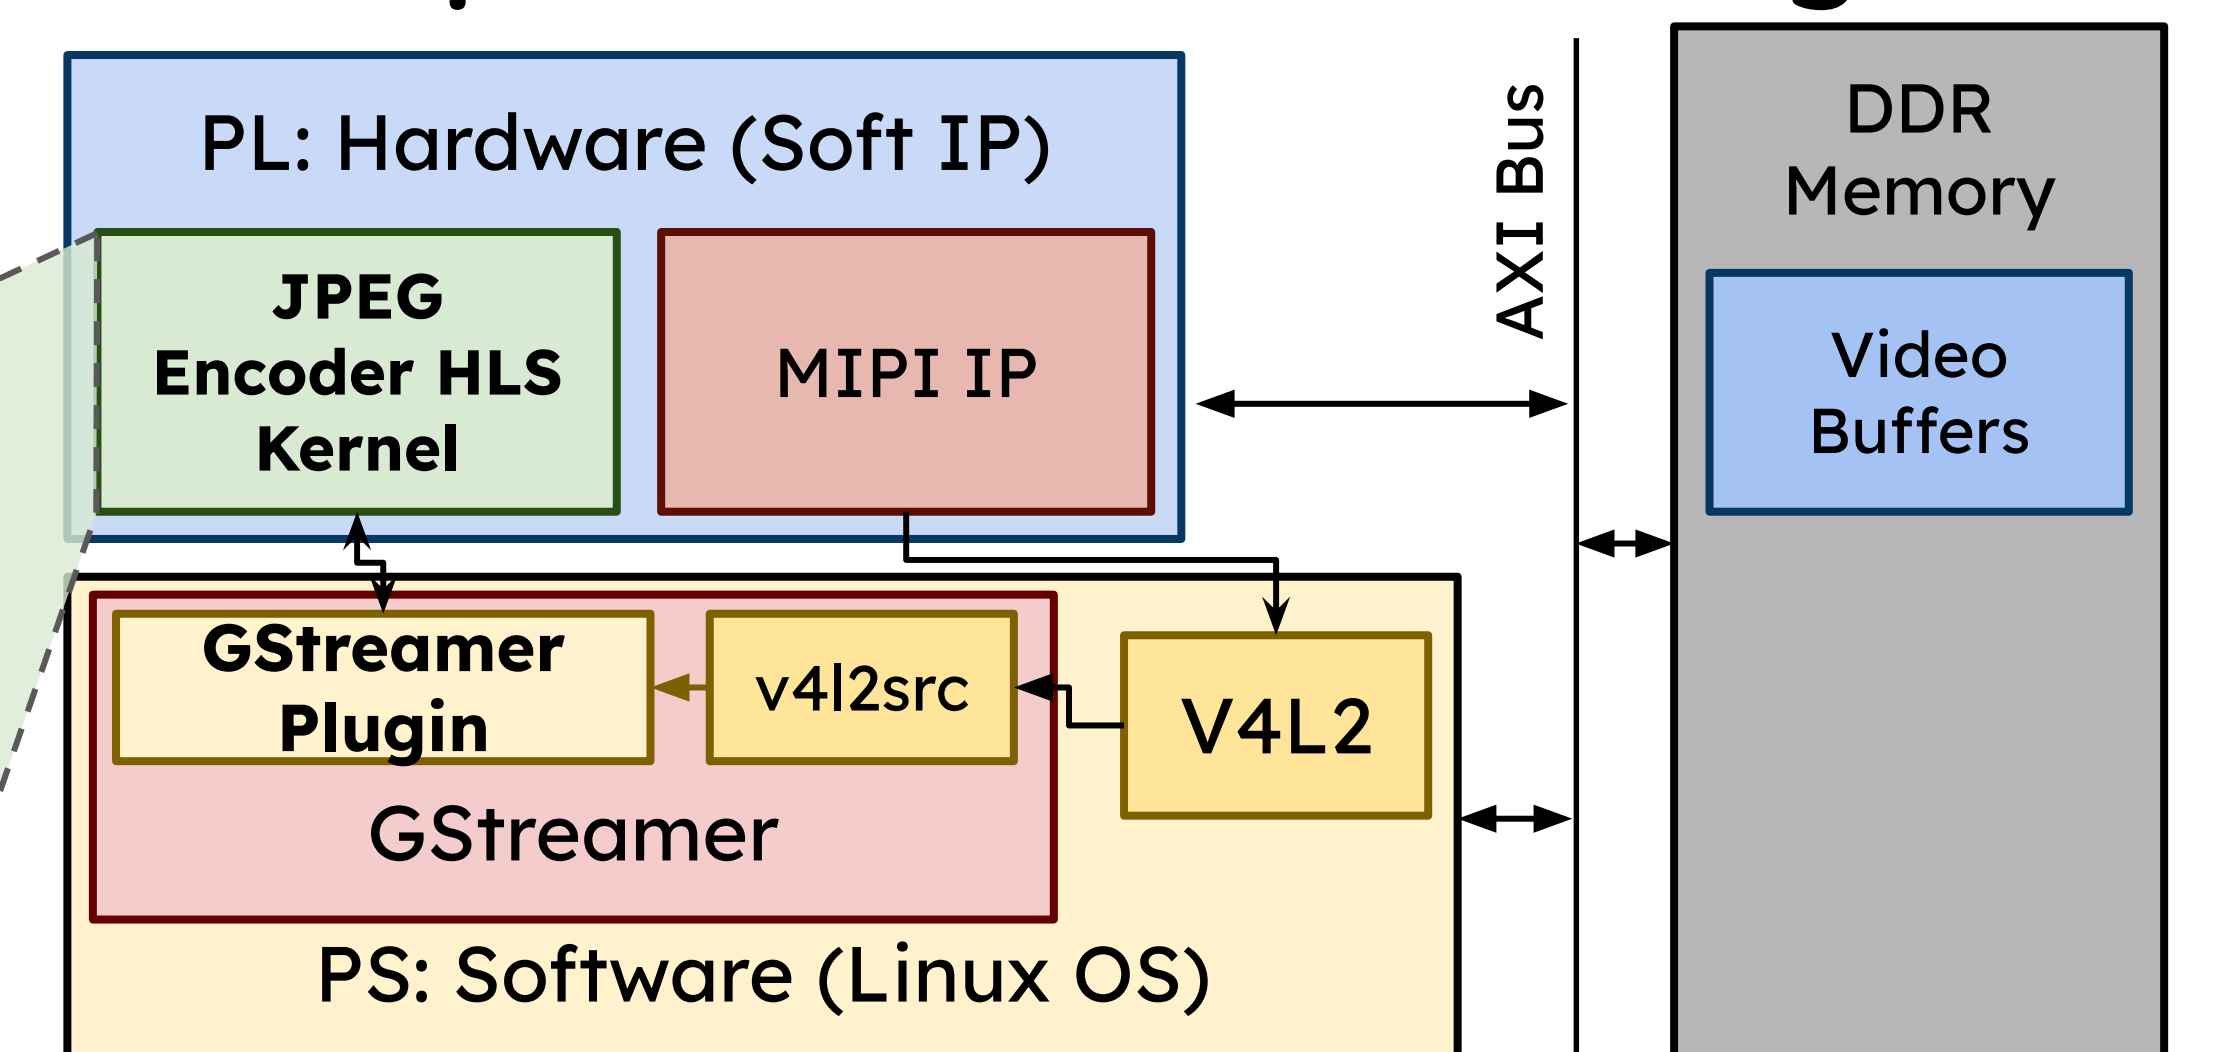
System Design



Contributions

- HW JPEG Encoder** implemented in *High-level Synthesis*.
- GStreamer plugin** interfaces with encoder.
 - Allows for mix and matching of pipeline elements for arbitrary media processing.
 - Easy usage of DMA capabilities.

Top-level Hardware Design

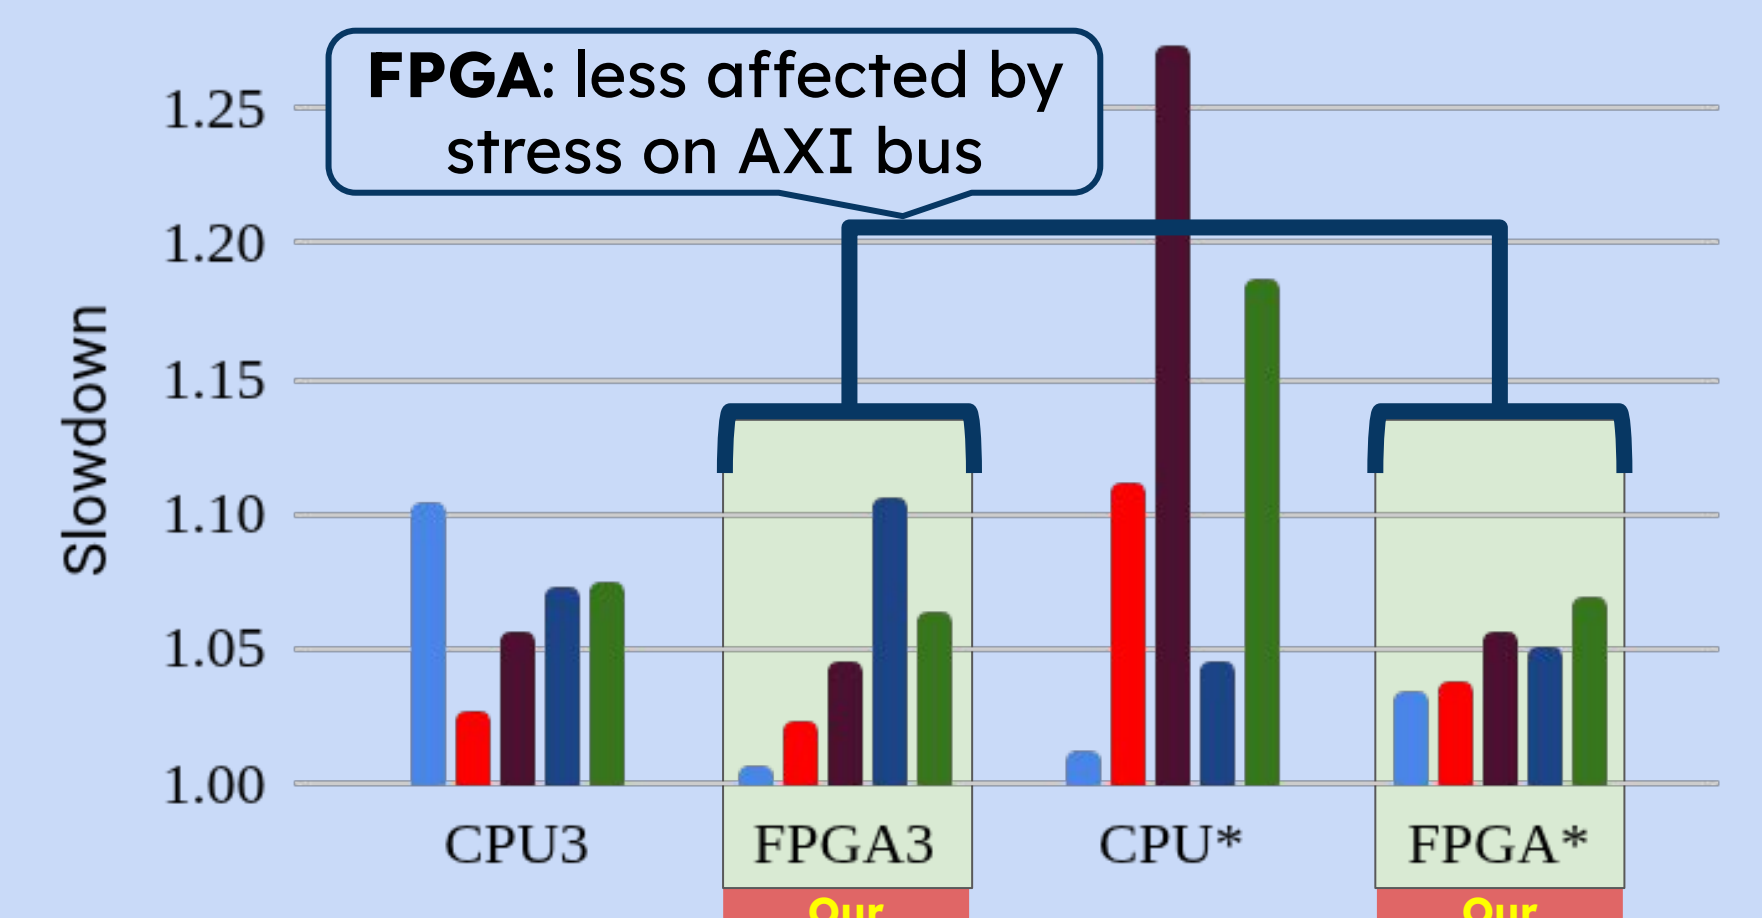
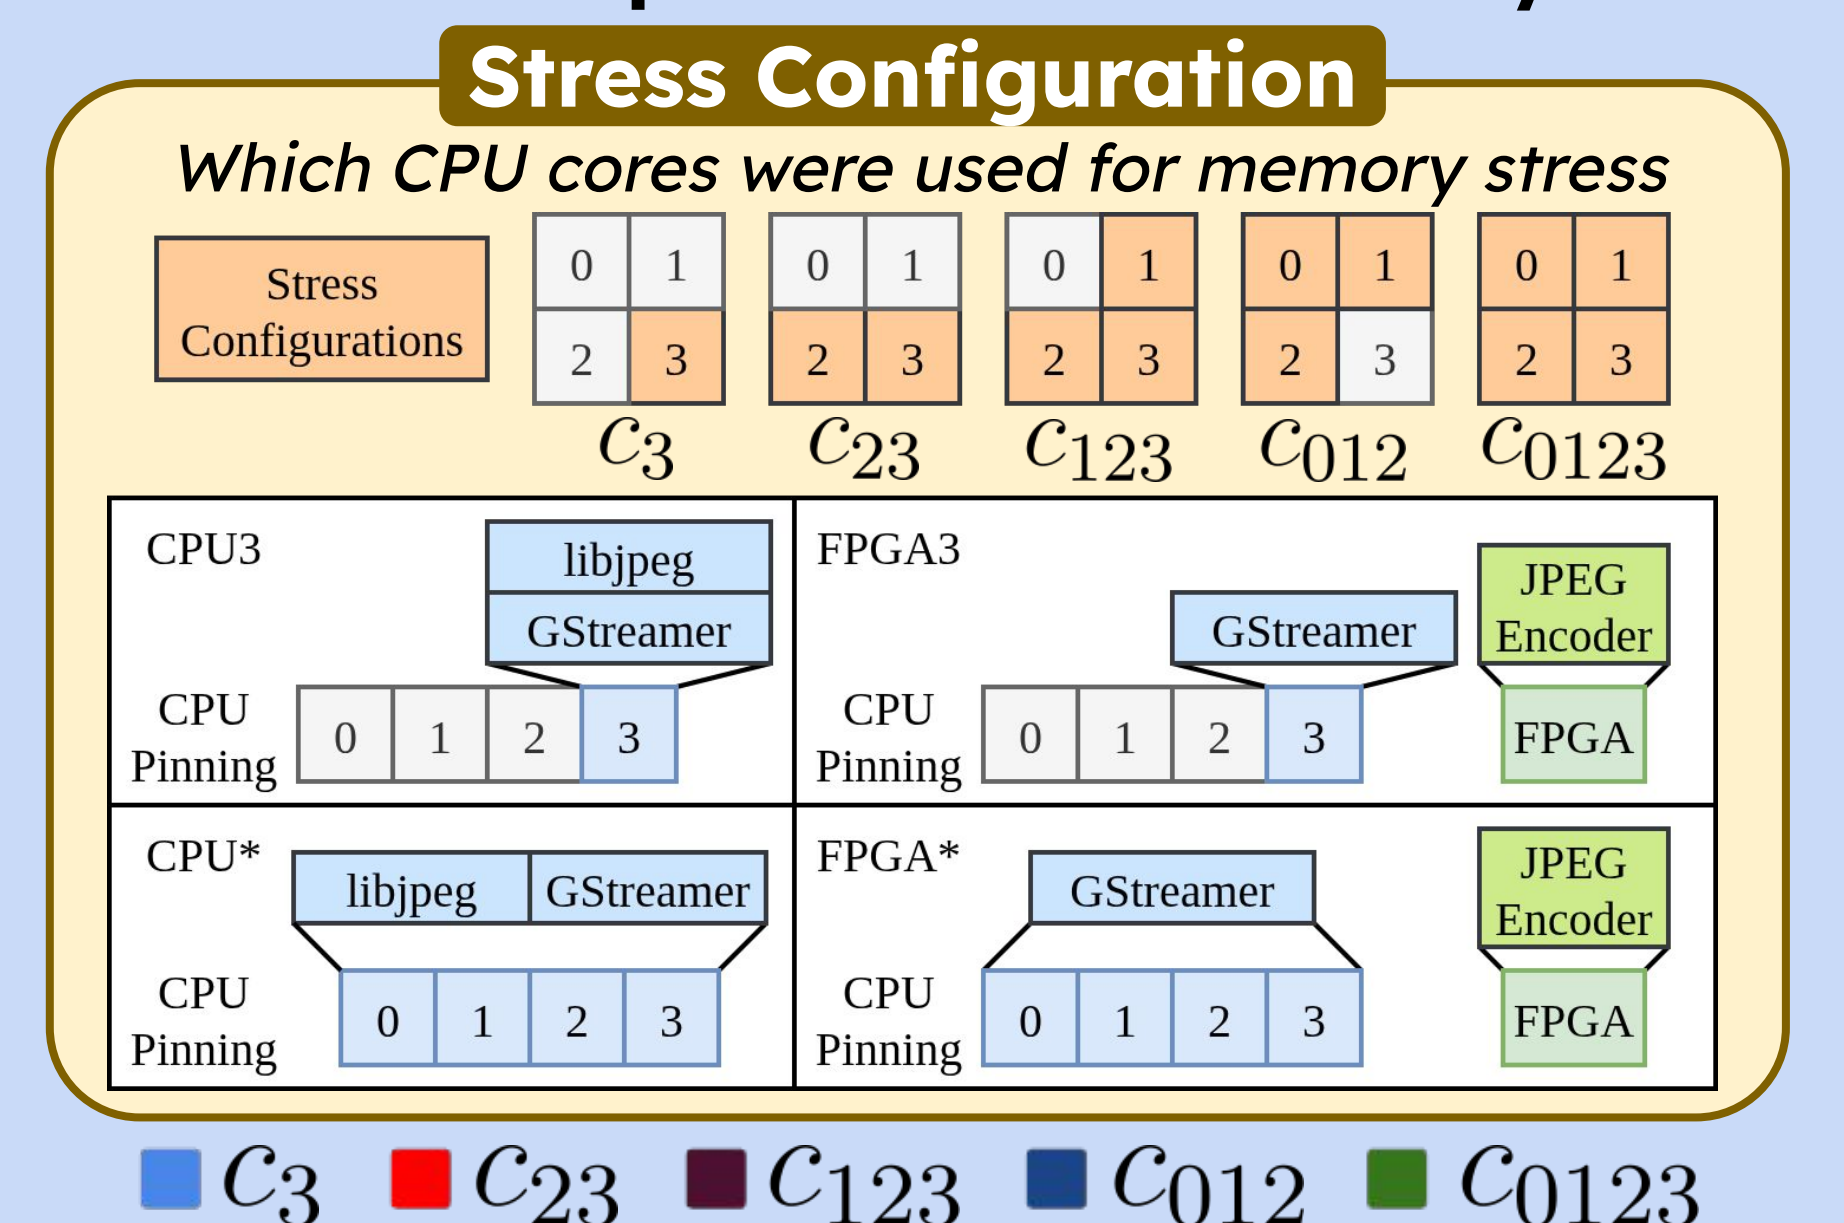


- Camera accessed through V4L2
 - Accesses must go through AXI bus

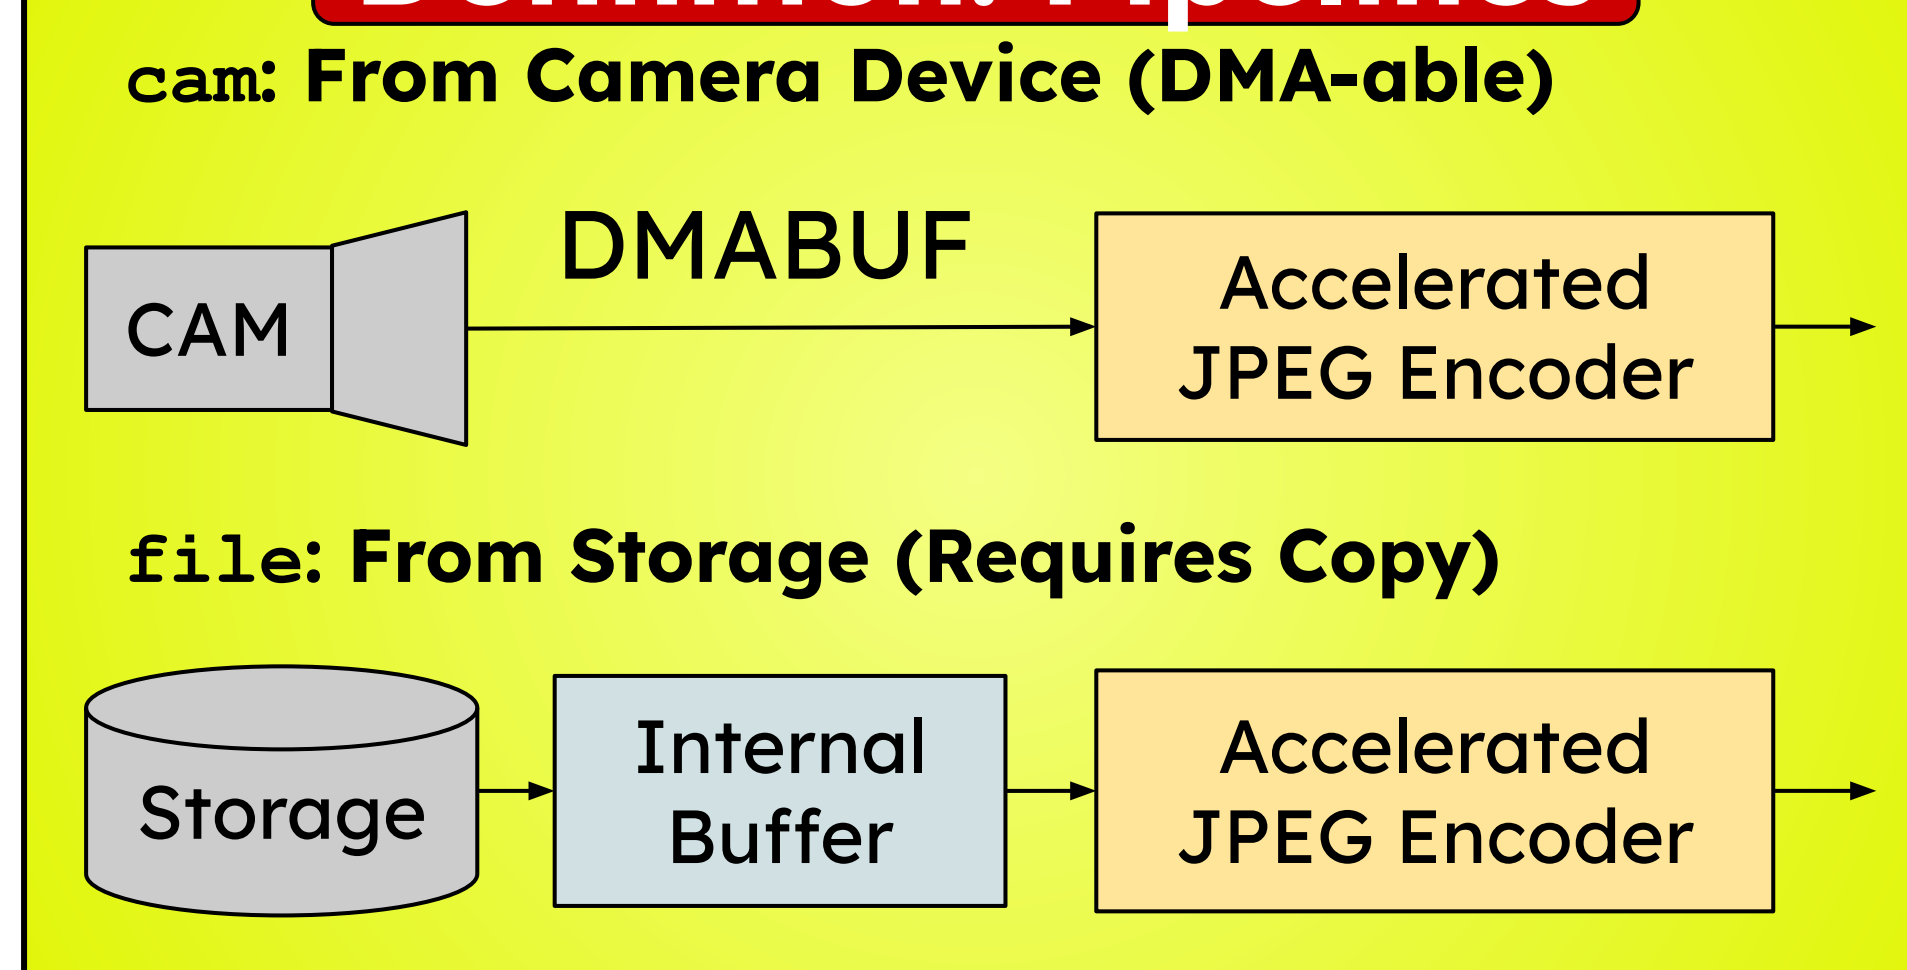
Experiments

Memory Stress Testing

Slowdown Comparison under Memory Stress



Definition: Pipelines



Performance

FPGA vs. CPU (jpegenc)

		CPU	Our FPGA	FPGA/CPU
cam	FPS	2.6	19.5	7.5
	CPU %	93	5	0.054
file	FPS	22.6	21.7	.96
	CPU %	107	14	0.13

Resource Utilization

JPEG (our) vs. Vitis Codec Libraries

	Freq.	BRAM	URAM	DSP	FF	LUT
Our JPEG ¹	250MHz	6	12	85	18150	13868
Webp ²	100MHz	81	46	834	71227	68906
Webp ²	250MHz	229	10	414	92030	68755
JXL ²	260MHz	584	122	644	270565	199537
PIK ²	200MHz	614	473	2398	549987	449995

¹Our work

²Vitis Codec Libraries

Webp Encoder²

- High Resource Usage**
 - Requires lowering target frequency to 100Hz

JPEG Encoder¹

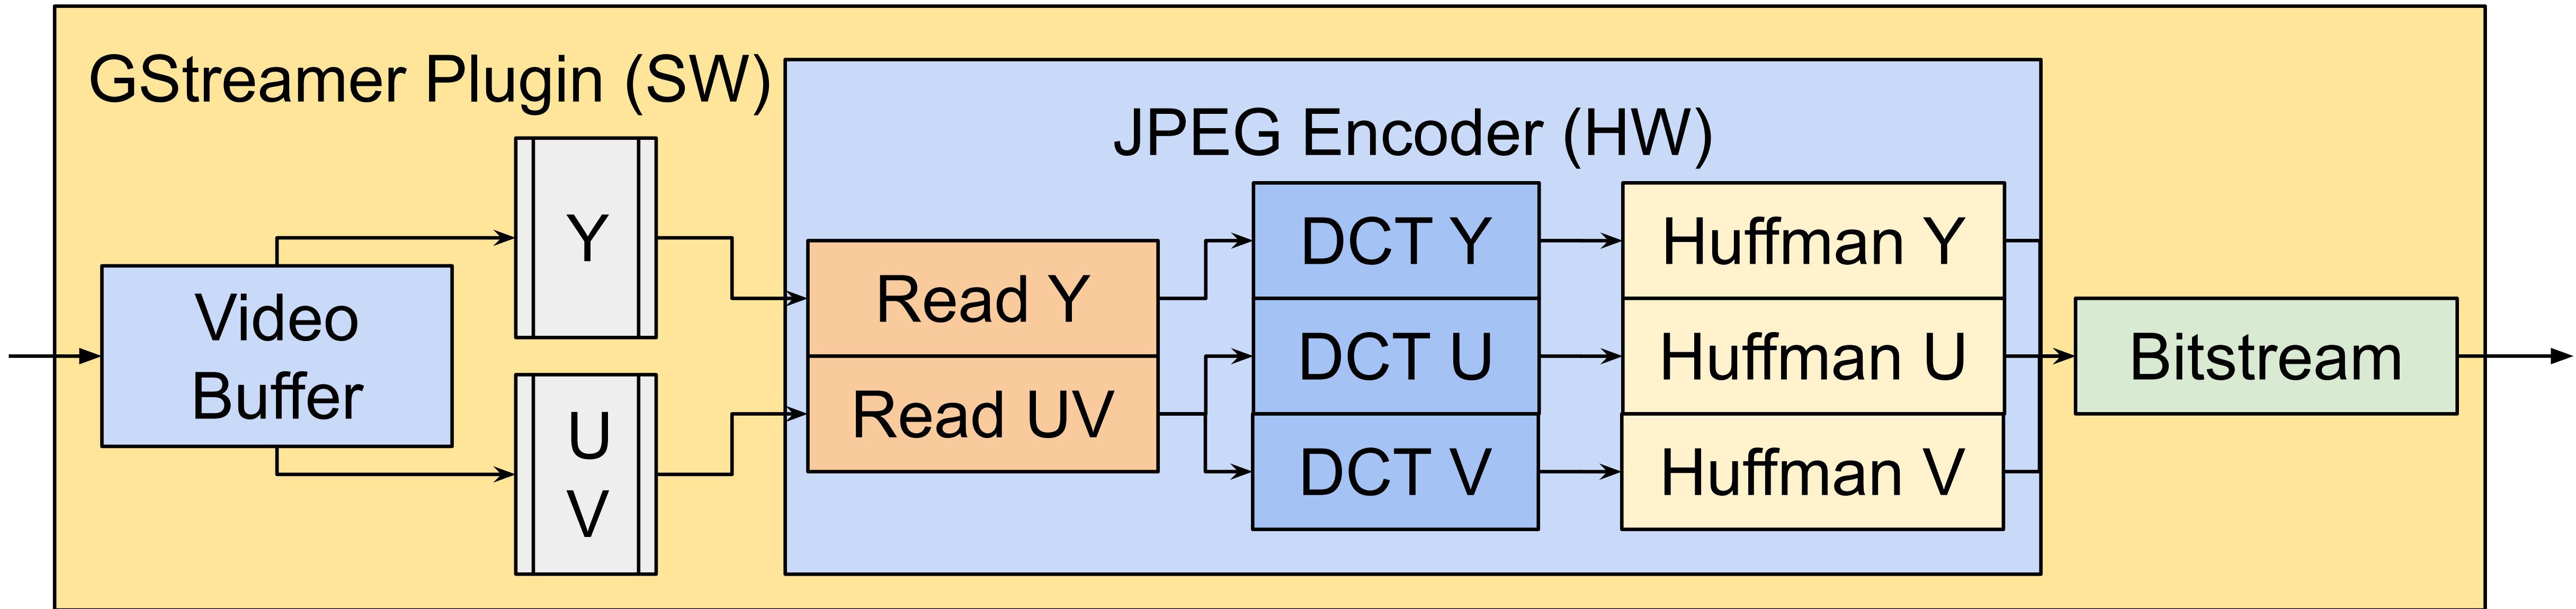
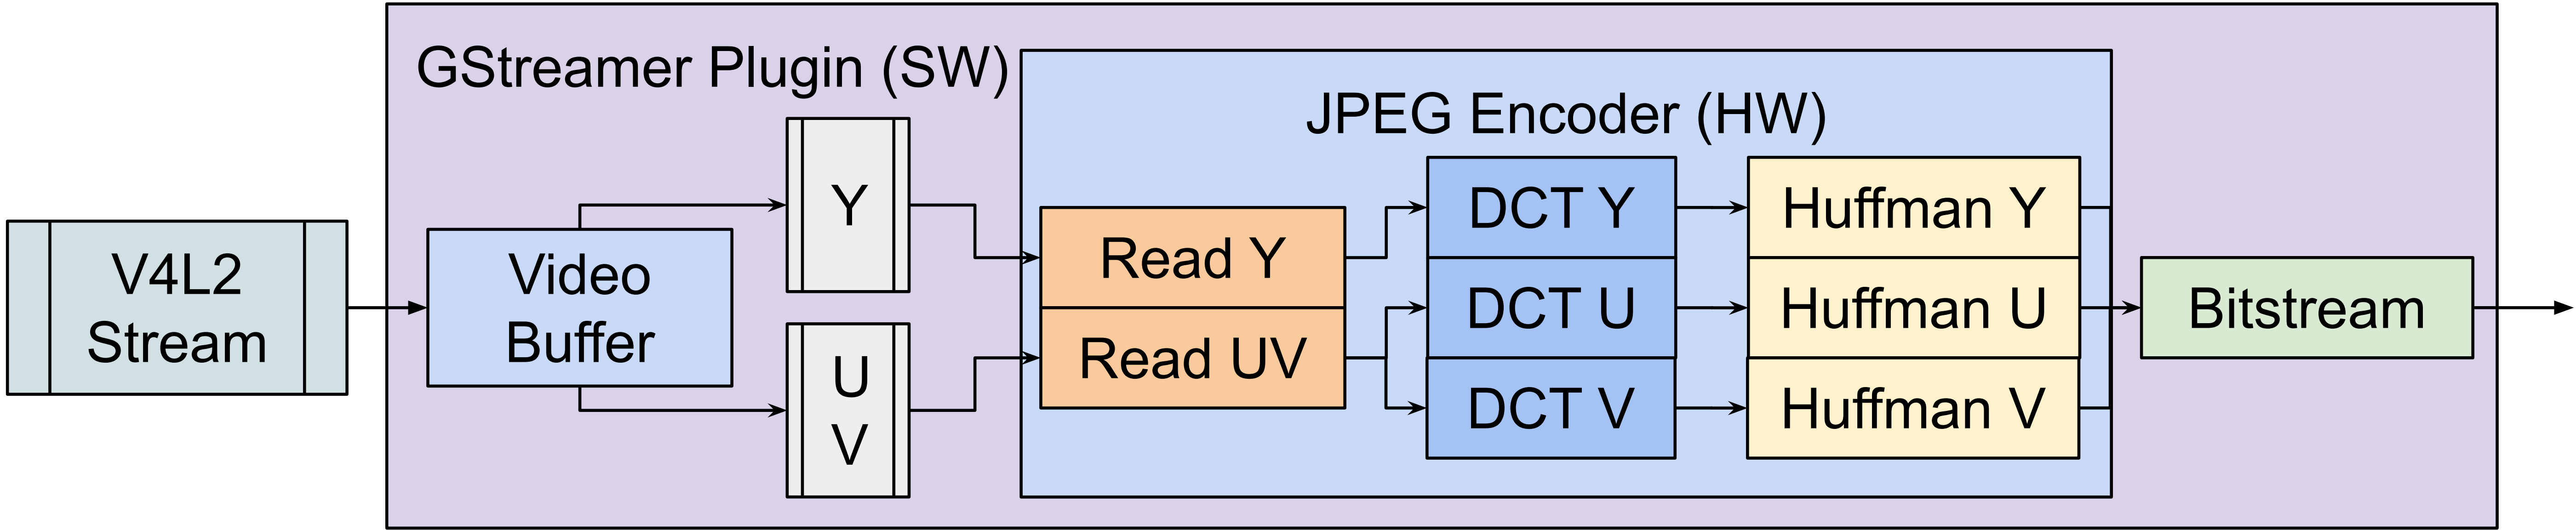
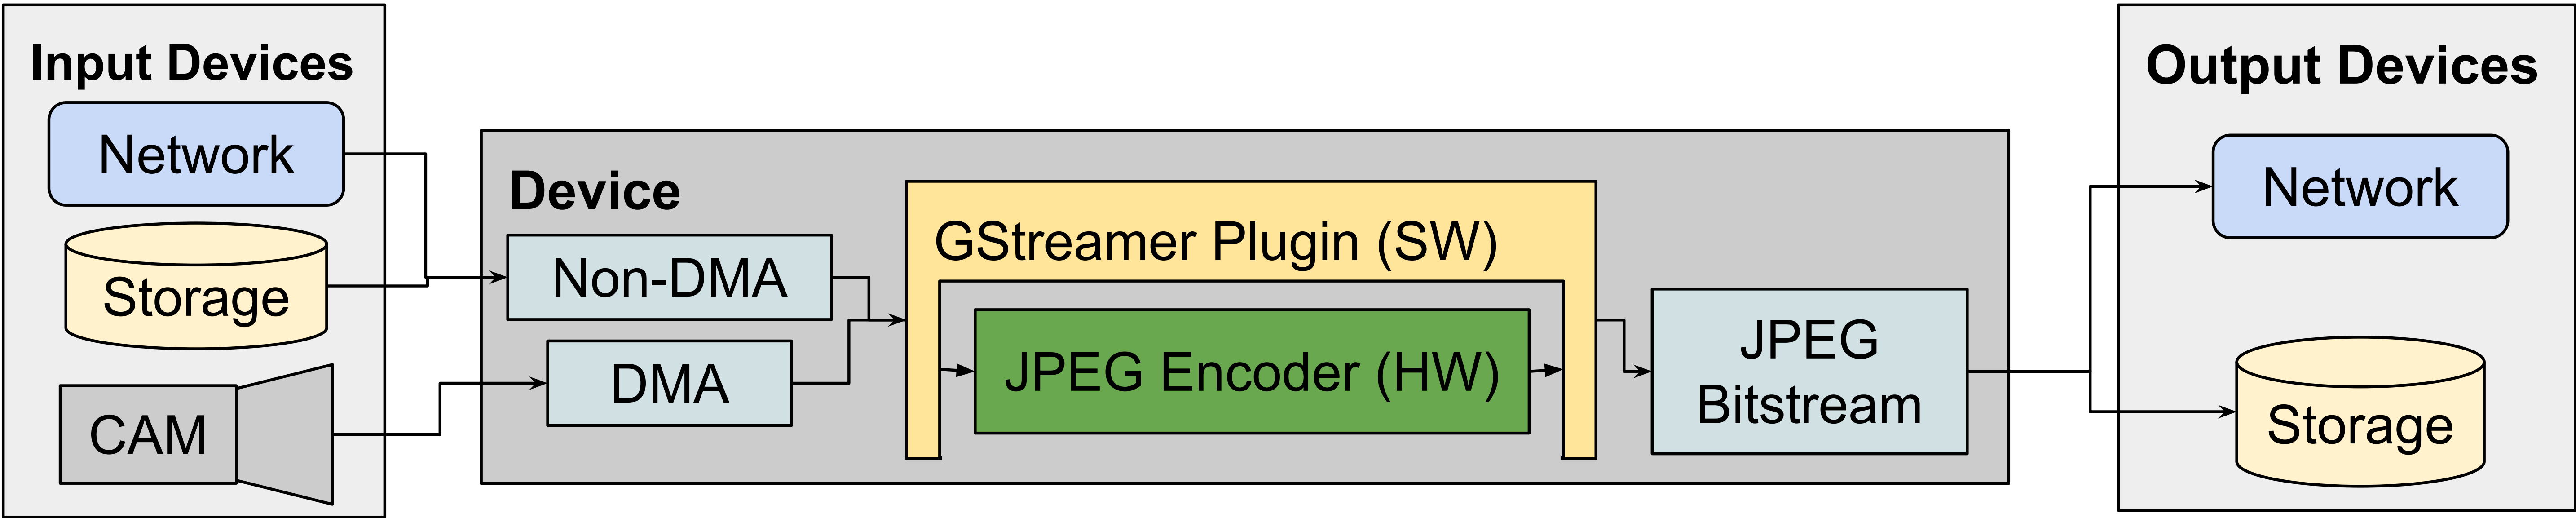
- 7~25% of Webp Resources**
 - Can fit **four** JPEG encoders in place of a **single** Webp Encoder

Discussion and Conclusions

- Suitability for Edge
 - Low resource usage
 - Realtime performance
- Usability
 - GStreamer integration
 - Easy composition of media pipelines
- Increased Predictability
 - Resilience against memory stress

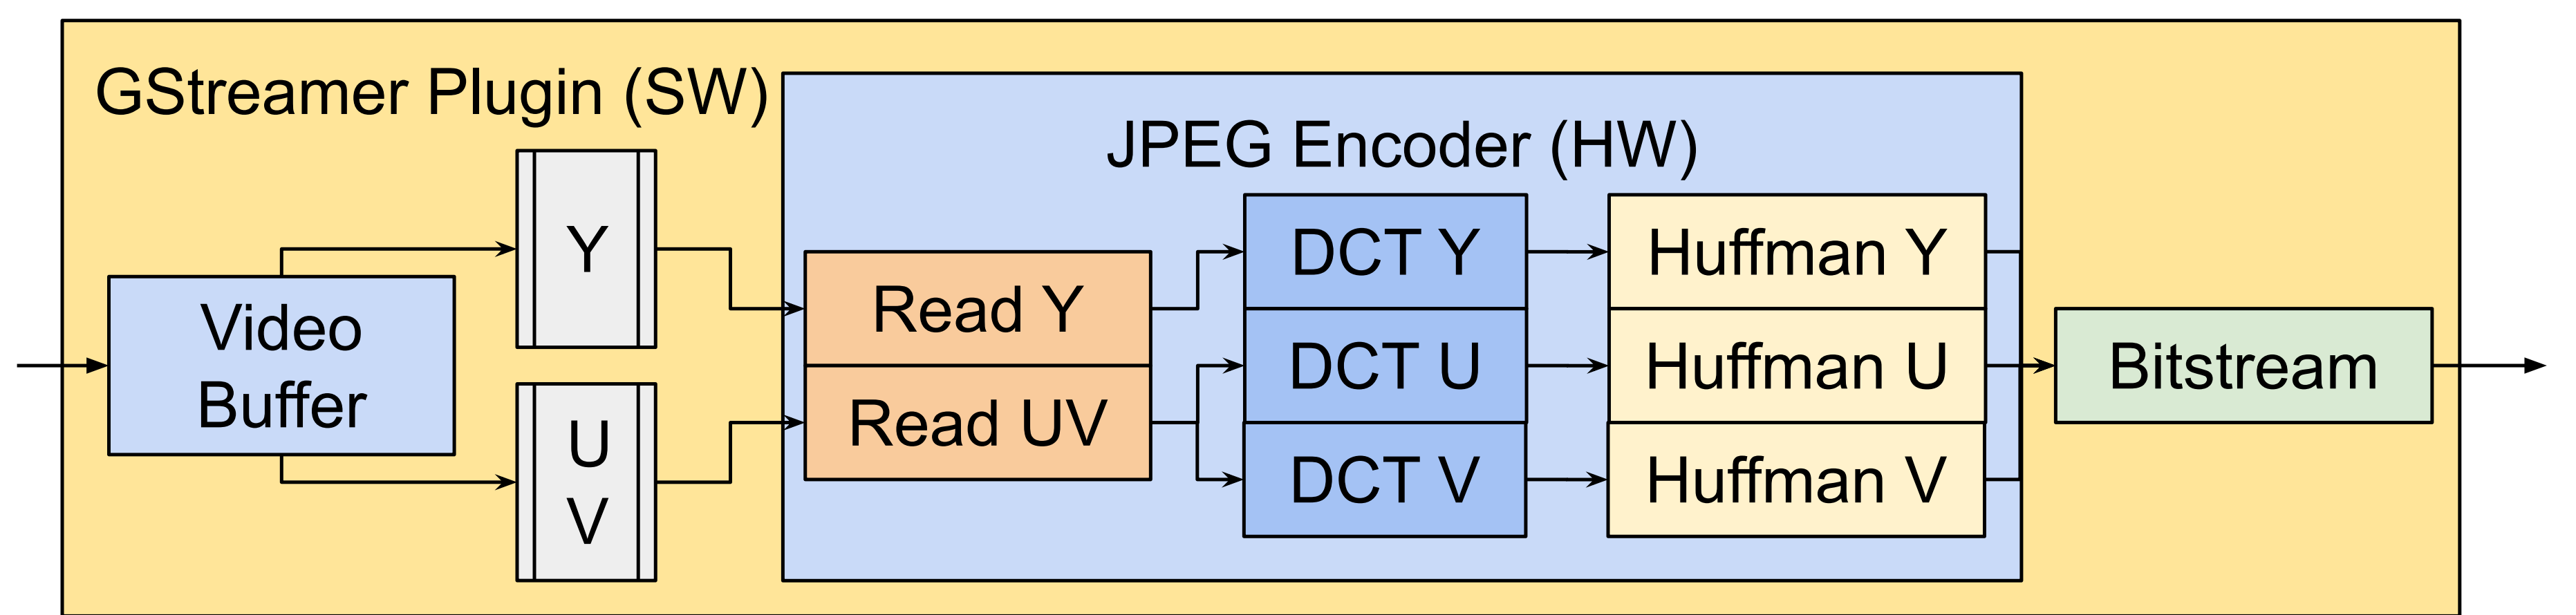
Download this Poster!



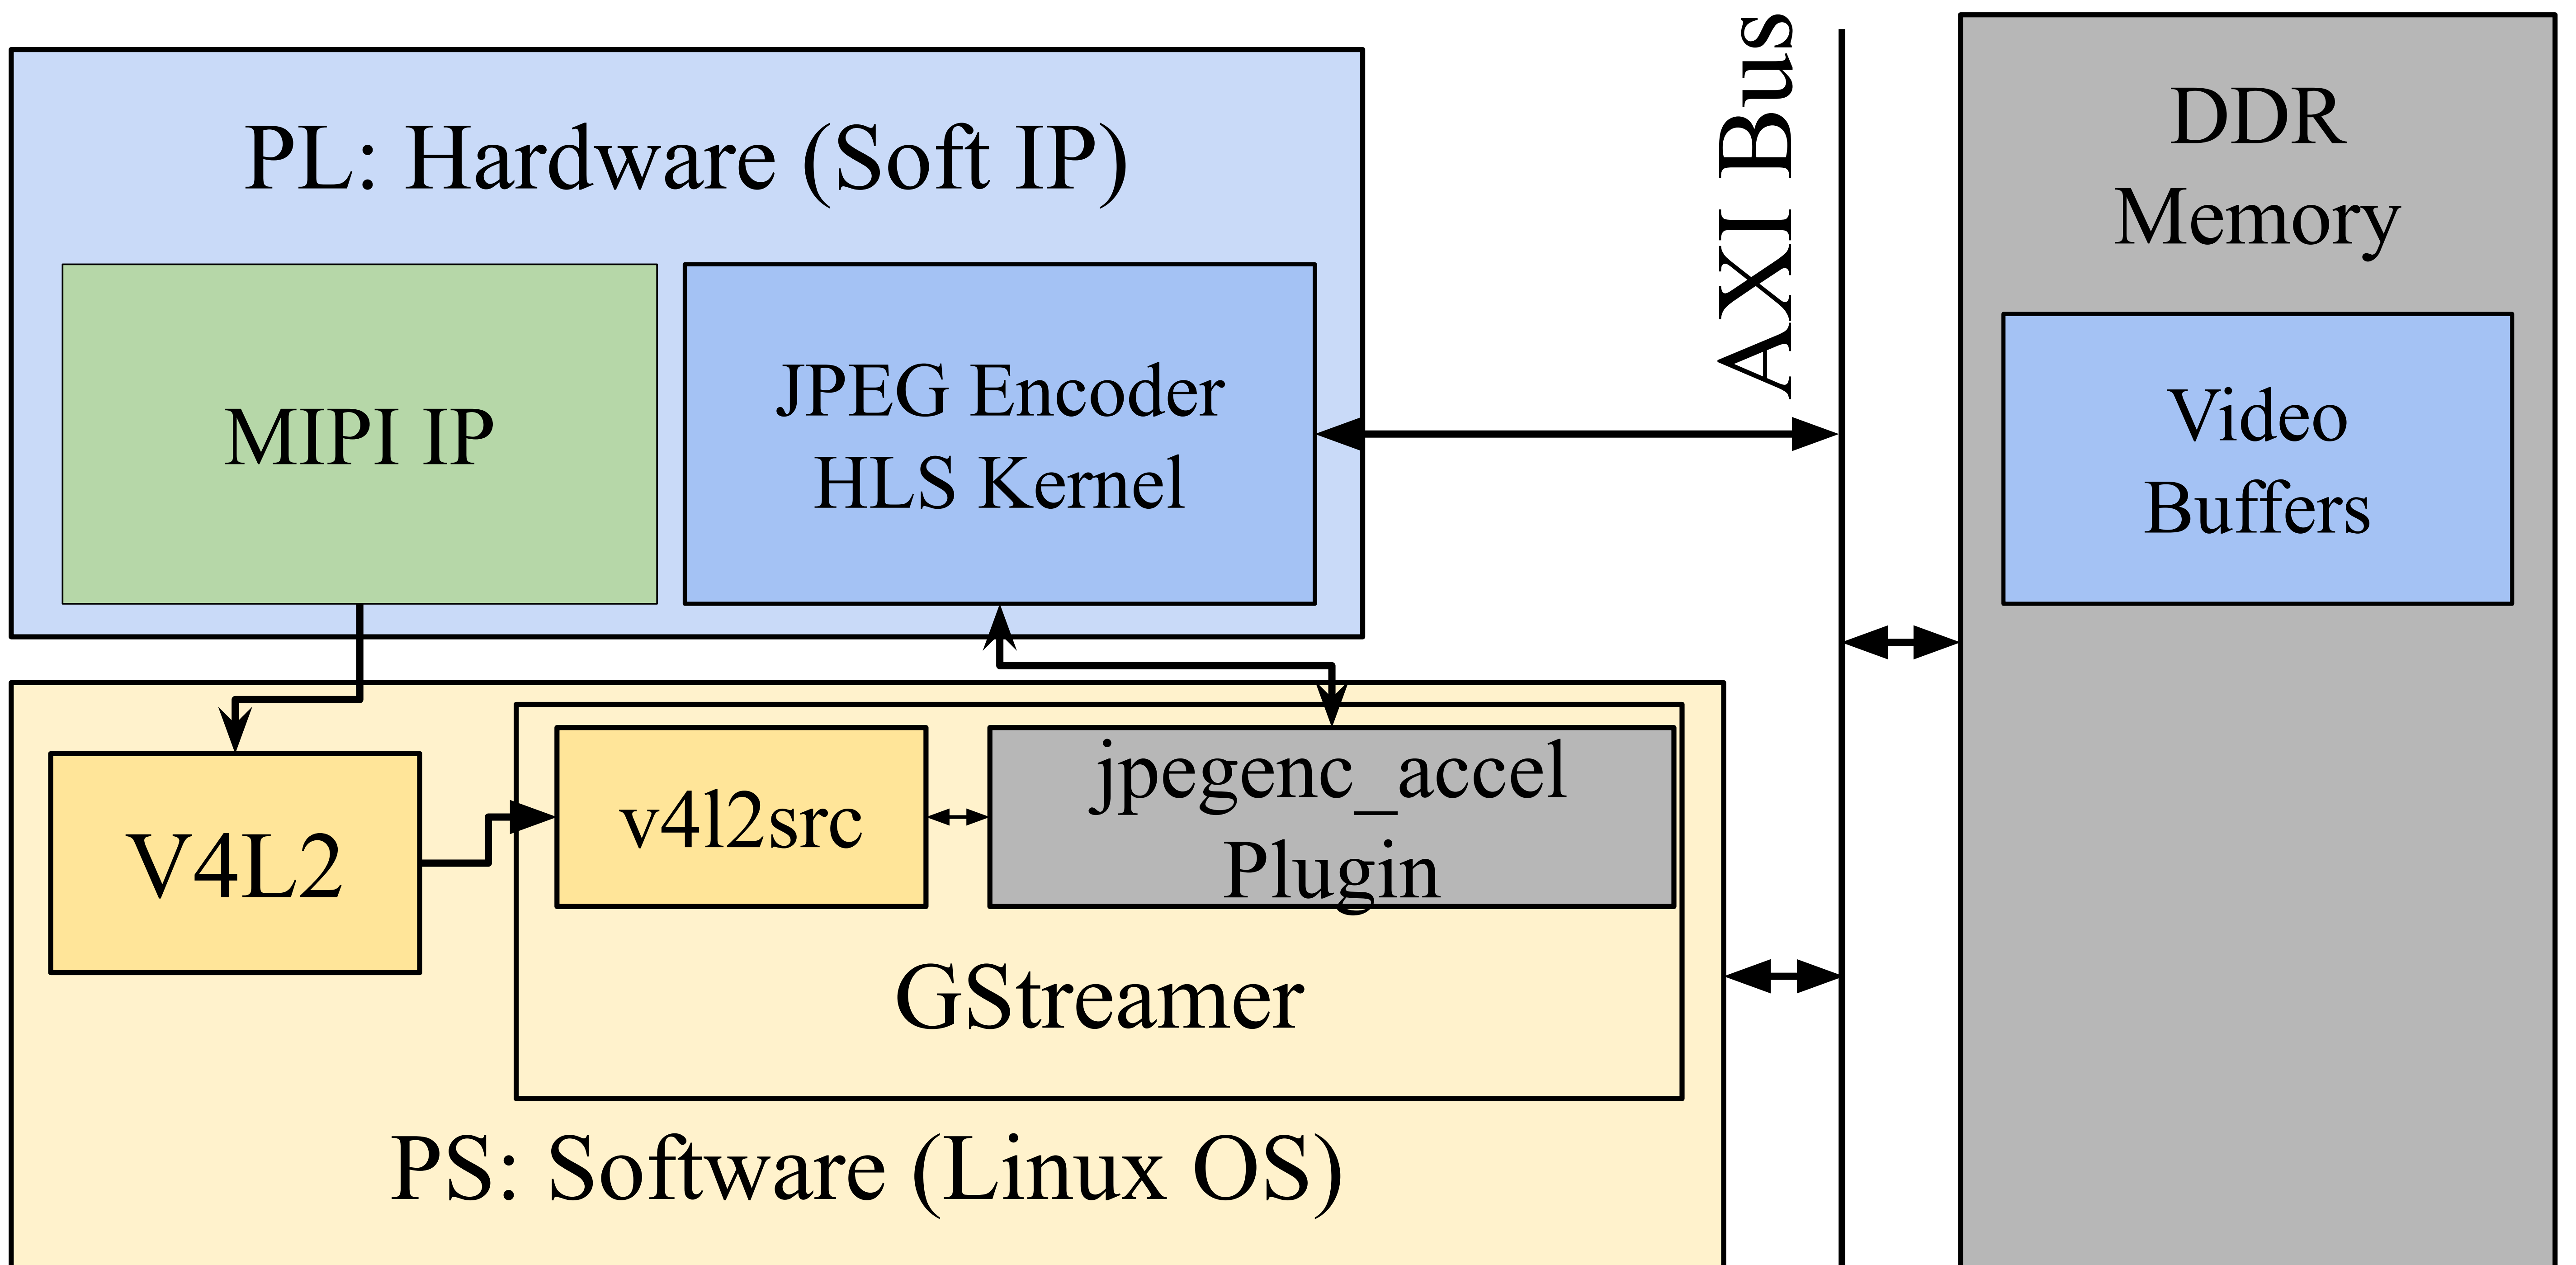


Overall Architecture

The design consists of a HW component integrated with a *GStreamer* plugin.



GStreamer allows for easy and quick composition of media pipelines, making the design easily testable in a wide range of scenarios.



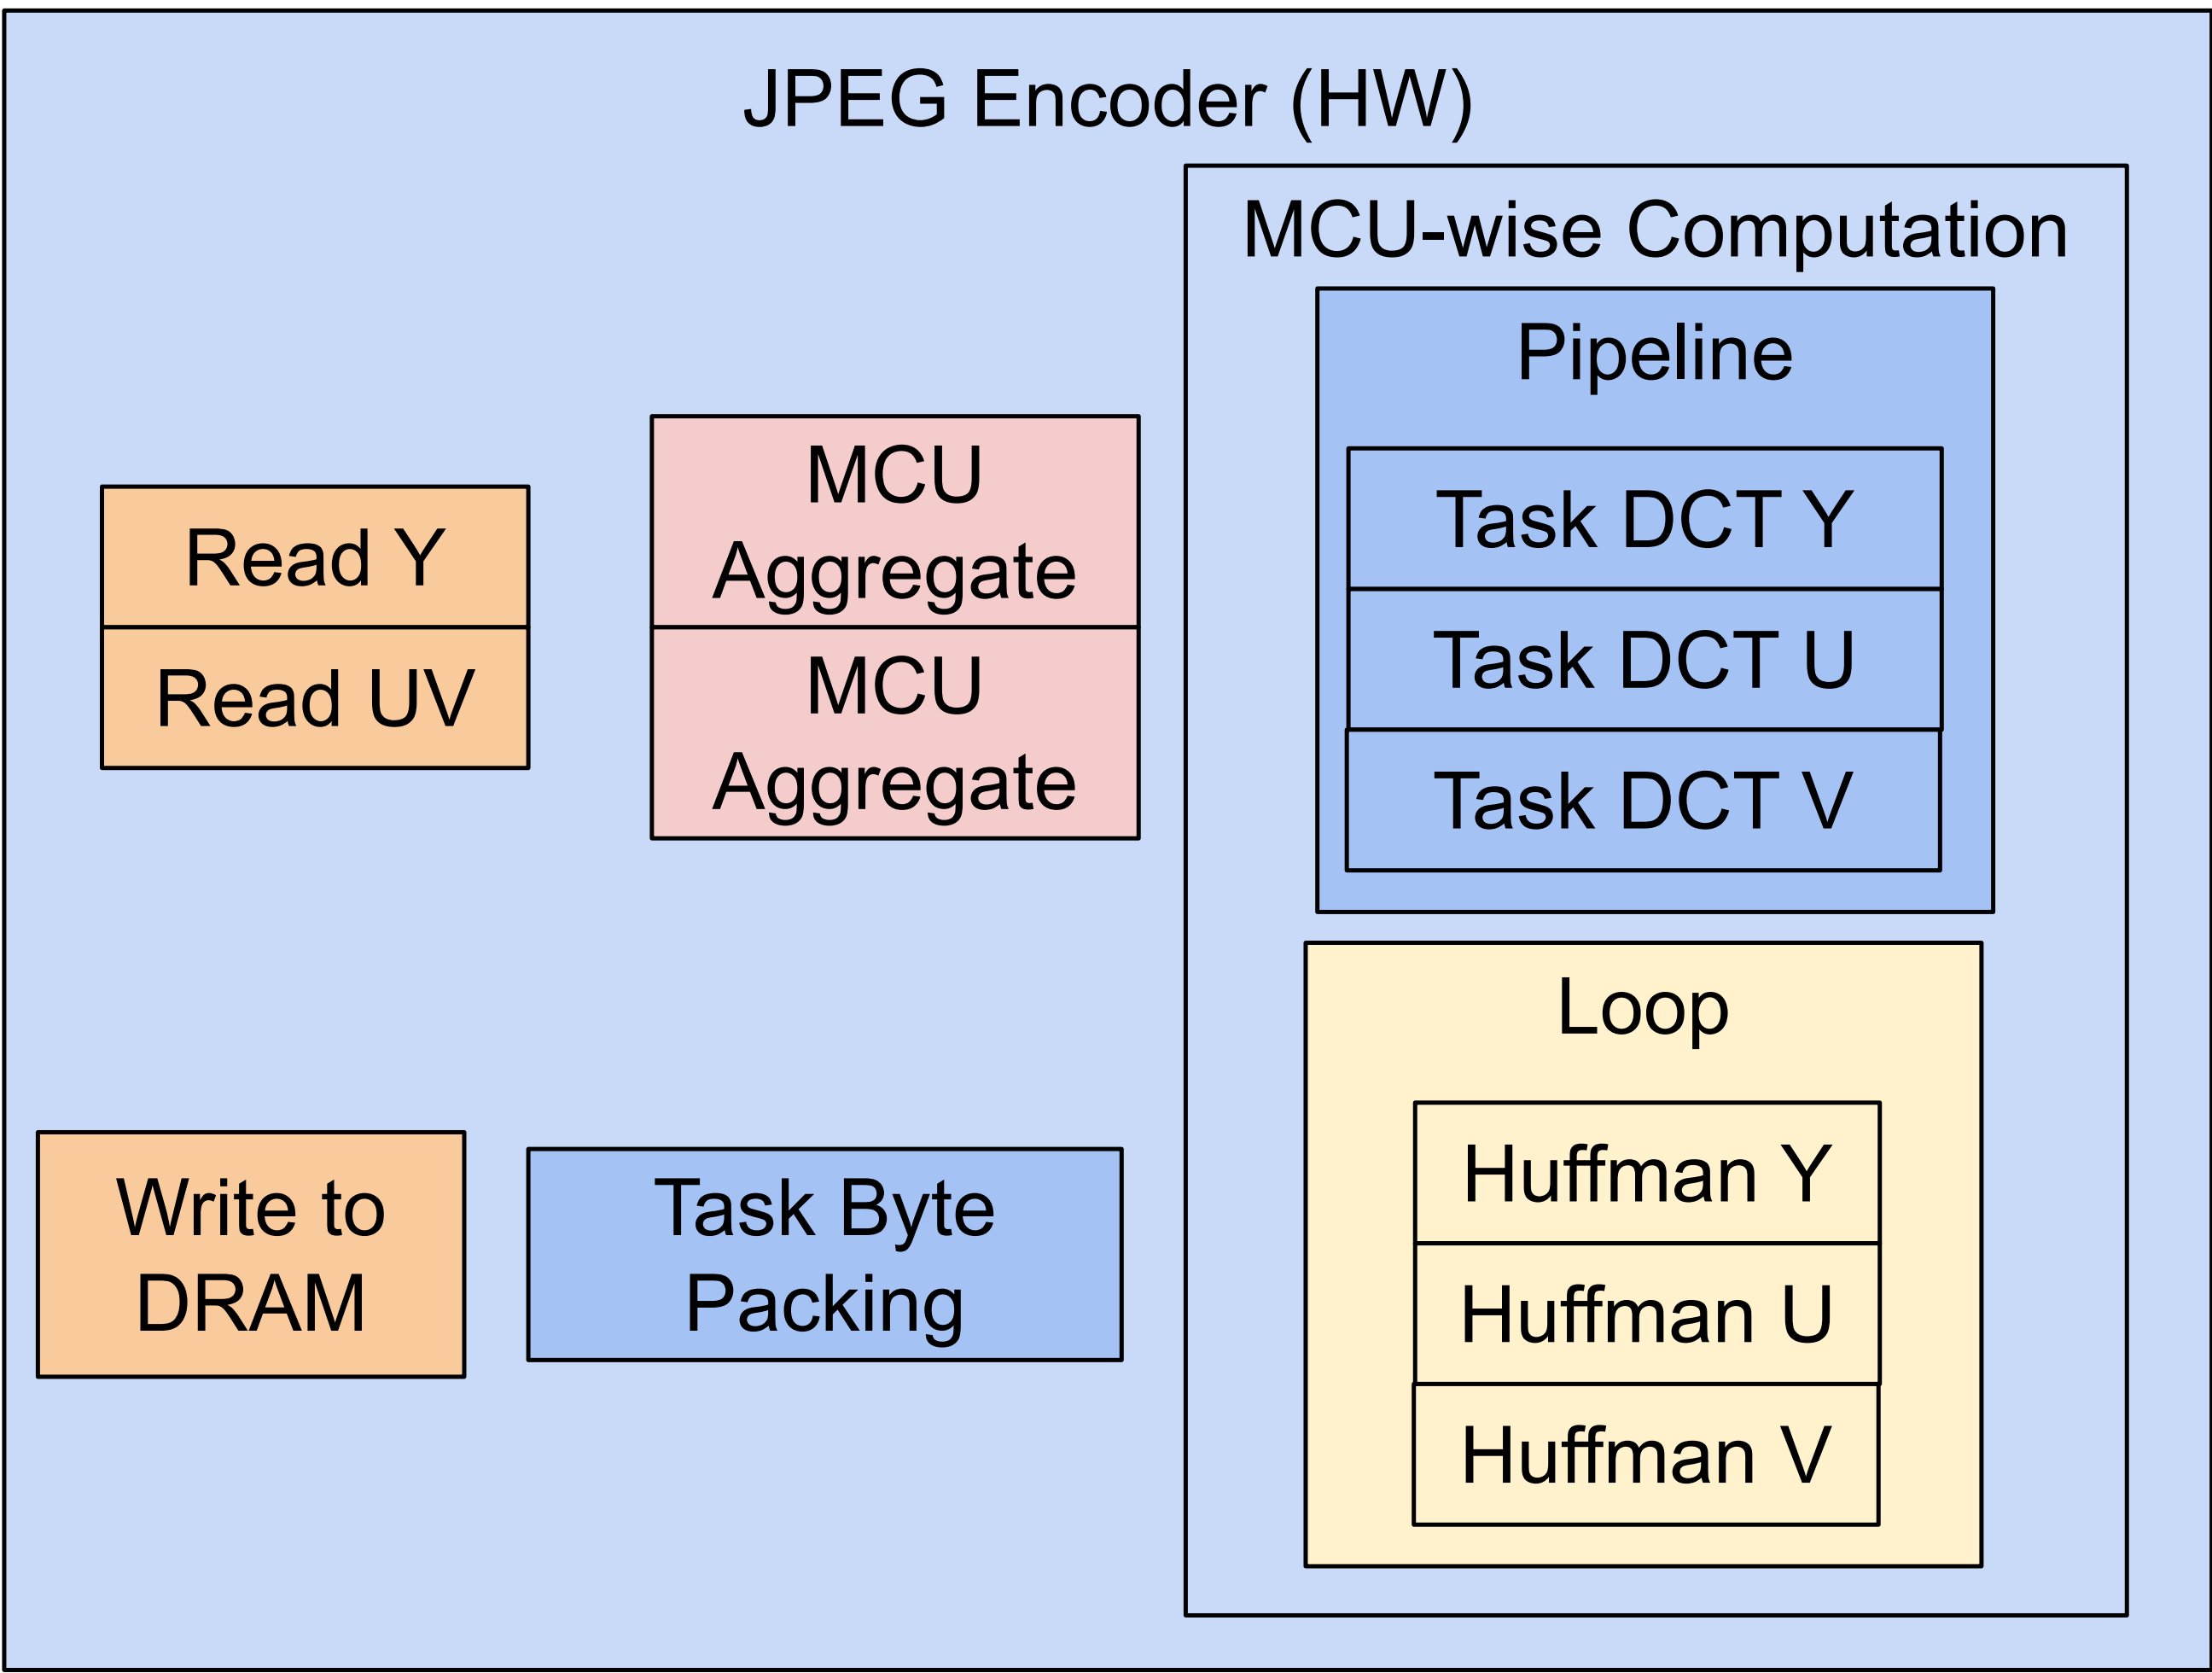
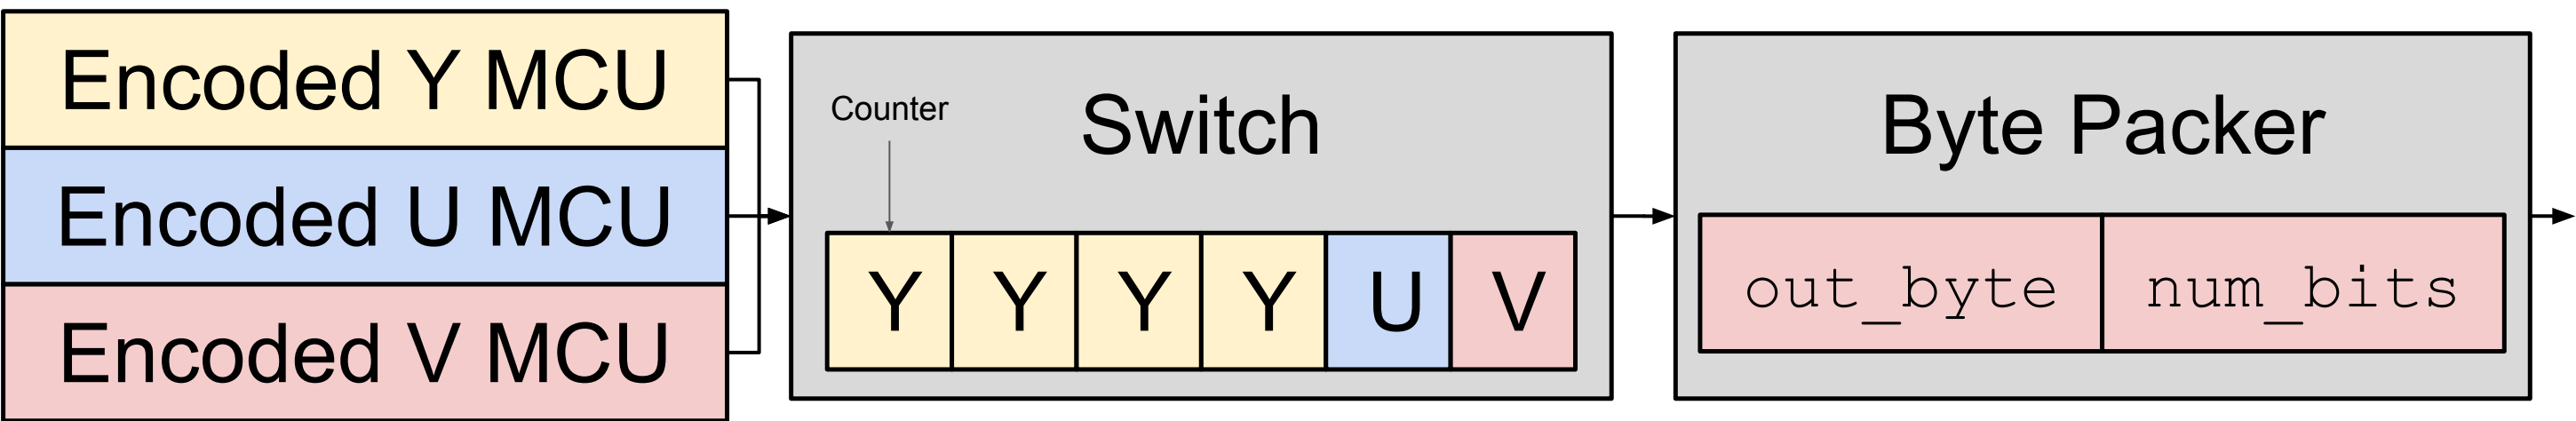
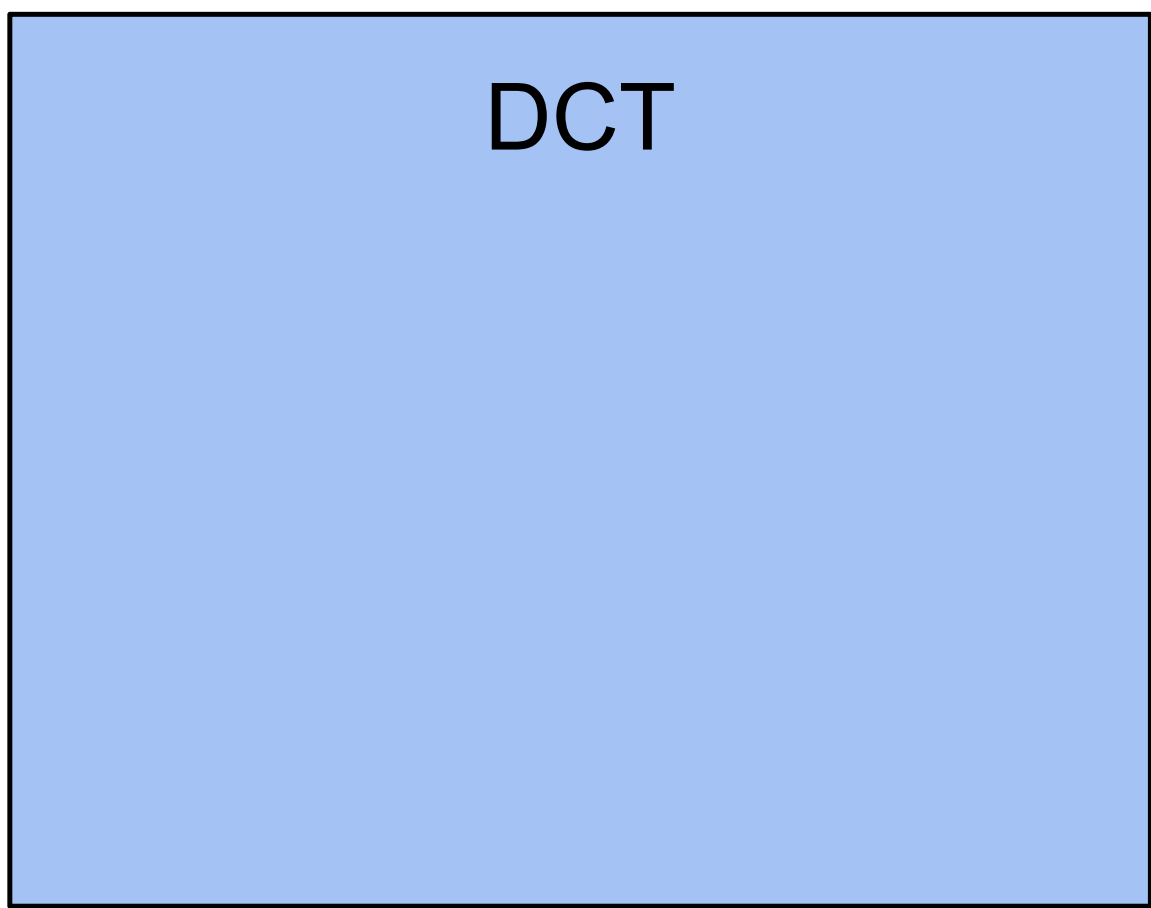
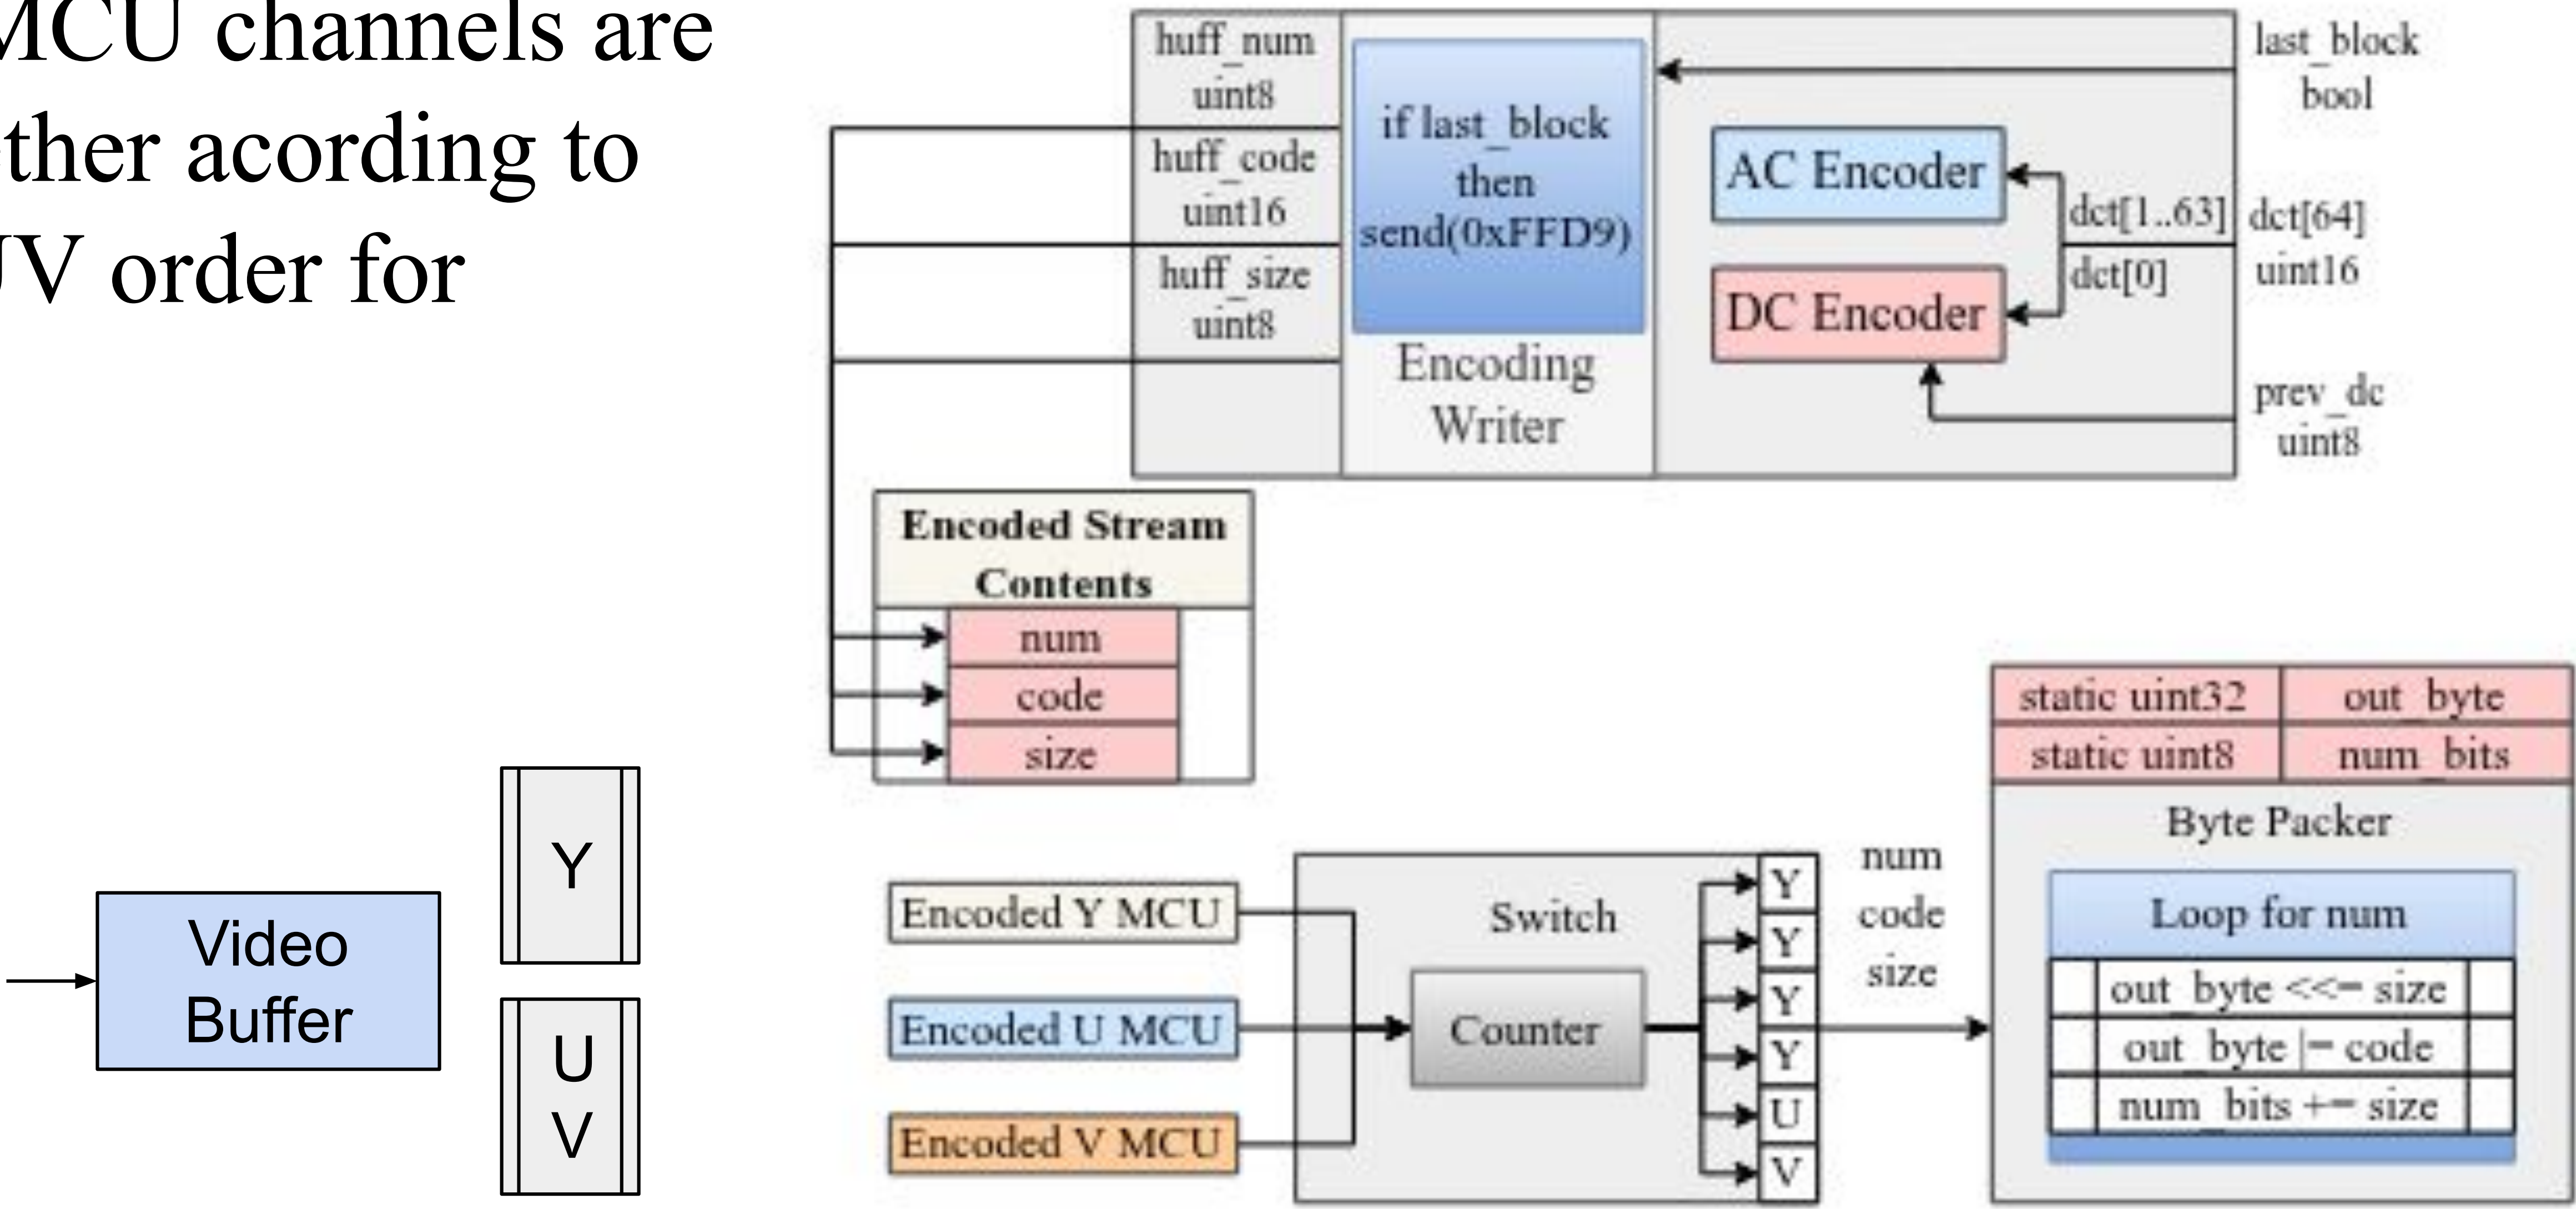
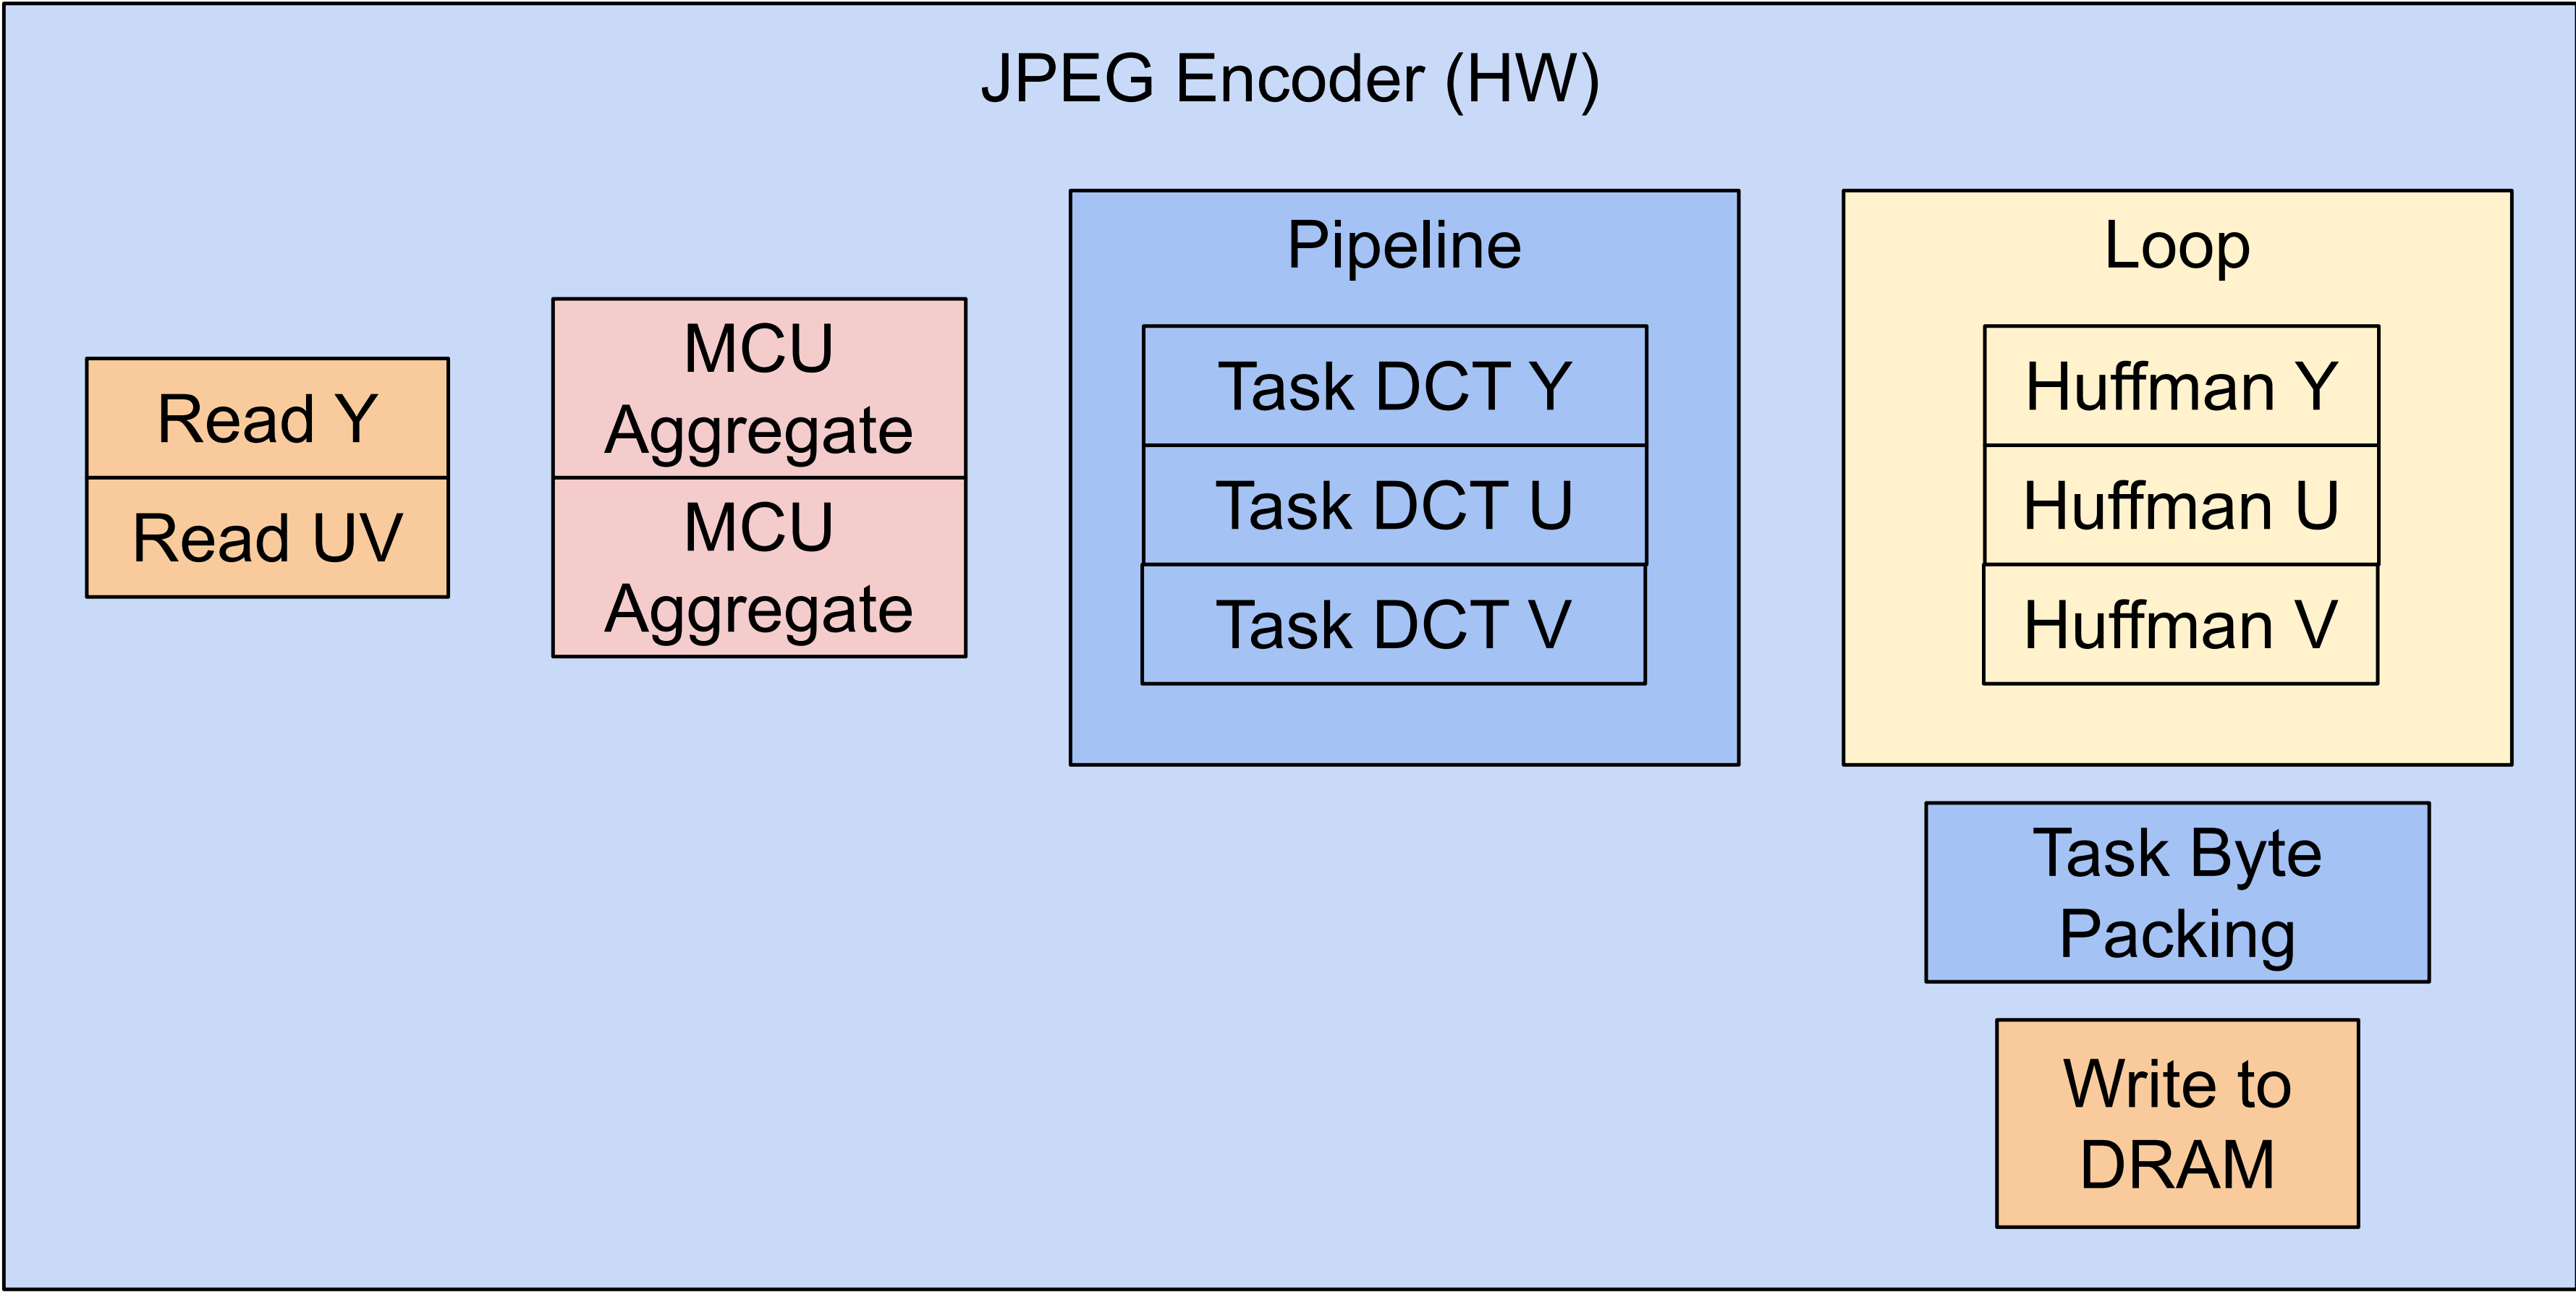
The camera-based pipeline consists of an AMD MIPI IP accessed through the v4l2src GStreamer plugin node. Separate buffer planes can be accessed directly and processed as individual channels on the JPEG IP, improving throughput.

Hardware Components

Y and UV channels are processed separately in order to aggregate MCUs into streams.

Afterwards, each MCU goes through DCT, RLE, Huffman encoding.

Finally, all MCU channels are merged together according to the YYYUYV order for YUV420.



Experiments

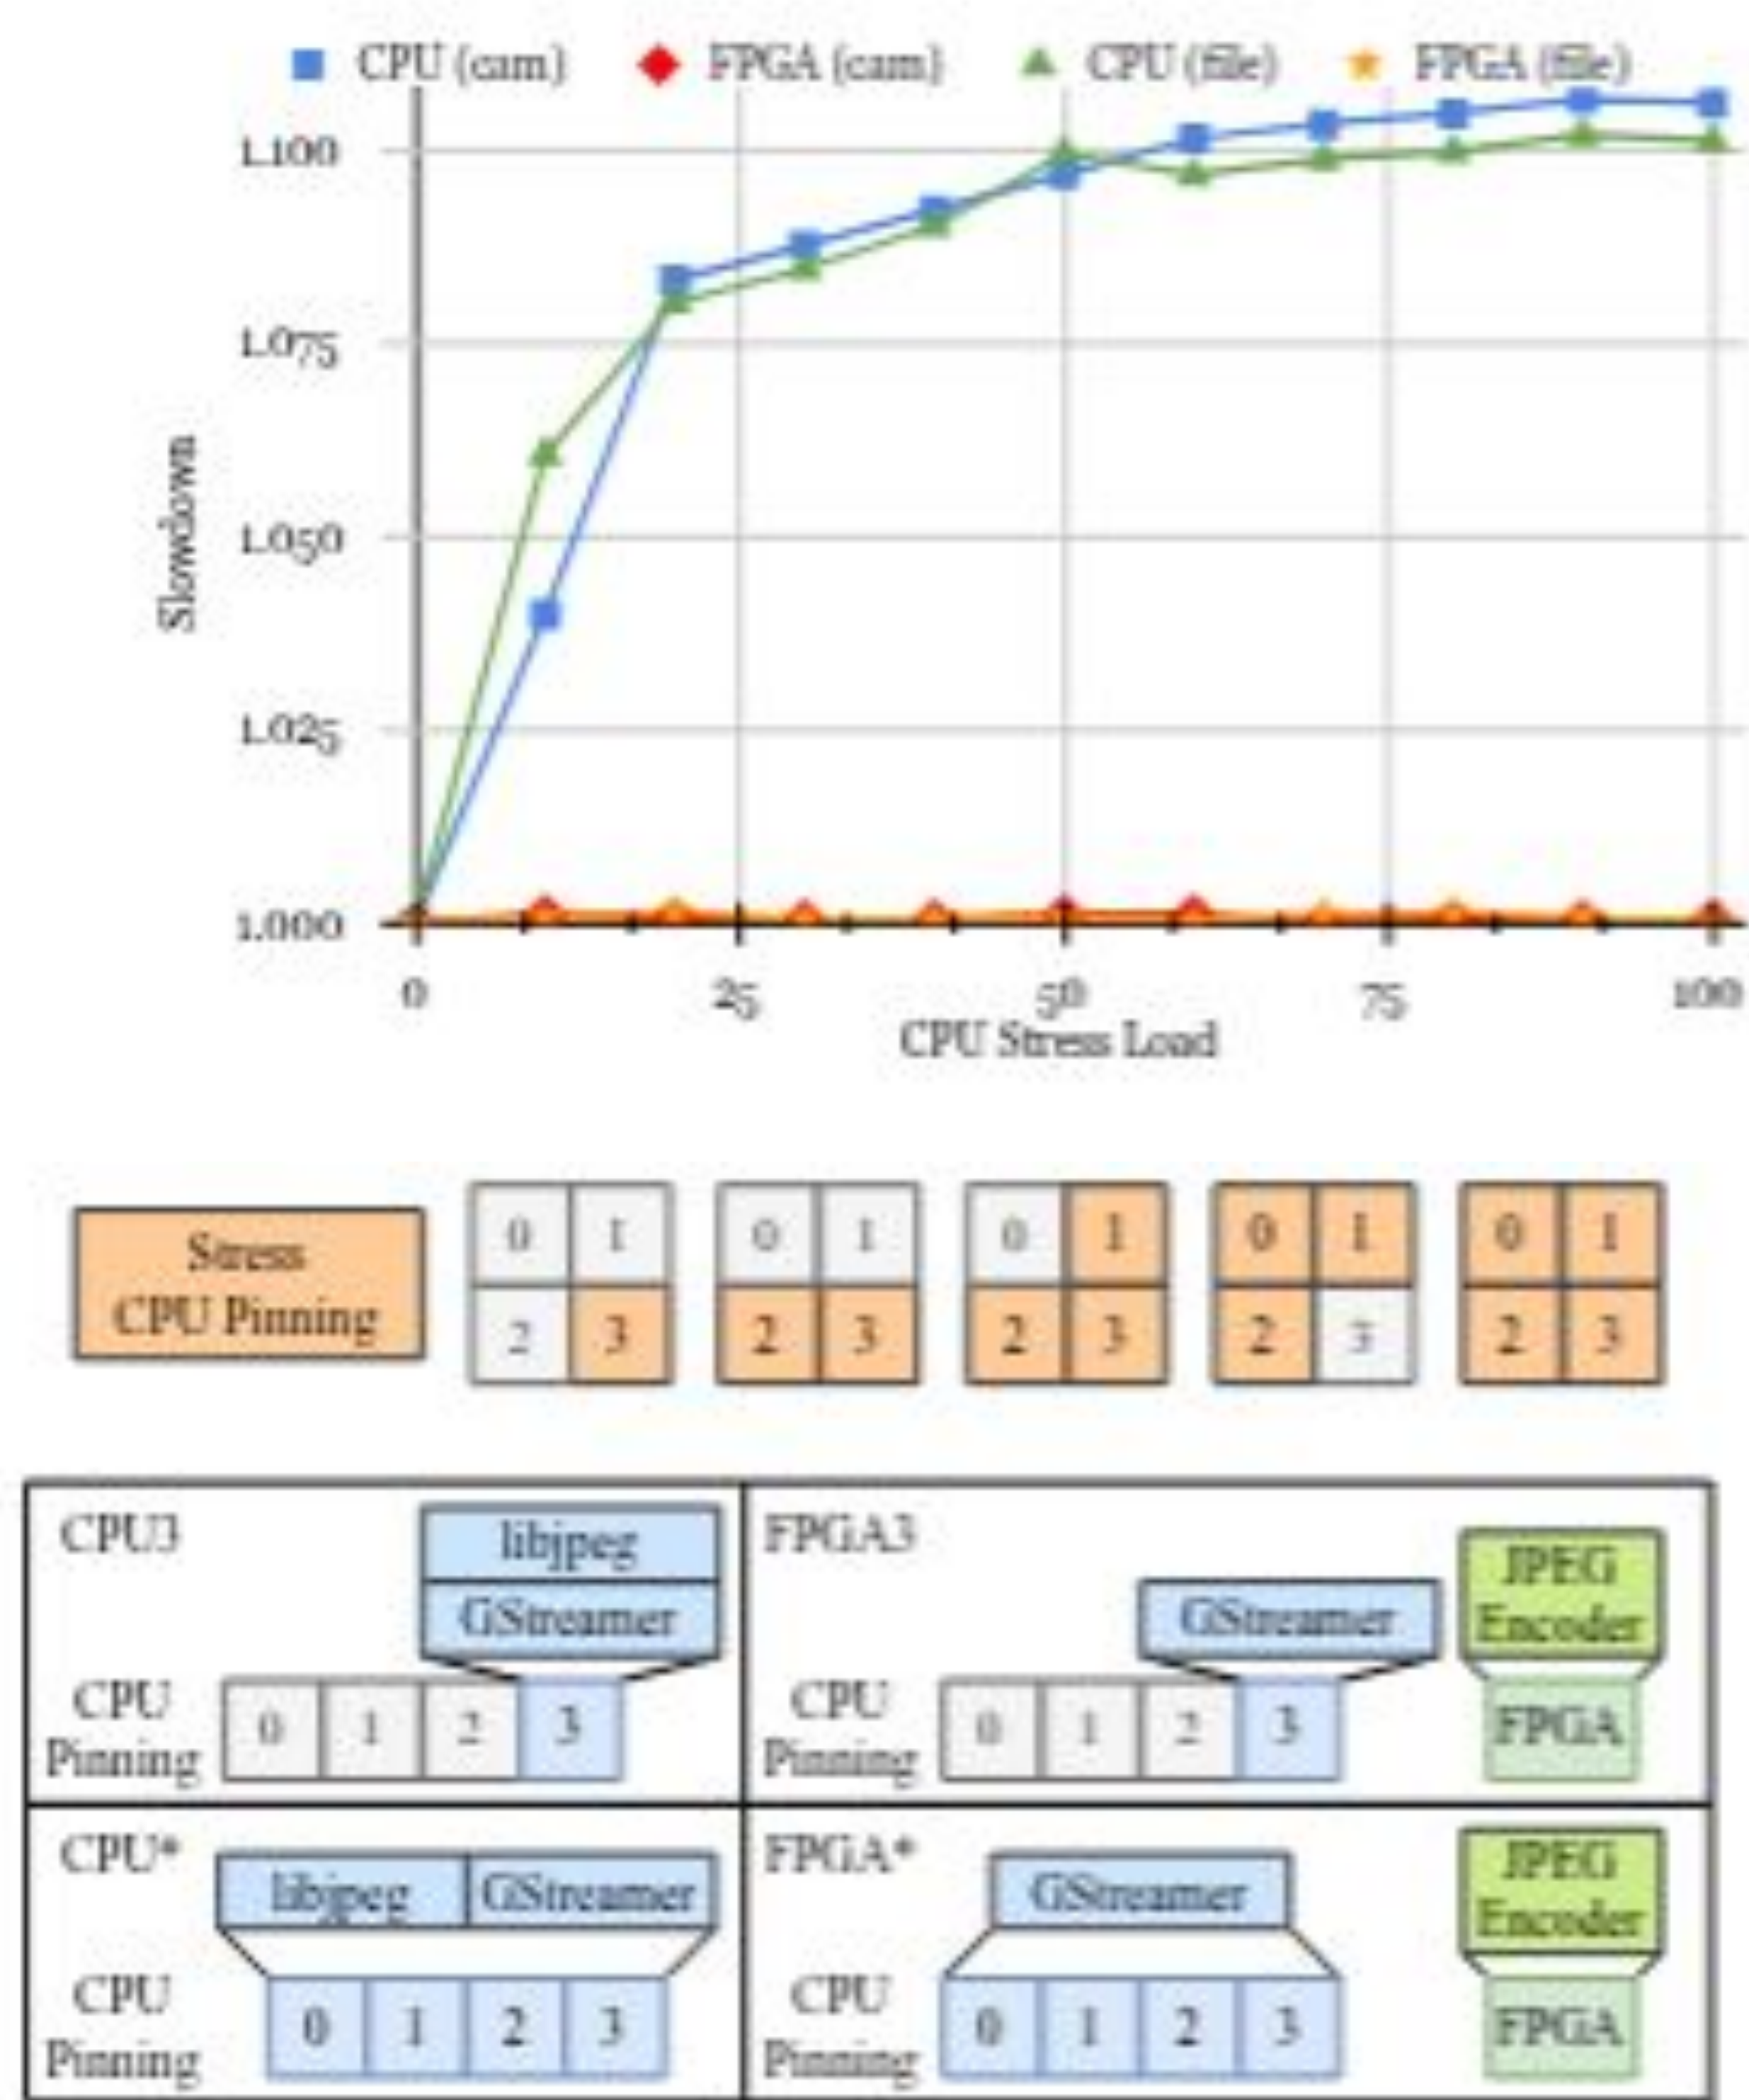


Fig. 6. **Memory Stress Configuration:** Four configurations for single-core or multi-core affinity of the pipeline (CPU3, CPU*, FPGA3, FPGA*), as well as five CPU affinity configurations for the memory stress.

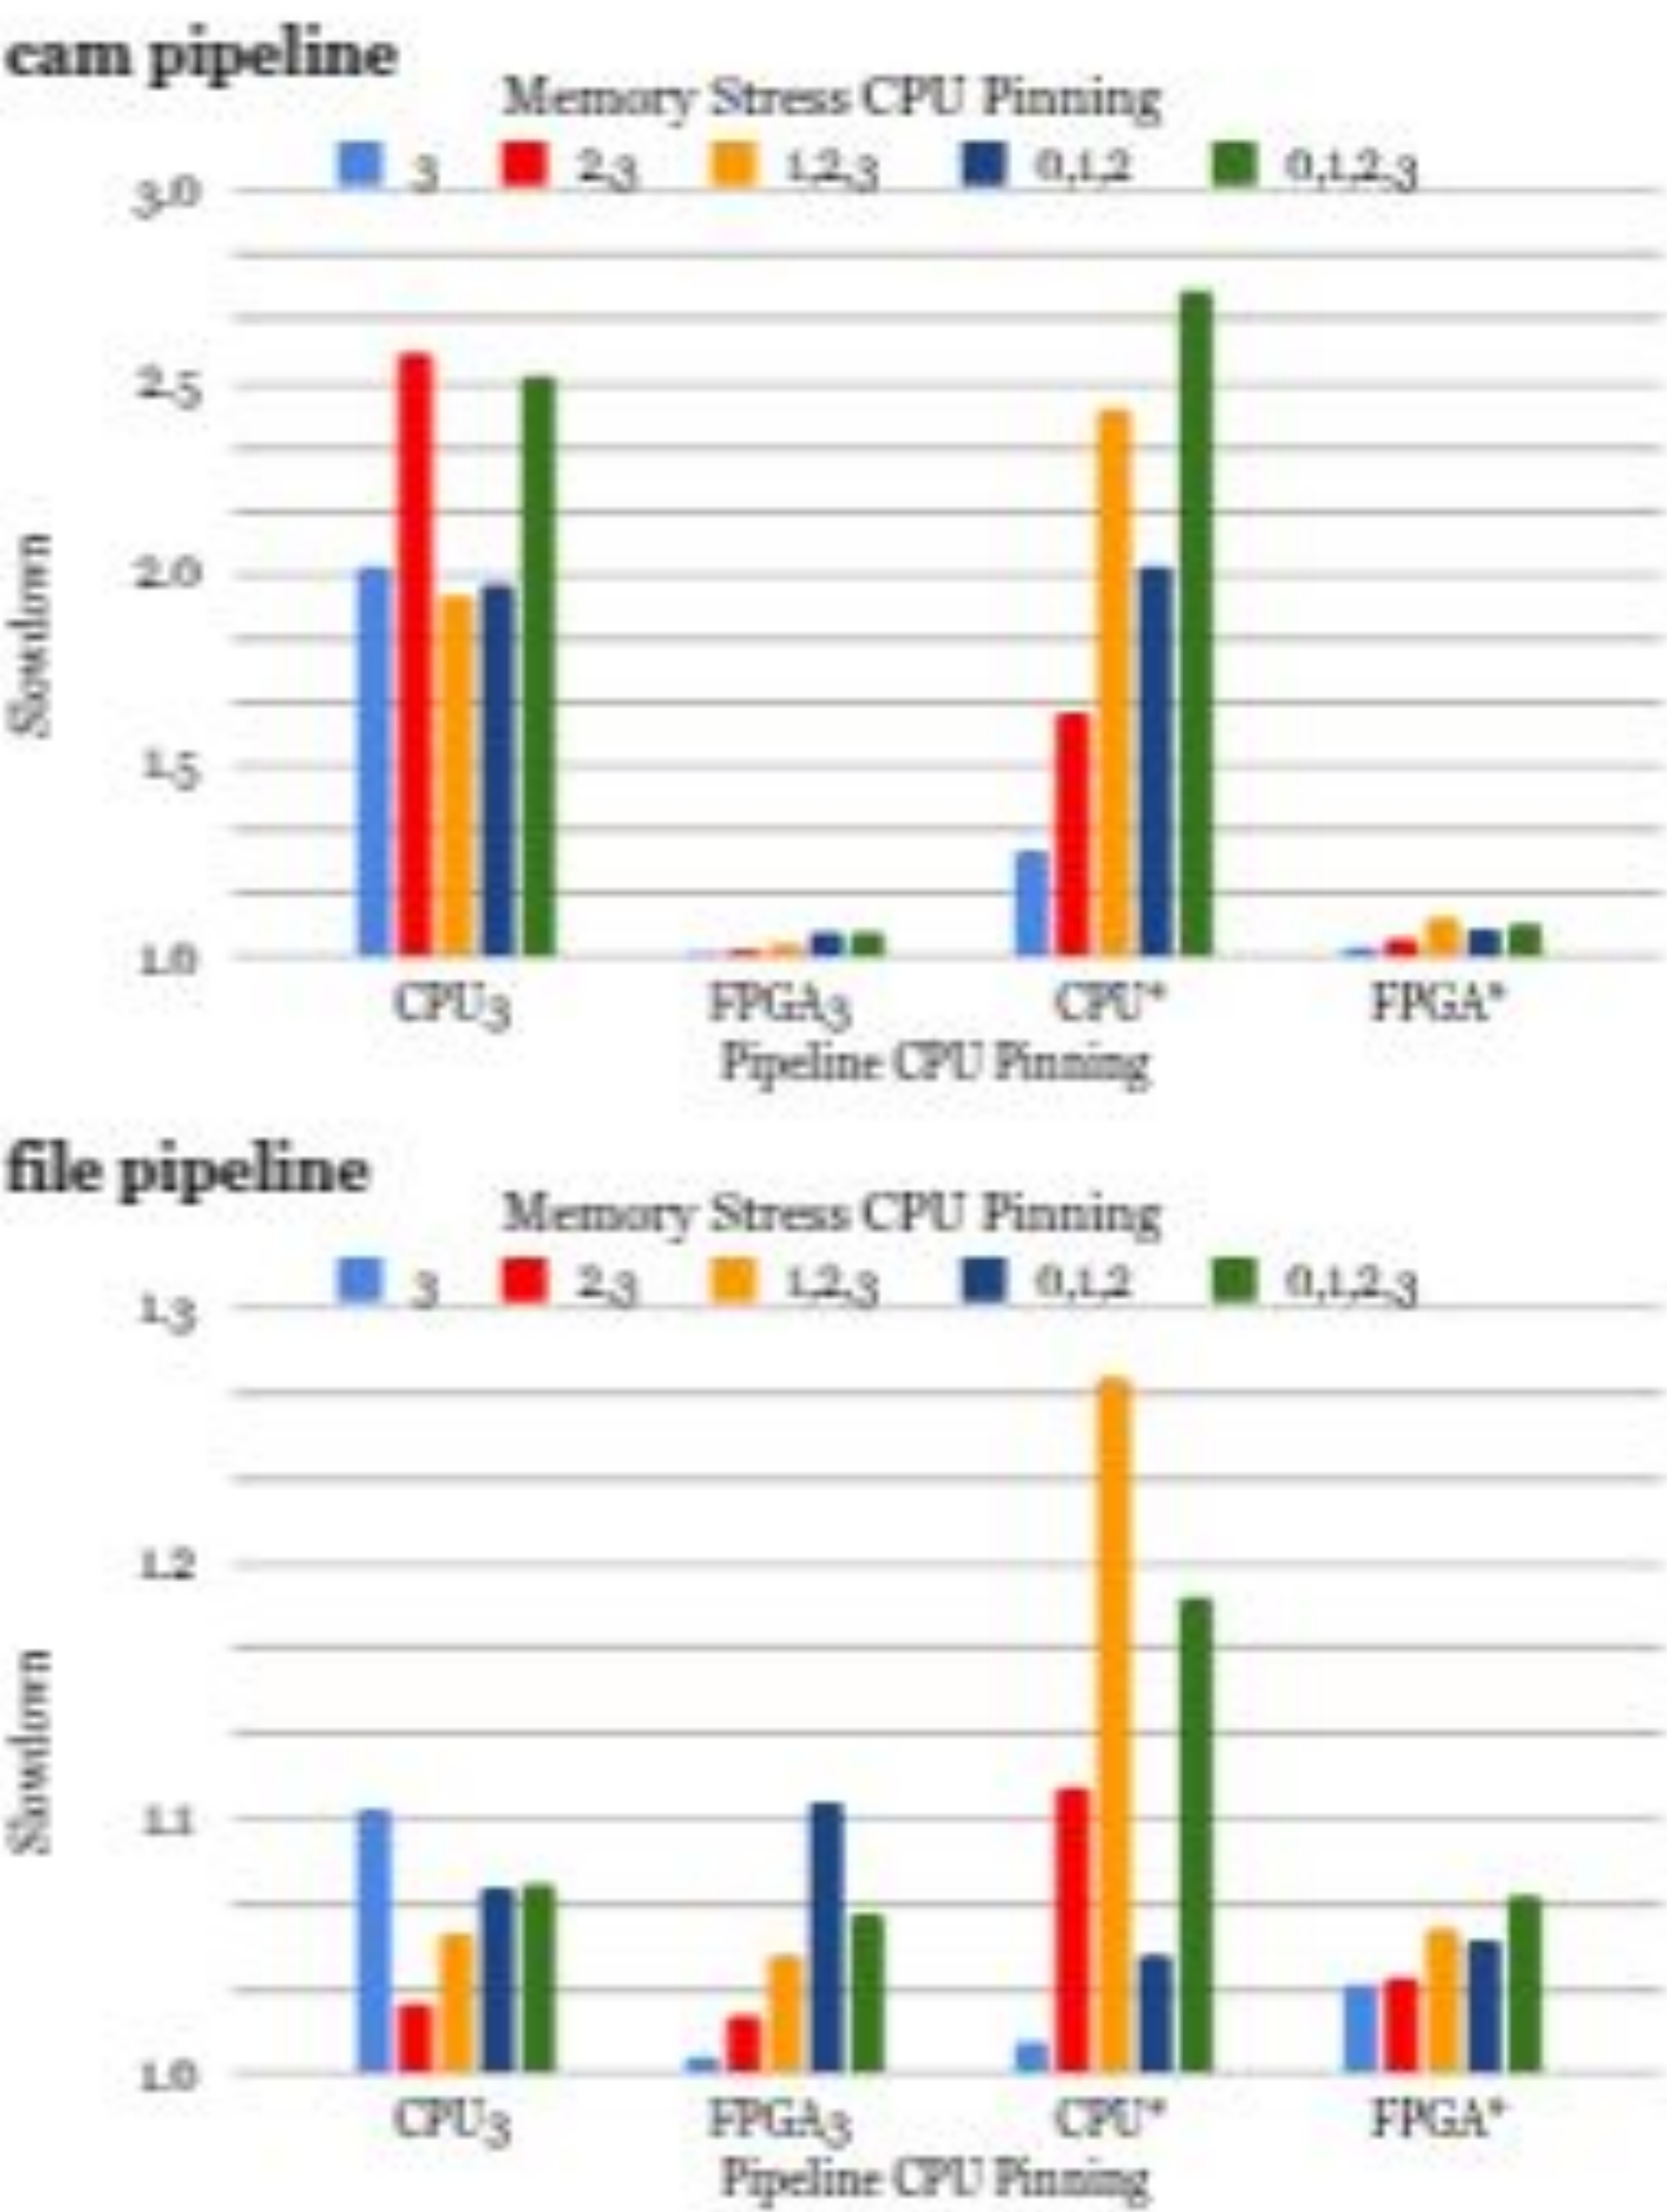


TABLE II
RESOURCE UTILIZATION COMPARISON OF
JPEG AND WebP ENCODERS ON KV260

	JPEG (our work)	WebP [3]	Total Available
Frequency	250 MHz	100 MHz	
Throughput	4.49 MPps	1.88 MPps	
LUT	11,513	62,296	92,832
LUTsMem	2,355	6,610	55,336
REG	18,150	71,227	191,197
BRAM	6	81	106
URAM	12	46	56
DSP	85	834	1248

TABLE III
PERFORMANCE EVALUATION OF THE JPEG ENCODER

		CPU	FPGA (Our)
cam	Throughput (fps)	2.6	19.5
	CPU Usage (%)	93	5
file	Throughput (fps)	22.6	21.7
	CPU Usage (%)	107	14
	Compression Ratio	3.202%	3.219%
	PSNR (dB)	27.72	27.45

